

# H<sub>2</sub>O on Hadoop

**[ H<sub>2</sub>O – The Open Source In-Memory  
Prediction Engine for Big Data ]**

Workshop Café, San Francisco, Dec. 12, 2013

Tom Kraljevic

## Outline

H<sub>2</sub>O as a Standalone HDFS client

H<sub>2</sub>O on Hadoop

Configuration settings

Stuff we learned

Questions

Using H<sub>2</sub>O as an HDFS client

## Why

I want to do run a Generalized Linear Model, and my data lives in HDFS.

## What

Use H<sub>2</sub>O standalone.

## Why


Let's also use the CPUs and Memory of my Hadoop cluster!

## What

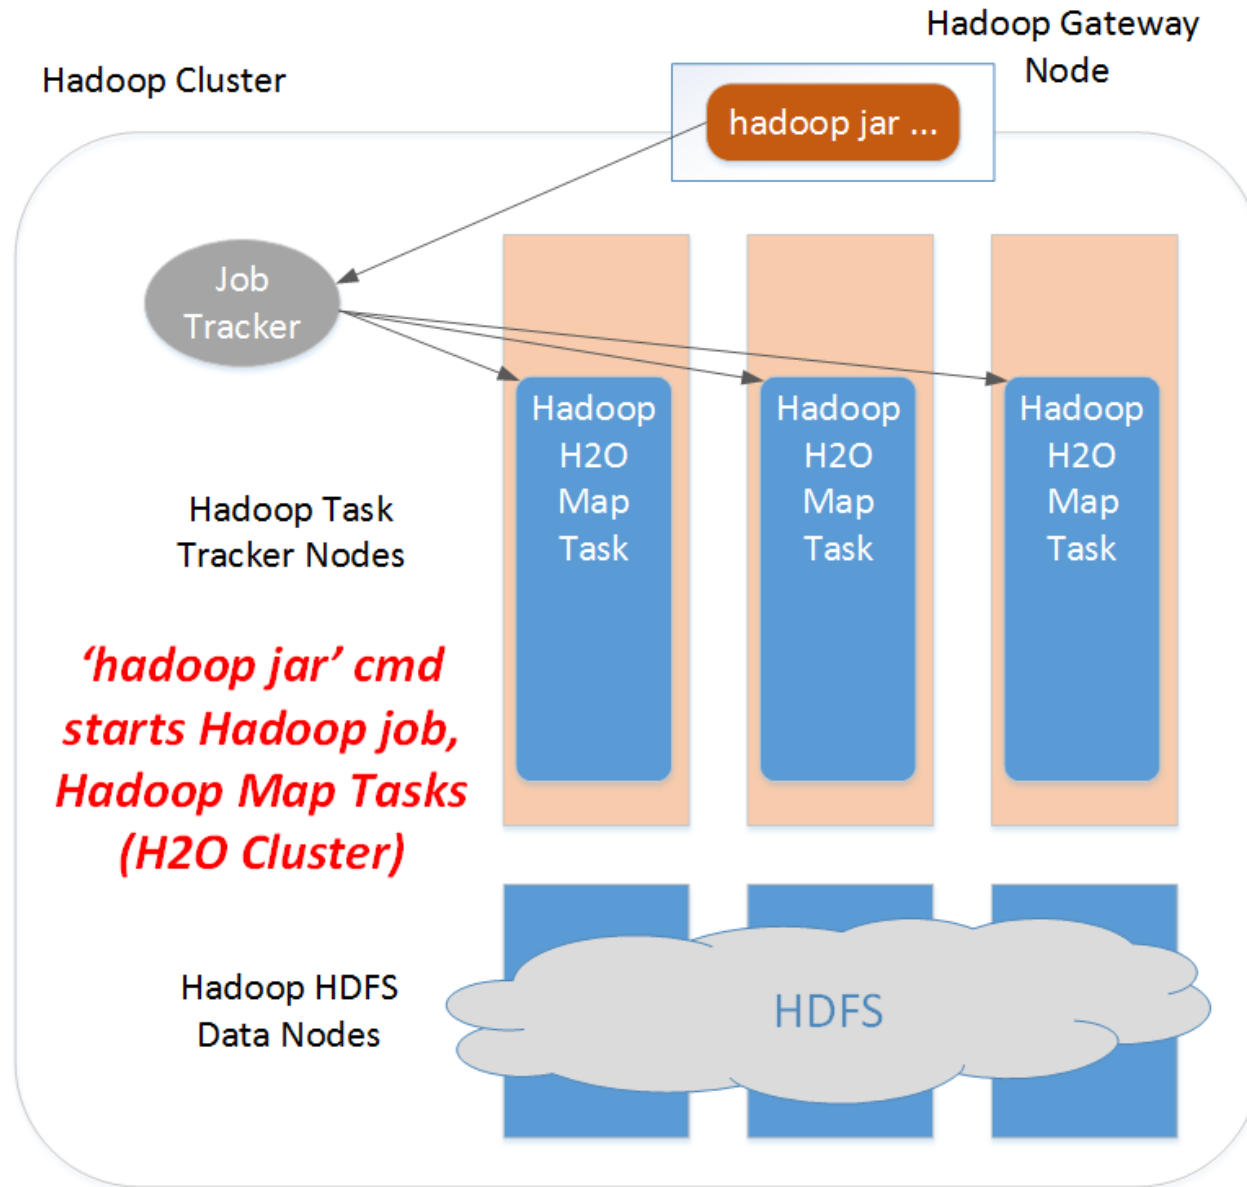
Run H<sub>2</sub>O on Hadoop.

**Oxdata**


## H<sub>2</sub>O Standalone Deployment Using HDFS




## H<sub>2</sub>O on Hadoop Deployment



## H<sub>2</sub>O on Hadoop Deployment



## H<sub>2</sub>O on Hadoop Deployment




## Standalone Mode Command-Line Invocation

```
h2o-2.1.0.1144 -- ubuntu@ip-10-07-32-113: /var/www/0xdata-website-files -- java -- 125x61
mbp2:Downloads tomk$ unzip h2o-2.1.0.1144.zip
Archive: h2o-2.1.0.1144.zip
  creating: h2o-2.1.0.1144/
  inflating: h2o-2.1.0.1144/h2o-sources.jar
  creating: h2o-2.1.0.1144/hadoop/
  inflating: h2o-2.1.0.1144/hadoop/h2odriver_mapr2.1.3.jar
  inflating: h2o-2.1.0.1144/hadoop/h2odriver_cdh4_yarn.jar
  inflating: h2o-2.1.0.1144/hadoop/README.txt
  inflating: h2o-2.1.0.1144/hadoop/h2odriver_cdh3.jar
  inflating: h2o-2.1.0.1144/hadoop/h2odriver_hdp1.3.2.jar
  inflating: h2o-2.1.0.1144/hadoop/h2odriver_cdh4.jar
  creating: h2o-2.1.0.1144/ec2/
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-distribute-h2o.sh
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-distribute-flatfile.sh
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-print-info.py
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-launch-instances.py
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-test-ssh.sh
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-distribute-aws-credentials.sh
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-download-h2o.sh
  inflating: h2o-2.1.0.1144/ec2/README.txt
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-stop-h2o.sh
  inflating: h2o-2.1.0.1144/ec2/h2o-cluster-start-h2o.sh
  inflating: h2o-2.1.0.1144/README.txt
  creating: h2o-2.1.0.1144/R/
  inflating: h2o-2.1.0.1144/R/README.txt
  inflating: h2o-2.1.0.1144/h2o.jar
  inflating: h2o-2.1.0.1144/LICENSE.txt
mbp2:Downloads tomk$ cd h2o-2.1.0.1144
mbp2:h2o-2.1.0.1144 tomk$ java -Xmx4g -jar h2o.jar
12:05:04.781 main INFO WATER: ----- H2O started -----
12:05:04.783 main INFO WATER: Build git hash: c206b694ba07e7da1b5ec97bac312ec30bf5b2d2
12:05:04.783 main INFO WATER: Build git describe: nn-2-1489-gc206b69
12:05:04.783 main INFO WATER: Build project version: 2.1.0.1144
12:05:04.783 main INFO WATER: Built by: 'jenkins'
12:05:04.783 main INFO WATER: Built on: 'Tue Dec 10 18:21:10 PST 2013'
12:05:04.783 main INFO WATER: Java availableProcessors: 8
12:05:04.786 main INFO WATER: Java heap totalMemory: 0.08 gb
12:05:04.786 main INFO WATER: Java heap maxMemory: 3.98 gb
12:05:04.786 main INFO WATER: Java version: Java 1.6.0_65 (from Apple Inc.)
12:05:04.786 main INFO WATER: OS version: Mac OS X 10.8.5 (x86_64)
12:05:04.787 main INFO WATER: ICE root: '/tmp/h2o-tomk'
12:05:04.790 main INFO WATER: Possible IP Address: vmnet8 (vmnet8), 172.16.175.1
12:05:04.790 main INFO WATER: Possible IP Address: vmnet1 (vmnet1), 192.168.183.1
12:05:04.790 main INFO WATER: Possible IP Address: en0 (en0), fe80:0:0:0:2acf:e9ff:fe1c:ccf%4
12:05:04.790 main INFO WATER: Possible IP Address: en0 (en0), 192.168.1.30
12:05:04.791 main INFO WATER: Possible IP Address: lo0 (lo0), 0:0:0:0:0:0:1
12:05:04.791 main INFO WATER: Possible IP Address: lo0 (lo0), fe80:0:0:0:0:0:0:1%1
12:05:04.791 main INFO WATER: Possible IP Address: lo0 (lo0), 127.0.0.1
12:05:04.824 main WARN WATER: Multiple local IPs detected:
+ /172.16.175.1 /192.168.183.1 /192.168.1.30
+ Attempting to determine correct address...
+ Using /192.168.1.30
12:05:04.876 main INFO WATER: Internal communication uses port: 54322
+ Listening for HTTP and REST traffic on http://192.168.1.30:54321/
12:05:04.906 main INFO WATER: H2O cloud name: 'tomk'
12:05:04.907 main INFO WATER: (v2.1.0.1144) 'tomk' on /192.168.1.30:54321, discovery address /225.54.105.89:57654
12:05:04.909 main INFO WATER: Cloud of size 1 formed [/192.168.1.30:54321]
12:05:04.909 main INFO WATER: Log dir: '/tmp/h2o-tomk/h2ologs'
```

unzip

java -jar





## H2O on Hadoop Client Invocation

```
tomk@mr-0xb1:~/h2o-2.1.0.1144/hadoop$ hadoop jar h2odriver_cdh4.jar water.hadoop.h2odriver  
-libjars ../h2o.jar -mapperXmx 30g -nodes 3 -output hdfsOutDir
```

```
Determining driver host interface for mapper->driver callback...
```

```
[Possible callback IP address: 192.168.1.161]
```

```
[Possible callback IP address: 127.0.0.1]
```

```
Using mapper->driver callback IP address and port: 192.168.1.161:60576
```

```
(You can override these with -driverif and -driverport.)
```

```
Driver program compiled with MapReduce V1 (Classic)
```

```
13/12/12 13:25:29 WARN conf.Configuration: mapred.map.child.java.opts is deprecated. Instead, use  
mapreduce.map.java.opts
```

```
Memory Settings:
```

```
mapred.child.java.opts: -Xms30g -Xmx30g
```

```
mapred.map.child.java.opts: -Xms30g -Xmx30g
```

```
Extra memory percent: 10
```

```
mapreduce.map.memory.mb: 33792
```

```
Job name 'H2O_49792' submitted
```

```
JobTracker job ID is 'job_1386713862878_0002'
```

```
Waiting for H2O cluster to come up...
```

```
H2O node 192.168.1.163:54321 requested flatfile
```

```
H2O node 192.168.1.161:54323 requested flatfile
```

```
H2O node 192.168.1.162:54339 requested flatfile
```

```
Sending flatfiles to nodes...
```

```
[Sending flatfile to node 192.168.1.163:54321]
```

```
[Sending flatfile to node 192.168.1.161:54323]
```

```
[Sending flatfile to node 192.168.1.162:54339]
```

```
H2O node 192.168.1.161:54323 reports H2O cluster size 1
```

```
H2O node 192.168.1.163:54321 reports H2O cluster size 1
```

```
H2O node 192.168.1.162:54339 reports H2O cluster size 1
```

```
H2O node 192.168.1.162:54339 reports H2O cluster size 3
```

```
H2O node 192.168.1.163:54321 reports H2O cluster size 3
```

```
H2O node 192.168.1.161:54323 reports H2O cluster size 3
```

```
H2O cluster (3 nodes) is up
```

```
(Note: Use the -disown option to exit the driver after cluster formation)
```

```
(Press Ctrl-C to kill the cluster)
```

```
Blocking until the H2O cluster shuts down...
```

H<sub>2</sub>O is not the typical Hadoop MapReduce job

- Long-lived Hadoop mapper tasks that take up CPUs and memory
- No Hadoop reduce tasks
- No mapper task retry (user needs to restart H<sub>2</sub>O cloud on failure)
- No HDFS “input file” to the Hadoop map phase
  - Null splits define number of Hadoop map tasks
- No HDFS “output files” produced by the Hadoop job.
- Multithreaded
- Big Java heap (e.g. 128 GB)
- Mappers talk to each other
  - Mappers need to be run at the same time

**Oxdata**

# Hadoop Configuration

## Client-side (Driver) Hadoop Configuration

mapreduce.job.ubertask.enable false

*[ mapreduceMapMemoryMb = Xmx + 10% fudge factor; ]*

mapreduce.map.memory.mb mapreduceMapMemoryMb

*[ mapChildJavaOpts = "-Xms" + mapperXmx + " -Xmx" + mapperXmx; ]*

mapred.child.java.opts mapChildJavaOpts

mapred.map.child.java.opts mapChildJavaOpts

mapreduce.map.speculative false

mapred.map.tasks.speculative.execution false

mapred.map.max.attempts 1

mapred.job.reuse.jvm.num.tasks 1

**0xdata**

## YARN Settings in CDH4

mapreduce.map.memory.mb	128 GB
mapreduce.map.java.opts.max.heap	100 GB
yarn.nodemanager.resource.memory-mb	128 GB

## ResourceManager Configuration Safety Valve for yarn-site.xml

```
<property>  
  <name>yarn.scheduler.maximum-allocation-mb</name>  
  <value>131072</value>  
</property>
```

0xdata

## Stuff We Learned

- Empty splits determine the mapper count
- Use Hadoop counters to avoid having Hadoop kill your job (which doesn't map anything) after 10 minutes
- Hook `System.exit()` to get mappers to exit properly
- YARN `/etc/alternatives` setting
- YARN memory settings
- YARN FIFO scheduler mapper clumping
- Hortonworks `-libjars` distribution unpacks the `.jar` file

Debugging a mapper process is an experience

- Finding the process is a challenge
- Finding the output is a challenge
- Convincing Hadoop not to delete the output on a failure is a challenge
  - Even if your logger setup code gets a chance to run, you still have to find the output.
- Attaching with standard JDWP debuggers is a challenge
  - Finding the process is (again) a challenge
  - Debugger doesn't understand the mapper environment; attach didn't work.
- Finally resorted to POST-code style printing techniques
  - “Log” output to /tmp, bypassing everything, so I get to keep it

## Possible Future Work

- Automated log collection upon failure
  - Fishing out logs from a cluster of 1000 Hadoop nodes is a great job for a tool.
- Driver vs. host checking of Hadoop version
  - Helps avoid subtle errors
- Custom YARN Application Master
- Experiment further with I/O optimization for H<sub>2</sub>O Node placement based on data locality

***Your Contributions Welcome***



## Key Takeaways

- H2O is the premier platform for *In-Memory* Predictive Analytics on Big Data
  - Open Source Apache 2 license
  - H<sub>2</sub>O mapreduce (in-memory) is not Hadoop mapreduce (HDFS)!
- H<sub>2</sub>O can read your data from HDFS
  - Standalone or running on Hadoop
- Run H<sub>2</sub>O on Hadoop to use CPUs and Memory from your existing Hadoop cluster
  - Get started easily with the gear you've already got

H<sub>2</sub>O on Hadoop in the Wild



K-Means on a Terabyte of data (Major Insurance Co.)

**Now Certified** on favorite  
hadoop vendors!

<http://0xdata.com/partners>

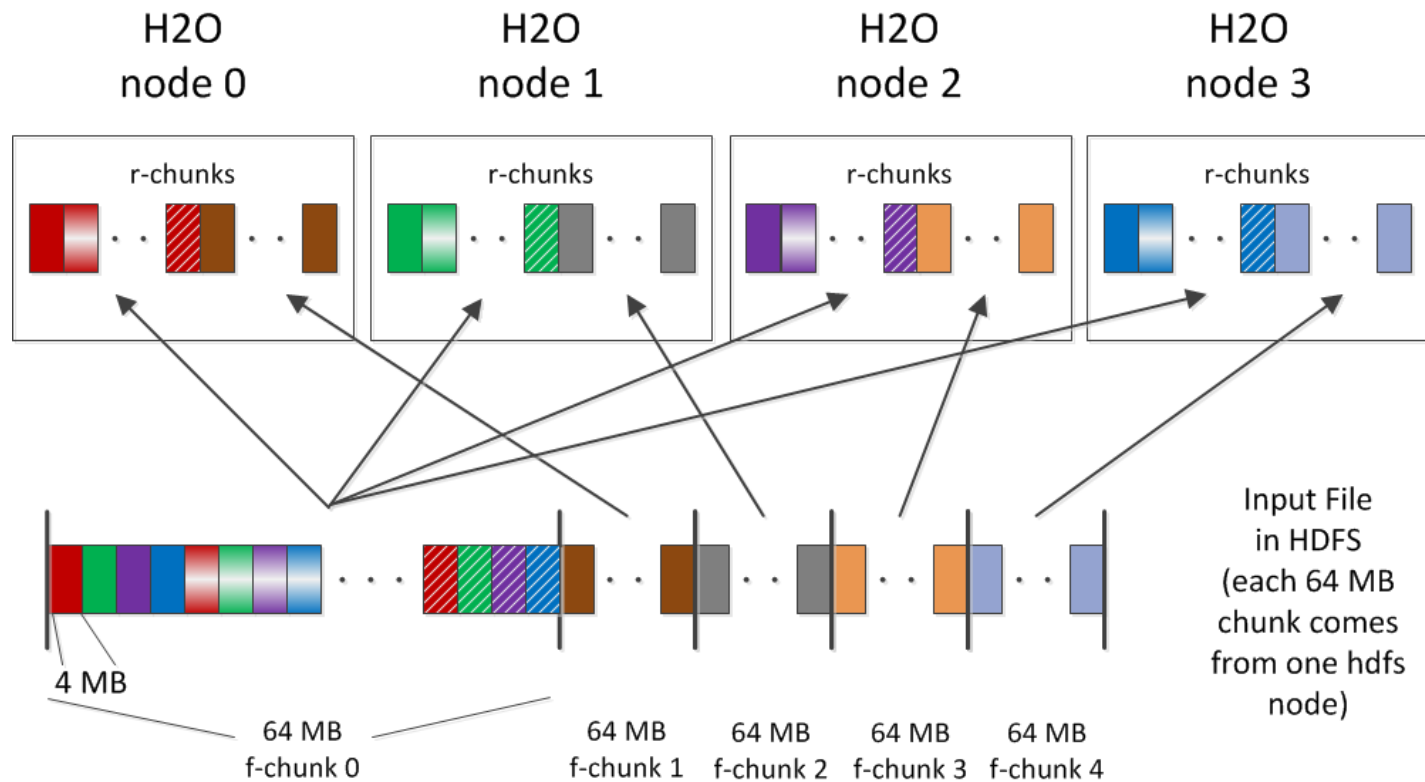
0xdata

Partners



0xdata

## Raw (Pre-Parse) Data Ingestion Pattern



Using H<sub>2</sub>O as an HDFS client

## Why

I want to do run a Generalized Linear Model, and my data lives in HDFS.

## What

Use H<sub>2</sub>O.

## How

(Command line) `java -jar h2o.jar`

(Web UI Menu) Data -> Import HDFS

**Oxdata**

Run H<sub>2</sub>O as a Hadoop job

## Why

I want to do Predictive Analytics on Big Data.

My data lives in HDFS.

***I have an existing Hadoop cluster. Let's utilize its CPUs and Memory!***

## What

Run H<sub>2</sub>O on Hadoop.

## How

```
hadoop jar h2odriver_<hadoop_version>.jar [...]
```

(Web UI Menu) Data -> Import HDFS

**Oxdata**