

Parallel and Distributed Joins in H2O

Lightning 12min Talk

Bay Area R User Group

09 Feb 2016

Matt Dowle

True radix sorting

Terdiman, 2000

<http://codercorner.com/RadixSortRevisited.htm>

Herf, 2001

<http://stereopsis.com/radix.html>

Dowle & Srinivasan, 2015

<http://user2015.math.aau.dk/presentations/234.pdf>

Now included in R itself thanks to time and skill of Michael Lawrence.

But

Single threaded

Single node

Limited to 2 billion rows (2^{31})

=> H2O

Every step has now been parallelized and distributed

data.table join

- Find order of the left join columns
- Find order of the right join columns
- Binary merge the two sorted indexes
- No hash table at all
- Fast ordered joins; e.g. rolling forwards, backwards, nearest and limited staleness

Further info

<https://github.com/Rdatatable/data.table/wiki>

Presentations

Articles

DataCamp course

Cardinality

Not just the number of rows, but what's in the rows

A 10 billion row file could contain :

500 stock tickers (low cardinality)

Millions of people (medium cardinality)

Billions of devices (high cardinality)

Create some test data

```
#!/usr/bin/awk -f
nrow = ARGV[2]
while(i<nrow)
{
    printf "%d,%d\n",
        rand()*nrow,
        rand()*nrow*2 - nrow;
    i++;
}
```

1e6 rows; 14MB

```
$ head X
```

```
KEY, X2
```

```
82967, -9233
```

```
550105, -819078
```

```
963516, -663146
```

```
706905, -128965
```

```
766103, 774695
```

```
$ head Y
```

```
KEY, Y2
```

```
610198, 322685
```

```
872395, -887505
```

```
340972, 535361
```

```
23067, 346231
```

```
295498, 918692
```


1e7 rows; 156MB

```
$ head X
```

```
KEY, X2
```

```
829673, -92335
```

```
5501052, -8190789
```

```
9635168, -6631465
```

```
7069052, -1289657
```

```
7661030, 7746956
```

```
$ head Y
```

```
KEY, Y2
```

```
6101982, 3226855
```

```
8723957, -8875053
```

```
3409724, 5353612
```

```
230673, 3462315
```

```
2954985, 9186925
```

1e8 rows; 1.8GB

```
$ head X
```

```
KEY,X2
```

```
8296733,-923350
```

```
55010523,-81907897
```

```
96351686,-66314650
```

```
70690522,-12896576
```

```
76610309,77469562
```

```
$ head Y
```

```
KEY,Y2
```

```
61019825,32268551
```

```
87239579,-88750532
```

```
34097244,53536122
```

```
2306734,34623153
```

```
29549857,91869251
```

1e9 rows; 19GB

```
$ head X
```

```
KEY, X2
```

```
82967333, -9233501
```

```
550105235, -819078974
```

```
963516860, -663146506
```

```
706905226, -128965762
```

```
766103099, 774695629
```

```
$ head Y
```

```
KEY, Y2
```

```
610198251, 322685514
```

```
872395790, -887505326
```

```
340972449, 535361227
```

```
23067343, 346231535
```

```
295498572, 918692512
```

1e10 rows; 200GB

```
$ head X
```

```
KEY,X2
```

```
829673335,-92335012
```

```
5501052357,-8190789743
```

```
9635168607,-6631465068
```

```
7069052265,-1289657629
```

```
7661030994,7746956291
```

```
$ head Y
```

```
KEY,Y2
```

```
6101982517,3226855142
```

```
8723957901,-8875053263
```

```
3409724497,5353612273
```

```
230673439,3462315357
```

```
2954985724,9186925123
```

H2O commands

```
library(h2o)
```

```
h2o.init(ip="mr-0xd6", port=55666)
```

```
X = h2o.importFile("hdfs://mr-  
0xd6/datasets/mattd/X1e9_2c.csv")
```

```
Y = h2o.importFile("hdfs://mr-  
0xd6/datasets/mattd/Y1e9_2c.csv")
```

```
ans = h2o.merge(X, Y, method="radix")
```

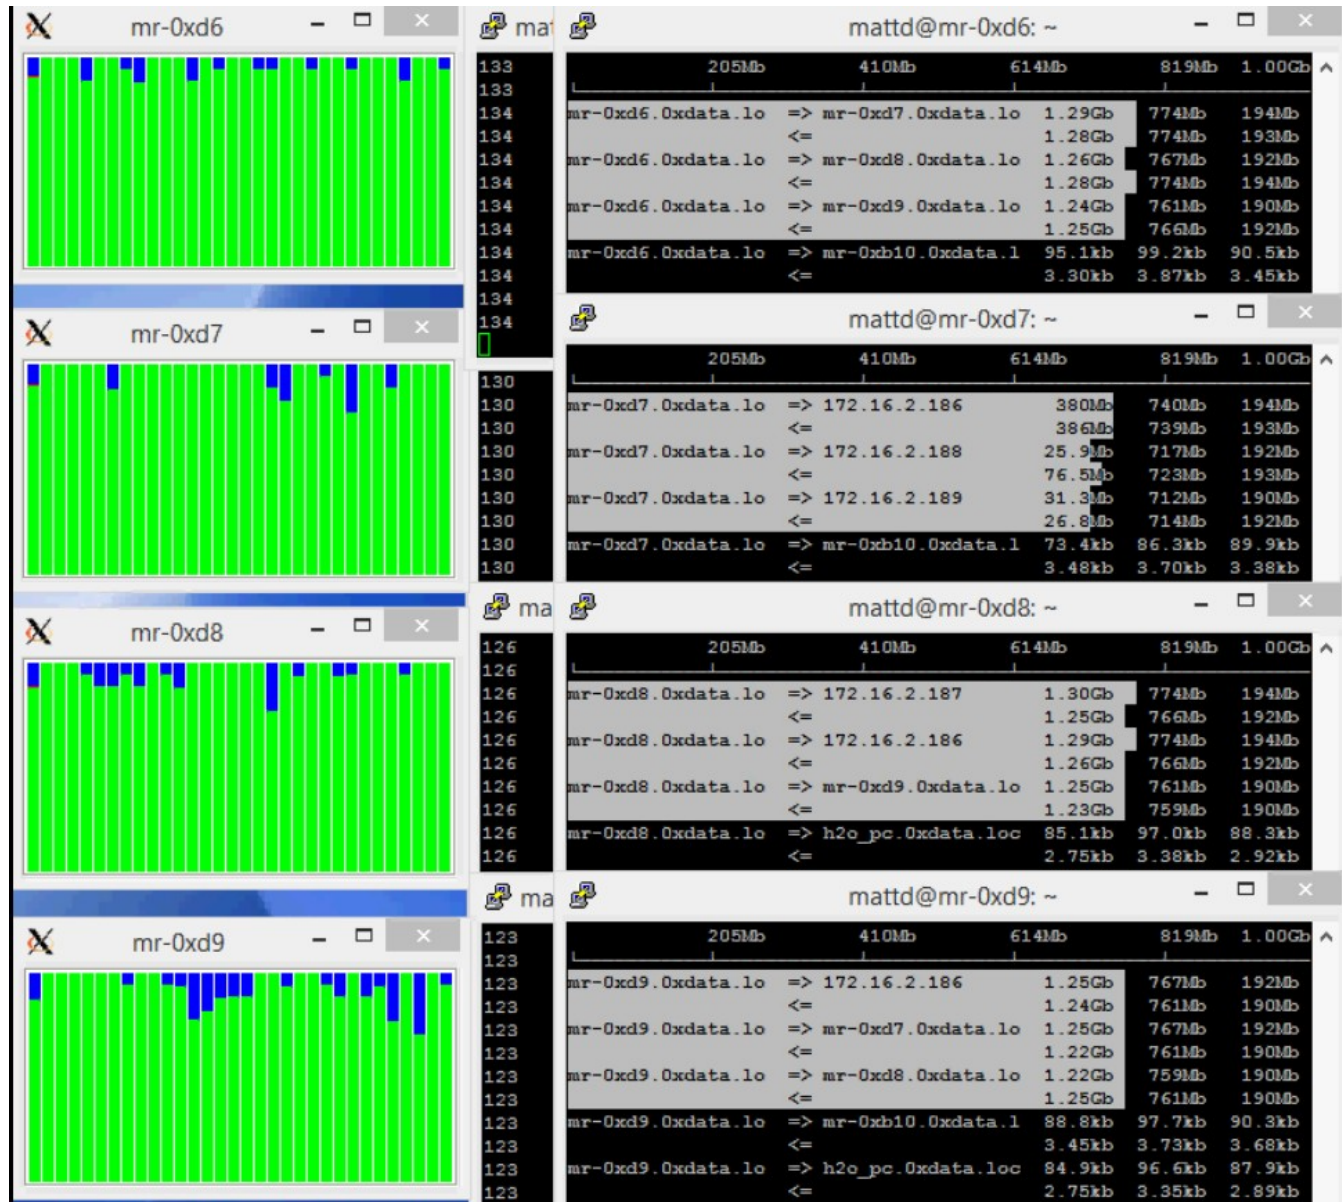
```
system.time(print(head(ans)))
```

Results on 4 nodes

4 physical 256GB / 32core

1e6	6s
1e7	7s
1e8	13s
1e9	49s
1e10	...

<https://www.youtube.com/watch?v=8VpzNibOme0>



Thank you.

Q & A