


H2O for Internet of Things



Jo-fai (Joe) Chow
Data Scientist
joe@h2o.ai
@matlabulous

Data Science Milan
Politecnico di Milano
10th October, 2016

Agenda


- First Talk (25 mins)
 - About H2O.ai
 - Demo
 - A Simple Classification Task
 - H2O's Web Interface
 - Why H2O?
 - Our Community
 - Our Customers
 - What's Next?
 - New H2O Features
- Second Talk (25 mins)
 - H2O for IoT
 - Predictive Maintenance
 - Anomaly Detection
 - H2O's R Interface
- Third Talk (25 mins)
 - Deep Water
 - Demo
 - H2O + mxnet on GPU
 - H2O's Python Interface

Data and Code

- Please go to bit.ly/h2o_milan_1
- subfolders
 - `iot_use_case_1`
 - `iot_use_case_2`

Use Case 1

Predictive Maintenance



Data for Use Case 1: SECOM

UCI



Machine Learning Repository

Center for Machine Learning and Intelligent Systems

SECOM Data Set

Download: [Data Folder](#), [Data Set Description](#)

Abstract: Data from a semi-conductor manufacturing process



| | | | | | |
|-----------------------------------|----------------------------------|------------------------------|------|----------------------------|------------|
| Data Set Characteristics: | Multivariate | Number of Instances: | 1567 | Area: | Computer |
| Attribute Characteristics: | Real | Number of Attributes: | 591 | Date Donated | 2008-11-19 |
| Associated Tasks: | Classification, Causal-Discovery | Missing Values? | Yes | Number of Web Hits: | 37895 |

Source:

<https://archive.ics.uci.edu/ml/datasets/SECOM>

Authors: Michael McCann, Adrian Johnston

Data Set Information:

A complex modern semi-conductor manufacturing process is normally under consistent surveillance via the monitoring of signals/variables collected from sensors and or process measurement points. However, not all of these signals are equally valuable in a specific monitoring system. The measured signals contain a combination of useful information, irrelevant information as well as noise. It is often the case that useful information is buried in the latter two. Engineers typically have a much larger number of signals than are actually required. If we consider each type of signal as a feature, then feature selection may be applied to identify the most relevant signals. The Process Engineers may then use these signals to determine key factors contributing to yield excursions downstream in the process. This will enable an increase in process throughput, decreased time to learning and reduce the per unit production costs.

To enhance current business improvement techniques the application of feature selection as an intelligent systems technique is being investigated.

The dataset presented in this case represents a selection of such features where each example represents a single production entity with associated measured features and the labels represent a simple pass/fail yield for in house line testing, figure 2, and associated date time stamp. Where -1 corresponds to a pass and 1 corresponds to a fail and the data time stamp is for that specific test point.

Using feature selection techniques it is desired to rank features according to their impact on the overall yield for the product, causal relationships may also be considered with a view to identifying the key features.

Results may be submitted in terms of feature relevance for predictability using error rates as our evaluation metrics. It is suggested that cross validation be applied to generate these results. Some baseline results are shown below for basic feature selection techniques using a simple kernel ridge classifier and 10 fold cross validation.

Baseline Results: Pre-processing objects were applied to the dataset simply to standardize the data and remove the constant features and then a number of different feature selection objects selecting 40 highest ranked features were applied with a simple classifier to achieve some initial results. 10 fold cross validation was used and the balanced error rate (*BER) generated as our initial performance metric to help investigate this dataset.

SECOM Dataset: 1567 examples 591 features, 104 fails

FSmethod (40 features) BER % True + % True - %
S2N (signal to noise) 34.5 +-2.6 57.8 +-5.3 73.1 +-2.1
Ttest 33.7 +-2.1 59.6 +-4.7 73.0 +-1.8
Relief 40.1 +-2.8 48.3 +-5.9 71.6 +-3.2
Pearson 34.1 +-2.0 57.4 +-4.3 74.4 +-4.9
Ftest 33.5 +-2.2 59.1 +-4.8 73.8 +-1.8
Gram Schmidt 35.6 +-2.4 51.2 +-11.8 77.5 +-2.3

We want to predict fails in the future.

The ML Problem – Pass/Fail

- Inputs
 - 591 features
- Output
 - Classification
 - -1 = pass
 - 1 = fail
- Size: 1567 Samples



| secom.csv - Excel | | | | | | | | | | | | |
|---|----|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| File Home Insert Draw Page Layout Formulas Data Review View Tell me what you want to do | | | | | | | | | | | | |
| A1 X ✓ fx ID | | | | | | | | | | | | |
| | A | B | C | D | E | F | G | H | I | J | K | L |
| 1 | ID | Feature 001 | Feature 002 | Feature 003 | Feature 004 | Feature 005 | Feature 006 | Feature 007 | Feature 008 | Feature 009 | Feature 010 | Feature 011 |
| 2 | 1 | 3030.93 | 2564 | 2187.7333 | 1411.1265 | 1.3602 | 100 | 97.6133 | 0.1242 | 1.5005 | 0.0162 | -0.0001 |
| 3 | 2 | 3095.78 | 2465.14 | 2230.4222 | 1463.6606 | 0.8294 | 100 | 102.3433 | 0.1247 | 1.4966 | -5.00E-04 | -0.0001 |
| 4 | 3 | 2932.61 | 2559.94 | 2186.4111 | 1698.0172 | 1.5102 | 100 | 95.4878 | 0.1241 | 1.4436 | 0.0041 | 0.0001 |
| 5 | 4 | 2988.72 | 2479.9 | 2199.0333 | 909.7926 | 1.3602 | 100 | 102.3433 | 0.1247 | 1.4882 | -0.0124 | -0.0001 |
| 6 | 5 | 3032.24 | 2502.87 | 2233.3667 | 1326.52 | 1.5334 | 100 | 100.3967 | 0.1235 | 1.5031 | -0.0031 | -0.0001 |
| 7 | 6 | 2946.25 | 2432.84 | 2233.3667 | 1326.52 | 1.5334 | 100 | 100.3967 | 0.1235 | 1.5287 | 0.0167 | 0.0001 |
| 8 | 7 | 3030.27 | 2430.12 | 2230.4222 | 1463.6606 | 0.8294 | 100 | 102.3433 | 0.1247 | 1.5816 | -0.027 | 0.0001 |
| 9 | 8 | 3058.88 | 2690.15 | 2248.9 | 1004.4692 | 0.7884 | 100 | 106.24 | 0.1185 | 1.5153 | 0.0157 | 7.00E-05 |
| 10 | 9 | 2967.68 | 2600.47 | 2248.9 | 1004.4692 | 0.7884 | 100 | 106.24 | 0.1185 | 1.5358 | 0.0111 | -0.0001 |
| 11 | 10 | 3016.11 | 2428.37 | 2248.9 | 1004.4692 | 0.7884 | 100 | 106.24 | 0.1185 | 1.5381 | 0.0159 | 0.0001 |
| 12 | 11 | 2994.05 | 2548.21 | 2195.1222 | 1046.1468 | 1.3204 | 100 | 103.34 | 0.1223 | 1.5144 | -0.019 | 0.0001 |
| 13 | 12 | 2928.84 | 2479.4 | 2196.2111 | 1605.7578 | 0.9959 | 100 | 97.9156 | 0.1257 | 1.469 | 0.017 | -0.0001 |
| 14 | 13 | 2920.07 | 2507.4 | 2195.1222 | 1046.1468 | 1.3204 | 100 | 103.34 | 0.1223 | 1.531 | -0.0259 | 0.0001 |
| 15 | 14 | 3051.44 | 2529.27 | 2184.4333 | 877.6266 | 1.4668 | 100 | 107.8711 | 0.124 | 1.5236 | -0.0209 | -0.0001 |
| 16 | 15 | 2963.97 | 2629.48 | 2224.6222 | 947.7739 | 1.2824 | 100 | 104.8489 | 0.1197 | 1.4474 | 0.0144 | -0.0001 |
| 17 | 16 | 2963.97 | 2629.48 | 2224.6222 | 947.7739 | 1.2824 | 100 | 104.8489 | 0.1197 | 1.5465 | 0.025 | -0.0001 |
| 18 | 17 | 2963.97 | 2629.48 | 2224.6222 | 947.7739 | 1.2824 | 100 | 104.8078 | 0.1207 | 1.4368 | 0.015 | -0.0001 |
| 19 | 18 | 3032.73 | 2517.79 | 2270.2556 | 1258.4558 | 1.395 | 100 | 104.8078 | 0.1207 | 1.5537 | 0.022 | -0.0001 |
| 20 | 19 | 3040.34 | 2501.16 | 2207.3889 | 962.5317 | 1.2043 | 100 | 104.0311 | 0.121 | 1.5481 | -0.0367 | 0.0001 |

Features (Numeric)


ID (excluded from modeling)

Features (Numeric)

Response (Classification)
-1 (Pass) or 1 (Fail)

Use Case 1: Predictive Maintenance

Step 1: R Packages



step_01_install_packages.R

Package 'h2o'

```
1 # -----
2 # Step 1: Install R Packages for Workshop
3 # -----
4
5 # -----
6 # Install "h2o" for machine learning
7 # Reference: http://www.h2o.ai/download/h2o/r
8 # -----
9
10 # The following two commands remove any previously installed H2O packages for R.
11 if ("package:h2o" %in% search()) { detach("package:h2o", unload=TRUE) }
12 if ("h2o" %in% rownames(installed.packages())) { remove.packages("h2o") }
13
14 # Next, we download packages that H2O depends on.
15 pkgs <- c("methods", "statmod", "stats", "graphics", "RCurl", "jsonlite", "tools", "utils")
16 for (pkg in pkgs) {
17   if (! (pkg %in% rownames(installed.packages()))) { install.packages(pkg) }
18 }
19
20 # Now we download, install and initialize the H2O package for R.
21 install.packages("h2o",
22                 type="source",
23                 repos=(c("http://h2o-release.s3.amazonaws.com/h2o/rel-turing/7/R")))
24
25 # Quick test
26 suppressPackageStartupMessages(library(h2o))
27 h2o.init(nthreads = -1)
```


```
> suppressPackageStartupMessages(library(h2o))
> h2o.init(nthreads = -1)
Connection successful!
```

R is connected to the H2O cluster:

| | |
|----------------------------|-------------------------------------|
| H2O cluster uptime: | 1 minutes 891 milliseconds |
| H2O cluster version: | 3.10.0.7 |
| H2O cluster version age: | 6 days |
| H2O cluster name: | H2O_started_from_R_jofaichow_ayn543 |
| H2O cluster total nodes: | 1 |
| H2O cluster total memory: | 3.28 GB |
| H2O cluster total cores: | 8 |
| H2O cluster allowed cores: | 8 |
| H2O cluster healthy: | TRUE |
| H2O Connection ip: | localhost |
| H2O Connection port: | 54321 |
| H2O Connection proxy: | NA |
| R Version: | R version 3.3.0 (2016-05-03) |

Use Case 1: Predictive Maintenance

Step 2: Exploratory Analysis



step_02_exploratory_analysis.R

Importing SECOM data

```
1- # -----
2 # Step 2: Data Exploration
3- # -----
4
5 # Start and connect to a local H2O cluster
6 library(h2o)
7 h2o.init(nthreads = -1)
8
9 # Import data from a local CSV file
10 # Source: https://archive.ics.uci.edu/ml/machine-learning-databases/secom/
11 secom <- h2o.importFile(path = "./data/secom.csv", destination_frame = "secom")
12
13 # (Optional) Demo - Importing files using URLs
14 secom <- h2o.importFile(
15   path = "https://github.com/woobe/H2O_London_Workshop/raw/master/data/secom.csv",
16   destination_frame = "secom")
17
18 # (Optional) Demo - Converting R data frame into H2O data frame
19 hdf_iris <- as.h2o(iris)
20
21 # (Optional) Turning off progress bar in R
22 h2o.no_progress()
```

Optional (different ways to import data)

```
java version "1.8.0_72"
Java(TM) SE Runtime Environment (build 1.8.0_72-b15)
Java HotSpot(TM) 64-Bit Server VM (build 25.72-b15, mixed mode)
```

Starting H2O JVM and connecting: ... Connection successful!

R is connected to the H2O cluster:

| | |
|----------------------------|-------------------------------------|
| H2O cluster uptime: | 2 seconds 721 milliseconds |
| H2O cluster version: | 3.10.0.7 |
| H2O cluster version age: | 7 days, 10 hours and 4 minutes |
| H2O cluster name: | H2O_started_from_R_jofaichow_cow128 |
| H2O cluster total nodes: | 1 |
| H2O cluster total memory: | 3.56 GB |
| H2O cluster total cores: | 8 |
| H2O cluster allowed cores: | 8 |
| H2O cluster healthy: | TRUE |
| H2O Connection ip: | localhost |
| H2O Connection port: | 54321 |
| H2O Connection proxy: | NA |
| R Version: | R version 3.3.0 (2016-05-03) |

```
>
> # Import data from a local CSV file
> # Source: https://archive.ics.uci.edu/ml/machine-learning-databases/secom/
> secom <- h2o.importFile(path = "./data/secom.csv", destination_frame = "secom")
|=====| 100%
```

step_02_exploratory_analysis.R

Basic exploratory analysis

```
25 # Basic exploratory analysis
26 print(dim(secom)) # 1567 x 599
27 print(summary(secom$Classification))
28 # alternatively, use H2O flow to look at data (localhost:54321)
29
30 # Convert Classification to factor
31 secom$Classification <- as.factor(secom$Classification)
32 print(summary(secom$Classification))
```

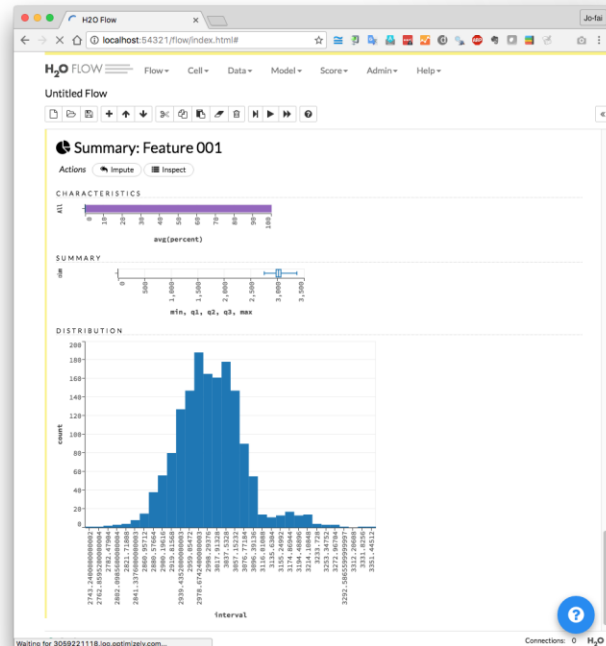
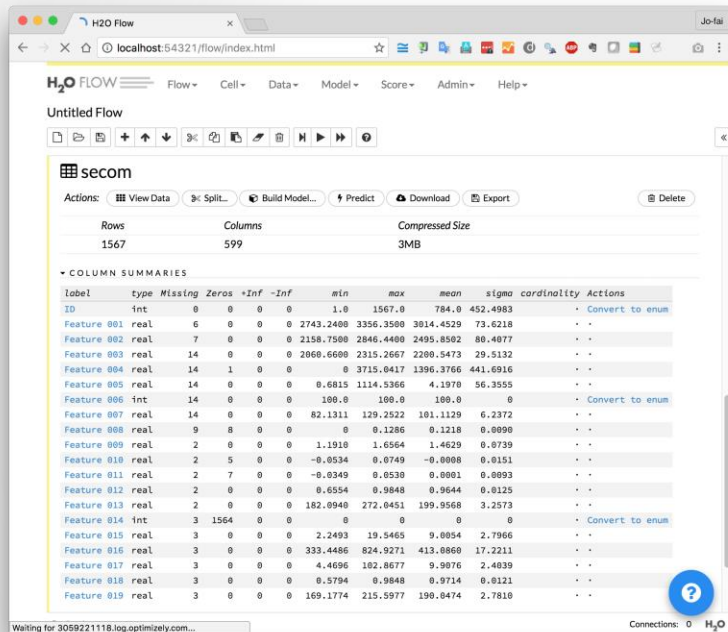
Convert -1 and 1 to categorical value

```
> # Basic exploratory analysis
> print(dim(secom)) # 1567 x 599
[1] 1567 599
> print(summary(secom$Classification))
Classification
Min.    :-1.0000
1st Qu. :-1.0000
Median  :-1.0000
Mean    :-0.8673
3rd Qu. :-1.0000
Max.    : 1.0000
```

```
> # Convert Classification to factor
> secom$Classification <- as.factor(secom$Classification)
> print(summary(secom$Classification))
Classification
-1:1463
1 : 104
```

Note: Imbalance dataset (only 104 fails)

Use H2O Flow (localhost:54321)



Use Case 1: Predictive Maintenance

Step 3: Building & Evaluating Models



step_03_basic_models.R

Define features & target

```
5 ~ # -----
6 # Loading data (same as previous steps)
7 ~ # -----
8
9 # Start and connect to a local H2O cluster
10 library(h2o)
11 h2o.init(nthreads = -1)
12
13 # Import data from a local CSV file
14 # Source: https://archive.ics.uci.edu/ml/machine-learning-databases/secom/
15 secom <- h2o.importFile(path = "./data/secom.csv", destination_frame = "secom")
16
17 # Convert Classification to factor
18 secom$Classification <- as.factor(secom$Classification)
19
20
21 ~ # -----
22 # Define Targets and Features
23 ~ # -----
24
25 target <- "Classification"
26 features <- setdiff(colnames(secom), c("ID", "Classification"))
27
28 print(target)
29 print(features)
```

```
> print(target)
[1] "Classification"
> print(features)
[1] "Feature 001" "Feature 002" "Feature 003" "Feature 004"
[5] "Feature 005" "Feature 006" "Feature 007" "Feature 008"
[9] "Feature 009" "Feature 010" "Feature 011" "Feature 012"
[13] "Feature 013" "Feature 014" "Feature 015" "Feature 016"
[17] "Feature 017" "Feature 018" "Feature 019" "Feature 020"
[21] "Feature 021" "Feature 022" "Feature 023" "Feature 024"
[25] "Feature 025" "Feature 026" "Feature 027" "Feature 028"
[29] "Feature 029" "Feature 030" "Feature 031" "Feature 032"
[33] "Feature 033" "Feature 034" "Feature 035" "Feature 036"
[37] "Feature 037" "Feature 038" "Feature 039" "Feature 040"
[41] "Feature 041" "Feature 042" "Feature 043" "Feature 044"
[45] "Feature 045" "Feature 046" "Feature 047" "Feature 048"
[49] "Feature 049" "Feature 050" "Feature 051" "Feature 052"
[53] "Feature 053" "Feature 054" "Feature 055" "Feature 056"
[57] "Feature 057" "Feature 058" "Feature 059" "Feature 060"
[61] "Feature 061" "Feature 062" "Feature 063" "Feature 064"
[65] "Feature 065" "Feature 066" "Feature 067" "Feature 068"
[69] "Feature 069" "Feature 070" "Feature 071" "Feature 072"
[73] "Feature 073" "Feature 074" "Feature 075" "Feature 076"
[77] "Feature 077" "Feature 078" "Feature 079" "Feature 080"
[81] "Feature 081" "Feature 082" "Feature 083" "Feature 084"
[85] "Feature 085" "Feature 086" "Feature 087" "Feature 088"
[89] "Feature 089" "Feature 090" "Feature 091" "Feature 092"
[93] "Feature 093" "Feature 094" "Feature 095" "Feature 096"
[97] "Feature 097" "Feature 098" "Feature 099" "Feature 100"
[101] "Feature 101" "Feature 102" "Feature 103" "Feature 104"
[105] "Feature 105" "Feature 106" "Feature 107" "Feature 108"
```

step_03_basic_models.R

Split data with a random seed

```
32 # -----  
33 # Split data into training / test  
34 # -----  
35  
36 # Split  
37 # i.e. using 60% of data for training and 40% for test  
38 secom_splits <- h2o.splitFrame(data = secom, ratios = 0.6, seed = 1234)  
39 secom_train <- secom_splits[[1]]  
40 secom_test  <- secom_splits[[2]]  
41  
42 # Check  
43 summary(secom_train$Classification) # 882 : 62 ... % of 1 = 0.07029478  
44 summary(secom_test$Classification) # 581 : 42 ... % of 1 = 0.07228916
```

```
> summary(secom_train$Classification)  
Classification  
-1:882  
1 : 62
```

```
> summary(secom_test$Classification)  
Classification  
-1:581  
1 : 42
```

Classification 1 samples \approx 7%

step_03_basic_models.R

Train H2O models with default values

```
47 # -----
48 # Train H2O models with default value
49 # -----
50
51 # Turn off progress bar (if you want to ...)
52 # h2o.no_progress()
53
54 # GBM
55 model_gbm <- h2o.gbm(x = features, y = target,
56                     training_frame = secom_train)
57
58 # Random Forest
59 model_drf <- h2o.randomForest(x = features, y = target,
60                             training_frame = secom_train)
61
62 # Deep Neural Network
63 model_dnn <- h2o.deeplearning(x = features, y = target,
64                              training_frame = secom_train)
```

```
> # GBM
> model_gbm <- h2o.gbm(x = features, y = target,
+                     training_frame = secom_train)
|=====| 100%

Warning message:
In .h2o.startModelJob(algo, params, h2oRestApiVersion) :
  Dropping constant columns: [Feature 516, Feature 234, Feature 233, Feature 236, Feature 235, Feature
510, Feature 238, Feature 513, Feature 237, Feature 479, Feature 515, Feature 514, Feature 193, Featur
e 192, Feature 195, Feature 194, Feature 075, Feature 230, Feature 232, Feature 231, Feature 529, Feat
ure 244, Feature 365, Feature 401, Feature 400, Feature 006, Feature 403, Feature 402, Feature 405, Fe
ature 404, Feature 241, Feature 482, Feature 243, Feature 242, Feature 180, Feature 179, Feature 459,
Feature 050, Feature 053, Feature 450, Feature 210, Feature 331, Feature 452, Feature 330, Feature 451
, Feature 191, Feature 070, Feature 190, Feature 506, Feature 505, Feature 508, Feature 507, Feature 5
09, Feature 465, Feature 343, Feature 464, Feature 467, Feature 466, Feature 227, Feature 348, Feature
502, Feature 504, Feature 503, Feature 463, Feature 187, Feature 462, Feature 399, Feature 277, Featur
e 398, Feature 315, Feature 314, Feature 316, Feature 150, Feature 395, Feature 39 [... truncated]
```

H2O automatically ignores
Columns with constant values

step_03_basic_models.R

summary(model_xxx)

```
> print(summary(model_gbm))
Model Details:
=====
H2OBinomialModel: gbm
Model Key:   GBM_model_R_1474971910804_1
Model Summary:
  number_of_trees number_of_internal_trees model_size_in_bytes min_depth max_depth mean_depth
1              50              50              10880              5              5      5.000000
  min_leaves max_leaves mean_leaves
1           6          21      12.38000

H2OBinomialMetrics: gbm
** Reported on training data. **

MSE: 0.002598174
RMSE: 0.05097229
Logloss: 0.0248643
Mean Per-Class Error: 0
AUC: 1
Gini: 1

Confusion Matrix for F1-optimal threshold:
      -1  1  Error  Rate
-1      882  0 0.000000 =0/882
1         62  0 0.000000 =0/62
Totals 882 62 0.000000 =0/944

Maximum Metrics: Maximum metrics at their respective thresholds
      metric threshold  value idx
1      max f1 0.628627 1.000000 61
2      max f2 0.628627 1.000000 61
3      max f0point5 0.628627 1.000000 61
4      max accuracy 0.628627 1.000000 61
5      max precision 0.956650 1.000000 0
6      max recall 0.628627 1.000000 61
7      max specificity 0.956650 1.000000 0
8      max absolute_mcc 0.628627 1.000000 61
9      max min_per_class_accuracy 0.628627 1.000000 61
10     max mean_per_class_accuracy 0.628627 1.000000 61
```

```
Scoring History:
      timestamp  duration number_of_trees training_rmse training_logloss training_auc
1 2016-09-27 11:25:31 0.021 sec              0      0.24772      0.24231      0.50000
2 2016-09-27 11:25:32 0.722 sec              1      0.24010      0.21926      0.83663
3 2016-09-27 11:25:32 1.001 sec              2      0.22980      0.19523      0.92679
4 2016-09-27 11:25:32 1.178 sec              3      0.22041      0.17966      0.94964
5 2016-09-27 11:25:32 1.320 sec              4      0.21468      0.16997      0.96445
```

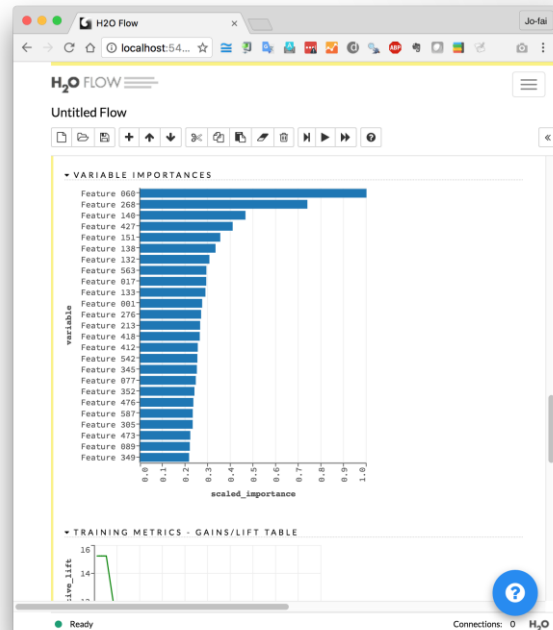
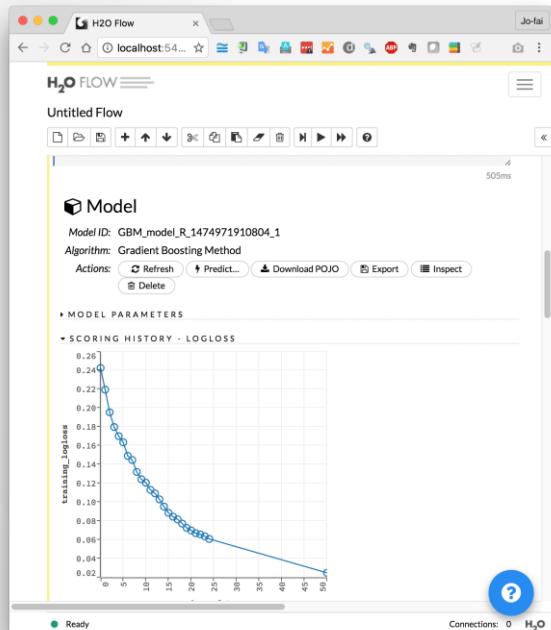
```
      training_lift training_classification_error
1      1.00000      0.93432
2      9.13548      0.08581
3     12.18065      0.04449
4     13.70323      0.03496
5     15.22581      0.03708
```

```
---
      timestamp  duration number_of_trees training_rmse training_logloss training_auc
21 2016-09-27 11:25:34 3.500 sec              20      0.12174      0.06986      0.99996
22 2016-09-27 11:25:34 3.631 sec              21      0.11797      0.06696      1.00000
23 2016-09-27 11:25:35 3.726 sec              22      0.11645      0.06571      1.00000
24 2016-09-27 11:25:35 3.841 sec              23      0.11373      0.06369      1.00000
25 2016-09-27 11:25:35 3.967 sec              24      0.10979      0.06082      1.00000
26 2016-09-27 11:25:37 6.596 sec              50      0.05097      0.02486      1.00000
      training_lift training_classification_error
21     15.22581      0.00106
22     15.22581      0.00000
23     15.22581      0.00000
24     15.22581      0.00000
25     15.22581      0.00000
26     15.22581      0.00000
```

Variable Importances: (Extract with `h2o.varimp`)

```
Variable Importances:
      variable relative_importance scaled_importance percentage
1 Feature 060      9.594806      1.000000      0.050880
2 Feature 268      7.089089      0.738847      0.037592
3 Feature 140      4.466067      0.465467      0.023683
4 Feature 427      3.924058      0.408977      0.020809
5 Feature 151      3.397748      0.354124      0.018018
```

Use H2O Flow (localhost:54321)



step_03_basic_models.R

Evaluate models with test data

```
72 # -----  
73 # Evaluate model performance on unseen data  
74 # -----  
75  
76 h2o.performance(model_gbm, newdata = secom_test)  
77 h2o.performance(model_drf, newdata = secom_test)  
78 h2o.performance(model_dnn, newdata = secom_test)
```

```
> h2o.performance(model_gbm, newdata = secom_test)  
H2OBinomialMetrics: gbm
```

```
MSE: 0.06408139  
RMSE: 0.253143  
LogLoss: 0.2678549  
Mean Per-Class Error: 0.2928858  
AUC: 0.7041841  
Gini: 0.4083682
```

Confusion Matrix for F1-optimal threshold:

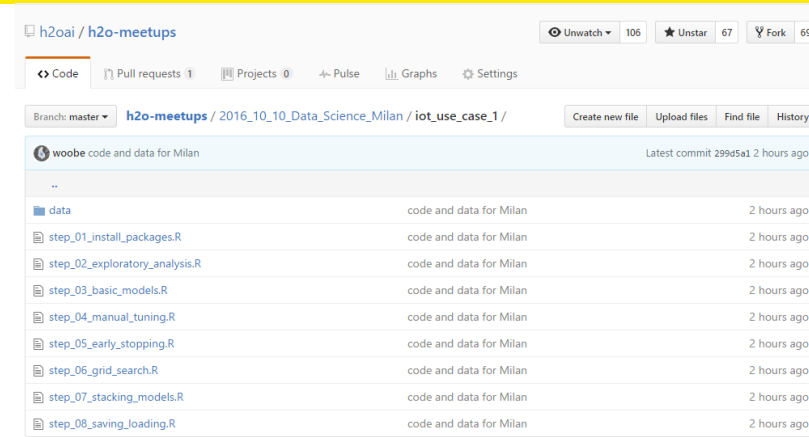
| | -1 | 1 | Error | Rate |
|--------|-----|-----|----------|----------|
| -1 | 379 | 202 | 0.347676 | =202/581 |
| 1 | 10 | 32 | 0.238095 | =10/42 |
| Totals | 389 | 234 | 0.340289 | =212/623 |

Maximum Metrics: Maximum metrics at their respective thresholds

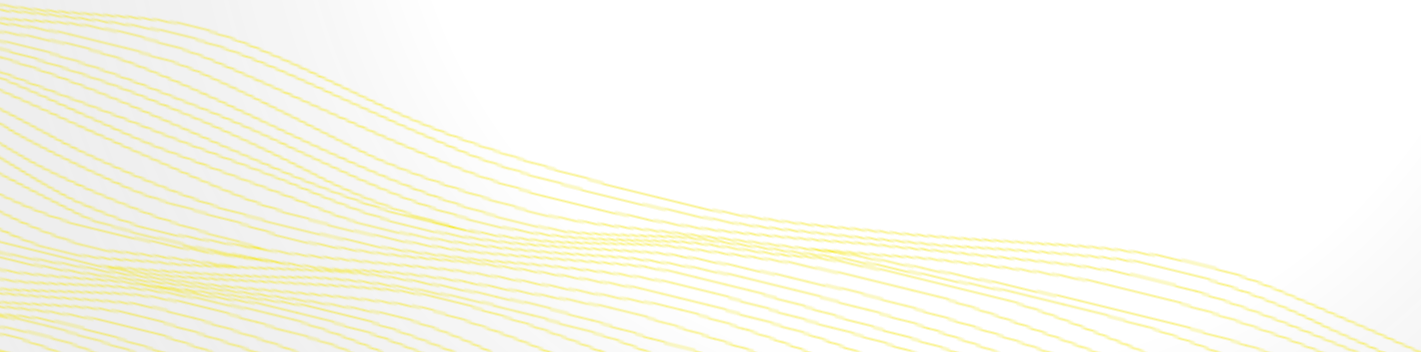
| | metric | threshold | value | idx |
|----|-----------------------------|-----------|----------|-----|
| 1 | max f1 | 0.016032 | 0.231884 | 201 |
| 2 | max f2 | 0.016032 | 0.398010 | 201 |
| 3 | max f0point5 | 0.053221 | 0.196078 | 63 |
| 4 | max accuracy | 0.492634 | 0.930979 | 0 |
| 5 | max precision | 0.244674 | 0.333333 | 5 |
| 6 | max recall | 0.005088 | 1.000000 | 378 |
| 7 | max specificity | 0.492634 | 0.998279 | 0 |
| 8 | max absolute_mcc | 0.016032 | 0.214472 | 201 |
| 9 | max min_per_class_accuracy | 0.016481 | 0.659208 | 196 |
| 10 | max mean_per_class_accuracy | 0.016032 | 0.707114 | 201 |

Advanced Procedures

- Step 4 – Manual Tuning
- Step 5 – Early Stopping
- Step 6 – Grid Search
- Step 7 – Stacking Models (“h2oEnsemble”)
- Step 8 – Saving/Loading Models
- Please try them out later (bit.ly/h2o_milan_1)



Use Case 2: Anomaly Detection



Anomaly (Outlier) Detection

- Definition

- Identification of items, events or observations which do not conform to an expected pattern or other items in a dataset.





Photo credit: www.dbta.com

- Use Cases


- Bank Fraud
- Monitoring Manufacturing Lines
- Machine Learning
 - Separate dataset and build different models

Deep Autoencoder for Anomaly Detection

- Consider the following three-layer neural network with one hidden layer and the same number of input neurons (features) as output neurons.
- The loss function is the mean squared error (MSE) between the input and the output. Hence, the network is forced to learn the identity via a nonlinear, reduced representation of the original data.
 - e.g. High MSE = potential outliers
- Such an algorithm is called a **deep autoencoder**.

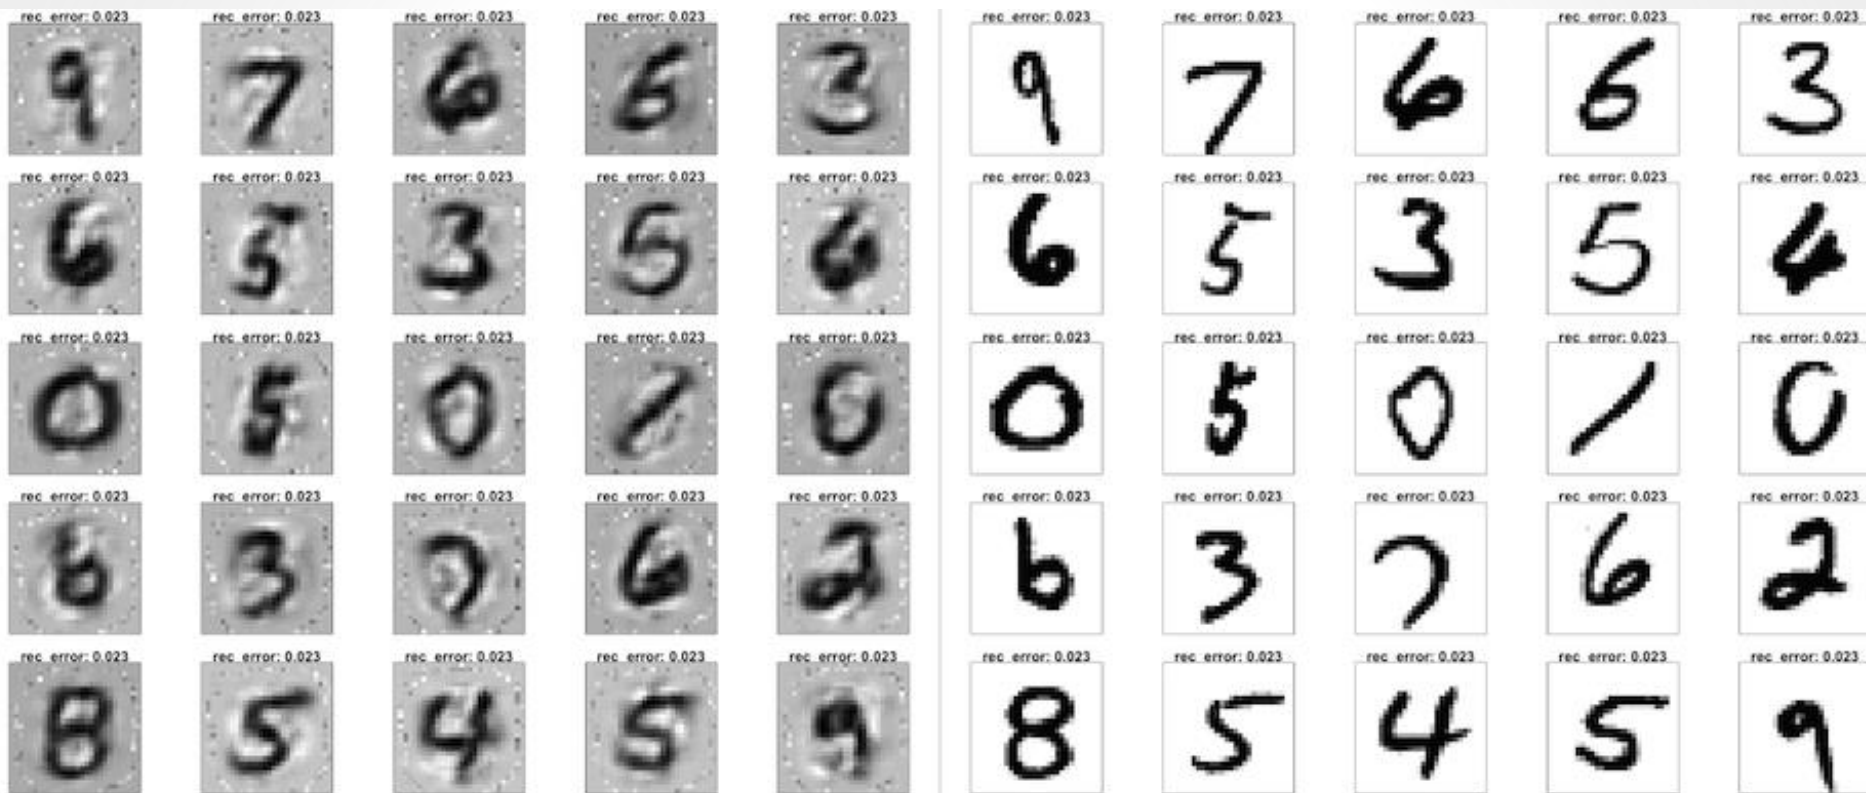


MNIST Example – The Good Ones



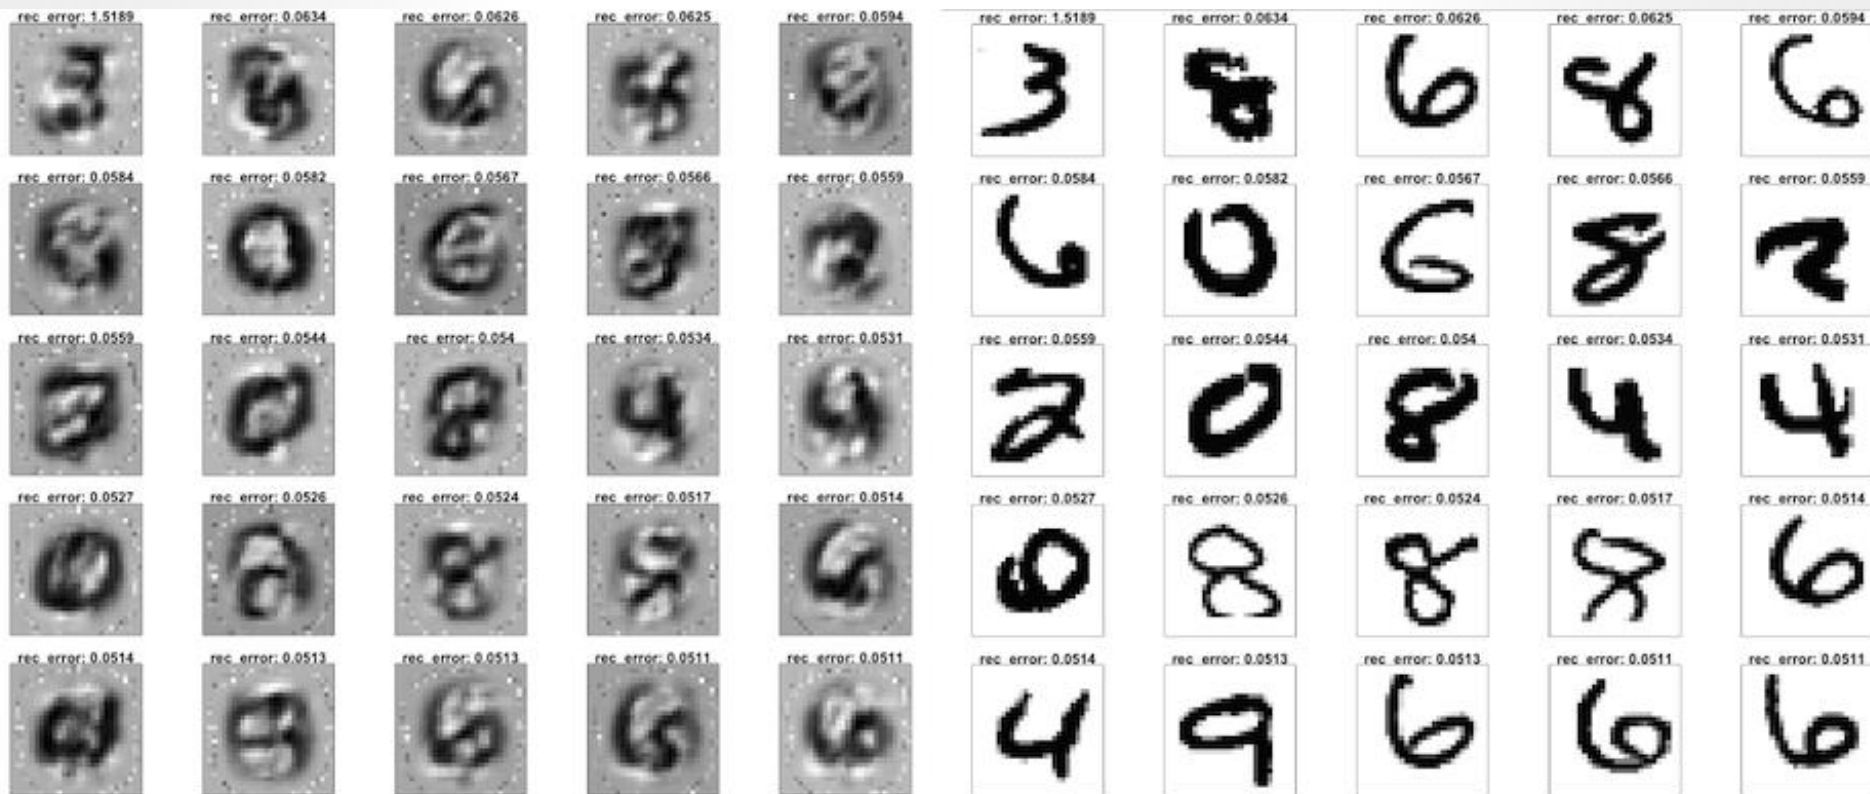
Samples with Low Mean Squared Error (MSE)

MNIST Example – The Bad Ones



Samples with High Mean Squared Error (MSE)

MNIST Example – The Ugly Ones



Samples with Highest Mean Squared Error (MSE)

use_case_2_anomaly_detection.R

```
1 # -----
2 # Step 8: Using Deep Learning for Anomaly Detection
3 # -----
4
5 # Start and connect to a local H2O cluster
6 library(h2o)
7 h2o.init(nthreads = -1)
8
9 # Import data from a local CSV file
10 mtcars <- read.csv("../data/auto_design.csv")
11 mtcars$gear <- as.factor(mtcars$gear)
12 mtcars$carb <- as.factor(mtcars$carb)
13 mtcars$cyl <- as.factor(mtcars$cyl)
14 mtcars$vs <- as.factor(mtcars$vs)
15 mtcars$am <- as.factor(mtcars$am)
16 mtcars$ID <- 1:nrow(mtcars)
17
18 # Print it out
19 print(mtcars)
20
21 # Convert R data frame into H2O data frame
22 h2o_mtcars <- as.h2o(mtcars)
```

```
> print(mtcars)
```

| | | X | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb | ID |
|----|---------------------|------|-----|-------|------|--------|---------|---------|------|----|----|------|------|----|
| 1 | Mazda RX4 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 | 1 | |
| 2 | Mazda RX4 Wag | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 | 2 | |
| 3 | Datsun 710 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 | 3 | |
| 4 | Hornet 4 Drive | 21.4 | 6 | 258.0 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 | 4 | |
| 5 | Hornet Sportabout | 18.7 | 8 | 360.0 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 | 5 | |
| 6 | Valiant | 18.1 | 6 | 225.0 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 | 6 | |
| 7 | Duster 360 | 14.3 | 8 | 360.0 | 245 | 3.21 | 3.570 | 15.84 | 0 | 0 | 3 | 4 | 7 | |
| 8 | Merc 240D | 24.4 | 4 | 146.7 | 62 | 3.69 | 3.190 | 20.00 | 1 | 0 | 4 | 2 | 8 | |
| 9 | Merc 230 | 22.8 | 4 | 140.8 | 95 | 3.92 | 3.150 | 22.90 | 1 | 0 | 4 | 2 | 9 | |
| 10 | Merc 280 | 19.2 | 6 | 167.6 | 123 | 3.92 | 3.440 | 18.30 | 1 | 0 | 4 | 4 | 10 | |
| 11 | Merc 280C | 17.8 | 6 | 167.6 | 210 | 800.00 | 900.000 | 1000.00 | 1 | 0 | 4 | 4 | 11 | |
| 12 | Merc 450SE | 16.4 | 8 | 275.8 | 180 | 3.07 | 4.070 | 17.40 | 0 | 0 | 3 | 3 | 12 | |
| 13 | Merc 450SL | 17.3 | 8 | 275.8 | 180 | 3.07 | 3.730 | 17.60 | 0 | 0 | 3 | 3 | 13 | |
| 14 | Merc 450SLC | 15.2 | 8 | 275.8 | 180 | 3.07 | 3.780 | 18.00 | 0 | 0 | 3 | 3 | 14 | |
| 15 | Cadillac Fleetwood | 10.4 | 8 | 472.0 | 205 | 2.93 | 5.250 | 17.98 | 0 | 0 | 3 | 4 | 15 | |
| 16 | Lincoln Continental | 10.4 | 8 | 460.0 | 215 | 3.00 | 5.424 | 17.82 | 0 | 0 | 3 | 4 | 16 | |
| 17 | Chrysler Imperial | 14.7 | 8 | 440.0 | 230 | 3.23 | 5.345 | 17.42 | 0 | 0 | 3 | 4 | 17 | |
| 18 | Fiat 128 | 32.4 | 4 | 780.0 | 2100 | 400.00 | 200.000 | 700.00 | 1 | 1 | 4 | 1 | 18 | |
| 19 | Honda Civic | 80.4 | 10 | 75.7 | 100 | 4.93 | 1.615 | 150.52 | 1 | 1 | 4 | 2 | 19 | |
| 20 | Toyota Corolla | 33.9 | 4 | 71.1 | 65 | 4.22 | 1.835 | 19.90 | 1 | 1 | 4 | 1 | 20 | |
| 21 | Toyota Corona | 21.5 | 4 | 120.1 | 97 | 3.70 | 2.465 | 20.01 | 1 | 0 | 3 | 1 | 21 | |
| 22 | Dodge Challenger | 15.5 | 8 | 318.0 | 150 | 2.76 | 3.520 | 16.87 | 0 | 0 | 3 | 2 | 22 | |
| 23 | AMC Javelin | 15.2 | 8 | 304.0 | 150 | 3.15 | 3.435 | 17.30 | 0 | 0 | 3 | 2 | 23 | |
| 24 | Camaro Z28 | 13.3 | 8 | 350.0 | 245 | 3.73 | 3.840 | 15.41 | 0 | 0 | 3 | 4 | 24 | |
| 25 | Pontiac Firebird | 19.2 | 8 | 400.0 | 175 | 3.08 | 3.845 | 17.05 | 0 | 0 | 3 | 2 | 25 | |
| 26 | Fiat X1-9 | 27.3 | 4 | 79.0 | 66 | 4.08 | 1.935 | 18.90 | 1 | 1 | 4 | 1 | 26 | |
| 27 | Porsche 914-2 | 26.0 | 4 | 120.3 | 91 | 4.43 | 2.140 | 16.70 | 0 | 1 | 5 | 2 | 27 | |
| 28 | Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 | 1.513 | 16.90 | 1 | 1 | 5 | 2 | 28 | |
| 29 | Ford Pantera L | 15.8 | 8 | 351.0 | 264 | 4.22 | 3.170 | 14.50 | 0 | 1 | 5 | 4 | 29 | |
| 30 | Ferrari Dino | 19.7 | 6 | 900.0 | 700 | 3.62 | 200.770 | 150.50 | 0 | 1 | 5 | 6 | 30 | |
| 31 | Maserati Bora | 15.0 | 8 | 301.0 | 335 | 3.54 | 3.570 | 14.60 | 0 | 1 | 5 | 8 | 31 | |
| 32 | Volvo 142E | 21.4 | 4 | 121.0 | 109 | 4.11 | 2.780 | 18.60 | 1 | 1 | 4 | 2 | 32 | |


use_case_2_anomaly_detection.R

Build a Deep Autoencoder

```
25 # -----
26 # Training an unsupervised deep neural network with autoencoder
27 # -----
28
29 # Use a bigger DNN
30 model <- h2o.deeplearning(x = 1:10,
31                           training_frame = h2o_mtcars,
32                           autoencoder = TRUE,
33                           activation = "RectifierWithDropout",
34                           hidden = c(50, 50, 50),
35                           epochs = 100)
36
37 # Calculate reconstruction errors (MSE)
38 errors <- h2o.anomaly(model, h2o_mtcars, per_feature = FALSE)
39 print(errors)
40 errors <- as.data.frame(errors)
41
42 # Plot
43 plot(sort(errors$Reconstruction.MSE), main = "Reconstruction Error")
44
45 # Outliers (define 0.09 as the cut-off point)
46 row_outliers <- which(errors > 0.09) # based on plot above
47 mtcars[row_outliers,]
```

Look at the MSE

Define cut-off



```
> row_outliers <- which(errors > 0.09) # based on plot above
> mtcars[row_outliers,]
```

| | X | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb | ID |
|----|------|------|------|------|-------|------|-----|------|------|----|------|------|------|
| 11 | Merc | 280C | 17.8 | 6 | 167.6 | 210 | 800 | 900 | 1000 | 1 | 0 | 4 | 4 11 |
| 18 | Fiat | 128 | 32.4 | 4 | 780.0 | 2100 | 400 | 200 | 700 | 1 | 1 | 4 | 1 18 |

Outliers identified

End of Second Talk – Thanks!

- Data Science Milan
- Gianmario Spacagna
- Politecnico di Milano
- Resources
 - bit.ly/h2o_milan_1
 - www.h2o.ai
 - docs.h2o.ai
- Contact
 - joe@h2o.ai
 - [@matlabulous](#)
 - github.com/woobe