

# Deep Water




Jo-fai (Joe) Chow  
Data Scientist  
joe@h2o.ai  
@matlabulous

Data Science Milan  
Politecnico di Milano  
10<sup>th</sup> October, 2016

# Agenda

- First Talk (25 mins)
  - About H2O.ai
  - Demo
    - A Simple Classification Task
    - H2O's Web Interface
  - Why H2O?
    - Our Community
    - Our Customers
  - What's Next?
    - New H2O Features
- Second Talk (25 mins)
  - H2O for IoT
    - Predictive Maintenance
    - Anomaly Detection
    - H2O's R Interface
- Third Talk (25 mins)
  - Deep Water
  - Demo
    - H2O + mxnet on GPU
    - H2O's Python Interface

# Deep Learning in H2O



# H2O Overview

Computer Science (CS)

Artificial Intelligence (A.I.)

Machine Learning (ML)

Deep Learning (DL)

hot hot hot hot hot

H<sub>2</sub>O.ai

# A Simple Neural Network




Image credit: <http://cs231n.github.io/>

# H2O Deep Learning in Action

116M rows, 6GB CSV file  
800+ predictors (numeric + categorical)

airlines\_all\_selected\_cols.hex

Actions: View Data Split Build Model Predict Download Export

| Rows      | Columns | Compressed Size |
|-----------|---------|-----------------|
| 116695259 | 12      | 2GB             |

## Job

Run Time 00:00:36.712

Remaining Time 00:00:17.188

Type Model

Key Q deeplearning-dd2f42f7-81f7-42e8-9d98-e34437309828

Description DeepLearning

Status RUNNING

Progress 69%

Iterations: 12. Epochs: 0.628821. Speed: 2,243,735 samples/sec. Estimated time left: 21.849 sec


Actions View Cancel Job

model trained in <1 min:  
2M+ samples/second

OUTPUT - STATUS OF NEURON LAYERS (PREDICTING ISDEPDELAYED, 2-CLASS CLASSIFICATION, BERNOULLI DISTRIBUTION, CROSSENTROPY LOSS, 17.462 WEIGHTS/BIASES, 221.9 KB, 106.585,365 TRAINING SAMPLES, MINI-BATCH SIZE 1)



| layer | units | type      | dropout | l1 | l2 | mean_rate | rate_RMS | momentum | mean_weight | weight_RMS | mean_bias | bias_RMS |
|-------|-------|-----------|---------|----|----|-----------|----------|----------|-------------|------------|-----------|----------|
| 1     | 887   | Input     | 0       |    |    |           |          |          |             |            |           |          |
| 2     | 20    | Rectifier | 0       | 0  | 0  | 0.0493    | 0.2020   | 0        | -0.0021     | 0.2111     | -0.0139   | 1.0036   |
| 3     | 20    | Rectifier | 0       | 0  | 0  | 0.0197    | 0.0227   | 0        | -0.1893     | 0.5362     | -1.3988   | 1.5259   |
| 4     | 20    | Rectifier | 0       | 0  | 0  | 0.0517    | 0.0446   | 0        | -0.1575     | 0.3068     | -0.0846   | 0.6046   |
| 5     | 20    | Rectifier | 0       | 0  | 0  | 0.0761    | 0.0844   | 0        | -0.0374     | 0.2275     | -0.2647   | 0.2481   |
| 6     | 2     | Softmax   | 0       | 0  | 0  | 0.0161    | 0.0083   | 0        | 0.0741      | 0.7260     | 0.4269    | 0.2056   |

ROC CURVE - VALIDATION METRICS, AUC = 0.702560



Threshold: Choose... Criterion: Choose...

VARIABLE IMPORTANCES



## Legend

Each bar represents one CPU.

Blue: idle time

Green: user time

Red: system time

White: other time (e.g. I/O)

10 nodes: all  
320 cores busy

H2O.ai Deep Learning Model

real-time, interactive  
model inspection in Flow



# H2O Deep Learning Community Quotes

**CIFAR-10 Competition  
Winners: Interviews with Dr.  
Ben Graham, Phil Culliton, &  
Zygmunt Zajac**

Triskelion | 01.02.2015

“I did really like H2O’s deep learning implementation in R, though - the interface was great, the back end extremely easy to understand, and it was scalable and flexible. Definitely a tool I’ll be going back to.”

[READ MORE](#)

**Kaggle challenge  
2nd place winner  
Colin Priest**

for creating this corpus. .  
do not contain Spanish sent-  
is a widespread major langu-  
reason was to create a corp-  
tasks. These tasks are com

Completed • Knowledge • 161 teams


**Denoising Dirty Documents**

Mon 1 Jun 2015 – Mon 5 Oct 2015 (3 months ago)

“For my final competition submission I used an ensemble of models, including 3 deep learning models built with R and h2o.”

[READ MORE](#)




# Why Deep Water?





# Deep Water: Best Open-Source Deep Learning




## Enterprise Deep Learning for Business Transformation

|  |   |   |
|--|---|---|
| Deep Water = THE Deep Learning Platform                            | H2O integrates the top open-source DL tools   |  |
| Native GPU support   |  is up to 100x faster than   |   |
| Enterprise Ready   | Easy to train and deploy, interactive, scalable, etc.<br>Flow, R, Python, Spark/Scala, Java, REST, POJO, <b>Steam</b>   |   |
| New Big Data Use Cases (previously impossible or difficult in H2O) | <b>Image</b> - social media, manufacturing, healthcare, ...<br><b>Video</b> - UX/UI, security, automotive, social media, ...<br><b>Sound</b> - automotive, security, call centers, healthcare, ...<br><b>Text</b> - NLP, sentiment, security, finance, fraud, ...<br><b>Time Series</b> - security, IoT, finance, e-commerce, ... |   |

# Deep Water opens the Floodgates for state-of-the-art Deep Learning

## H2O Deep Learning: simple multi-layer neural networks

1-5 layers  
MBs/GBs of data

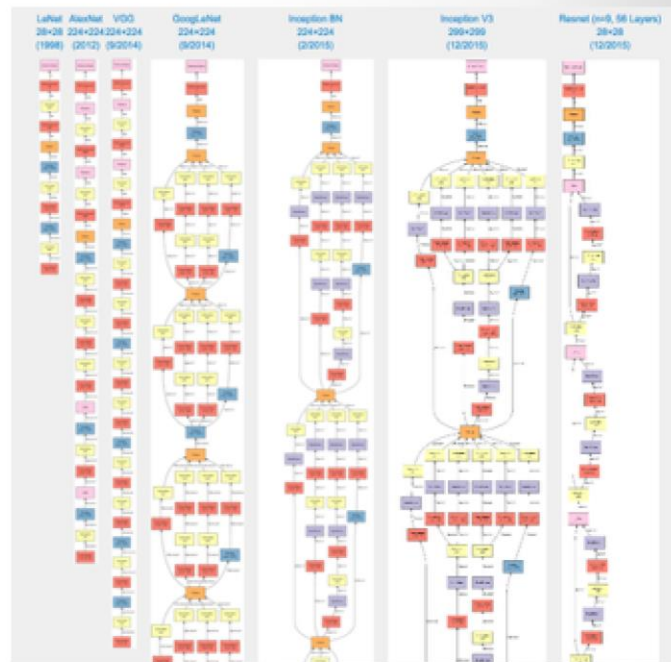



Limited to business analytics,  
statistical models (CSV data)

H<sub>2</sub>O.ai

## Deep Water: deep complex networks

5-1000 layers  
GBs/TBs of data



Large networks for big data  
(e.g. image 1000x1000x3 -> 3m inputs)

## Current Contributors (more H2O.ai folks joining soon)



Fabrizio Milo



Cyprien Noel



Qiang Kou



Arno Candel



Caffe



H<sub>2</sub>O.ai




This repository

Search

h2oai / deepwater

# Deep Water Demo

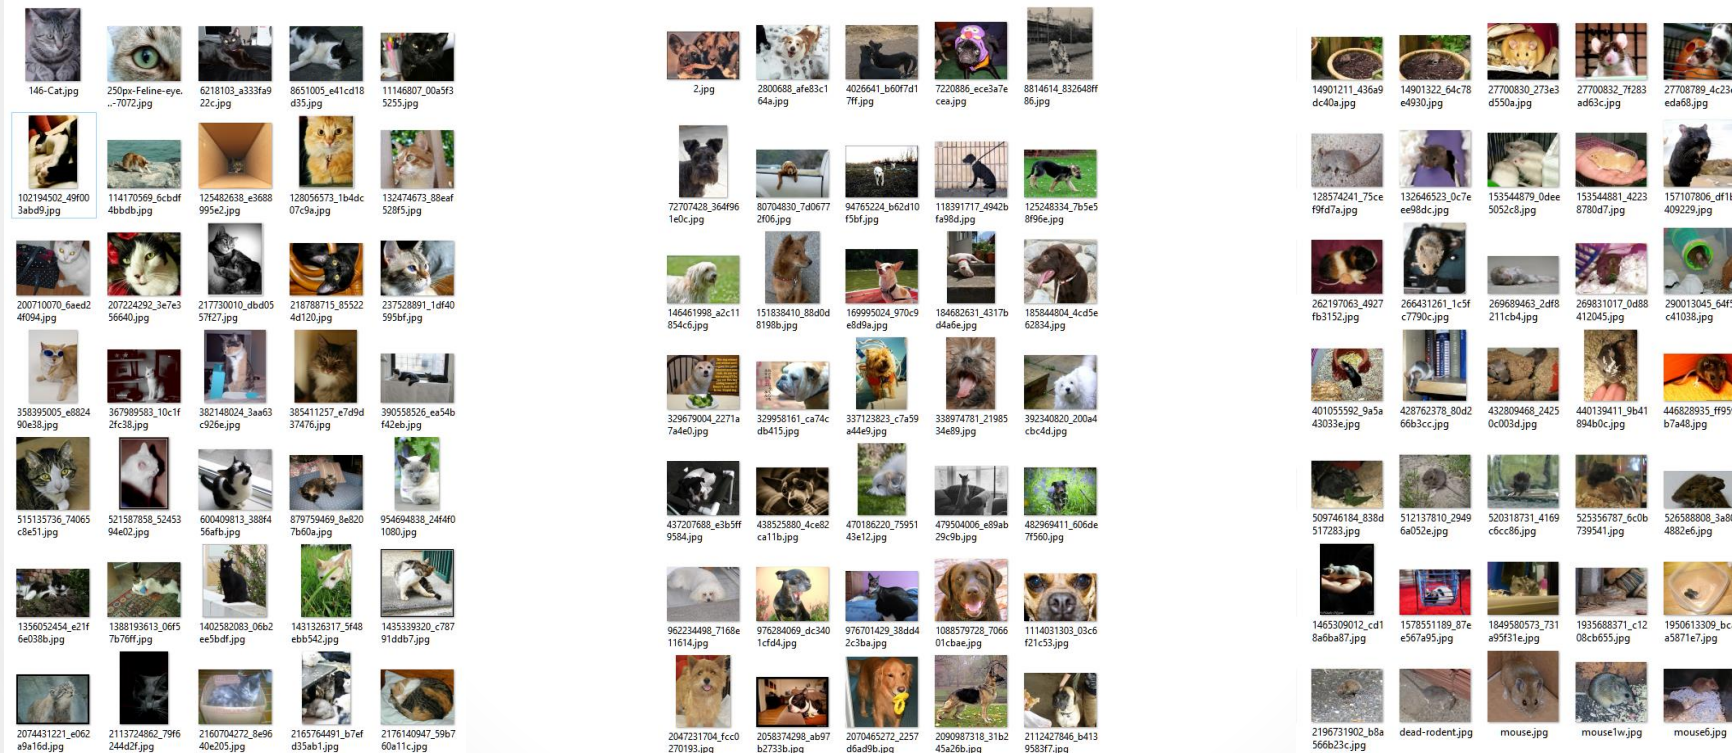


# Deep Water Demo

- H2O + mxnet
  - Dataset:
    - Cat / Dog / Mouse
  - H2O Python interface
  - mxnet GPU backend
  - Train a LeNet (CNN) model
  - Explore model in Flow
- Code and Data
  - [bit.ly/h2o\\_milan\\_1](https://bit.ly/h2o_milan_1)
  - subfolder
    - deep\_water\_demo



# Data – Cat/Dog/Mouse Images



# Data - CSV

|    | A  | B   |
|----|--|-----|
| 1  | bigdata/laptop/deepwater/imagenet/cat/102194502_49f003abd9.jpg   | cat |
| 2  | bigdata/laptop/deepwater/imagenet/cat/11146807_00a5f35255.jpg    | cat |
| 3  | bigdata/laptop/deepwater/imagenet/cat/1140846215_70e326f868.jpg  | cat |
| 4  | bigdata/laptop/deepwater/imagenet/cat/114170569_6cbdf4bbdb.jpg   | cat |
| 5  | bigdata/laptop/deepwater/imagenet/cat/1217664848_de4c7fc296.jpg  | cat |
| 6  | bigdata/laptop/deepwater/imagenet/cat/1241603780_5e8c8f1ced.jpg  | cat |
| 7  | bigdata/laptop/deepwater/imagenet/cat/1241612072_27ececbbdef.jpg | cat |
| 8  | bigdata/laptop/deepwater/imagenet/cat/1241613138_ef1d82973f.jpg  | cat |
| 9  | bigdata/laptop/deepwater/imagenet/cat/1244562192_35becd66bd.jpg  | cat |
| 10 | bigdata/laptop/deepwater/imagenet/cat/125482638_e3688995e2.jpg   | cat |
| 11 | bigdata/laptop/deepwater/imagenet/cat/128056573_1b4dc07c9a.jpg   | cat |
| 12 | bigdata/laptop/deepwater/imagenet/cat/12945197_75e607e355.jpg    | cat |
| 13 | bigdata/laptop/deepwater/imagenet/cat/132474673_88eaf528f5.jpg   | cat |
| 14 | bigdata/laptop/deepwater/imagenet/cat/1350530984_ecf3039cf0.jpg  | cat |
| 15 | bigdata/laptop/deepwater/imagenet/cat/1351606235_c9fbeb634.jpg   | cat |
| 16 | bigdata/laptop/deepwater/imagenet/cat/1356052454_e21f6e038b.jpg  | cat |
| 17 | bigdata/laptop/deepwater/imagenet/cat/1388193613_06f57b76ff.jpg  | cat |

# H2O + mxnet Demo



# End of Third Talk – Thanks!

- Data Science Milan
- Gianmario Spacagna
- Politecnico di Milano
- Resources
  - [bit.ly/h2o\\_milan\\_1](https://bit.ly/h2o_milan_1)
  - [www.h2o.ai](http://www.h2o.ai)
  - [docs.h2o.ai](http://docs.h2o.ai)
- Contact
  - [joe@h2o.ai](mailto:joe@h2o.ai)
  - [@matlabulous](#)
  - [github.com/woobe](https://github.com/woobe)