

# Kaggle tips

Dmitry Larko, Sr. Data Scientist @ H2O.ai  
dmitry@h2o.ai

December 15<sup>th</sup>, 2016



**H<sub>2</sub>O.ai**

# About me



## Dmitry Larko

Sr. Data Scientist at H2O.ai

San Francisco Bay Area, CA, United States

Joined 4 years ago · last seen in the past day



<http://h2o.ai>



**Competitions Grandmaster**

[Home](#)

[Competitions \(32\)](#)

[Kernels \(0\)](#)

[Discussion \(31\)](#)

[Datasets \(0\)](#)

[More](#)

[Edit Profile](#)

Competitions Grandmaster



Current Rank

**36**

of 52,561

Highest Rank

**25**



9



8



6

Kernels Contributor



Unranked



0



0



0

Discussion Contributor



Unranked



0



4



12

# Bio

- About 10 years working experience in DB/Data Warehouse field. Until 4 years ago I learned about Kaggle from my dad, who competing on Kaggle as well (and does that better than me)
- My first competition was [Amazon.com - Employee Access Challenge](#), I placed 10th out of 1687, learned a tons of new techniques/algorithms in ML field in a month and I got addicted to Kaggle.
- Until now I participated in 32 completions, was 2nd twice, won once and I am in Kaggle top- 100 Data Scientists.
- Currently I'm working as Senior Data Scientist at H2O.ai

# Motivation. Why to participate?

- **To win** – this is the best possible motivation for competition;
- **To learn** – Kaggle is a good place to learn and the best places to learn on Kaggle are forum and kernels from past and current competitions;
- **Looks good in resume** – well... only if you're constantly winning 😊

# How to start?

- Learn Python
- MOOC for Machine learning ([Coursera](#), [Udacity](#), [Harvard](#))
- Participate in Kaggle “Getting started” and “Playground” competitions
- Visit Kaggle finished competitions and go through winner’s solution posted at competition’s forum

# Kaggle's Toolset (1 of 2)

- Scikit-Learn ([scikit-learn.org](https://scikit-learn.org)). Simple and efficient tools for data mining and data analysis. A lot of tutorials, many ML algorithms have scikit-learn implementation.
- XGBoost ([github.com/dmlc/xgboost](https://github.com/dmlc/xgboost)). An optimized general purpose gradient boosting library. The library is parallelized (OpenMP). It implements machine learning algorithm under gradient boosting framework, including generalized linear model and gradient boosted regression tree (GBDT).
  - Theory behind XGBoost:
    - <https://homes.cs.washington.edu/~tqchen/pdf/BoostedTree.pdf>
  - Tutorial:
    - <https://www.kaggle.com/tqchen/otto-group-product-classification-challenge/understanding-xgboost-model-on-otto-data>

# Kaggle's Toolset (2 of 2)

- H2O ([h2o.ai](https://h2o.ai)). Fast Scalable Machine Learning API. Has state-of-the-art models like Random Forest and Gradient Boosting Trees. Allows you to work with really big datasets on Hadoop cluster. It also works on Spark! Check out Sparkling Water: [h2o.ai/product/sparkling-water/](https://h2o.ai/product/sparkling-water/)
- Neural Nets/Deep Learning
  - Keras ([github.com/fchollet/keras](https://github.com/fchollet/keras))
  - MXNet([github.com/dmlc/mxnet](https://github.com/dmlc/mxnet))

# Kaggle's Toolset - Advanced

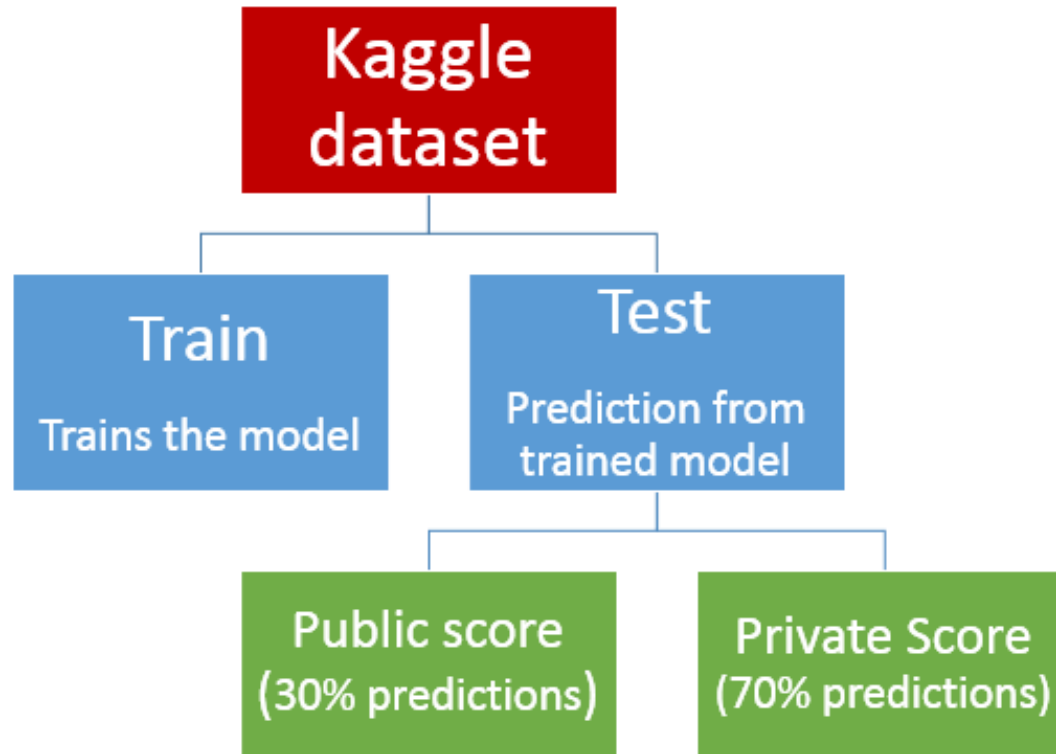
- Vowpal Wabbit (GitHub) – fast and state-of-the-art online learner. A great tutorial how to use VW for NLP is available here: [github.com/hal3/vwnlp](https://github.com/hal3/vwnlp)
- LibFM ([libfm.org](https://libfm.org)) - Factorization machines (FM) are a generic approach that allows to mimic most factorization models by feature engineering. Works great for sparse wide datasets, has a few competitors:
  - FastFM - [ibayer.github.io/fastFM/index.html](https://ibayer.github.io/fastFM/index.html)
  - pyFM - [github.com/coreylynch/pyFM](https://github.com/coreylynch/pyFM)
- Regularized Greedy Forest ([github.com/baidu/fast\\_rgf](https://github.com/baidu/fast_rgf)) - tree ensemble learning method, can be better than XGBoost, but you need to know how to cook it.
- Light GBM ([github.com/Microsoft/LightGBM](https://github.com/Microsoft/LightGBM)) – FAST gradient boosting framework based on decision tree algorithms.
- Four leaf clover, gives +100 to Luck



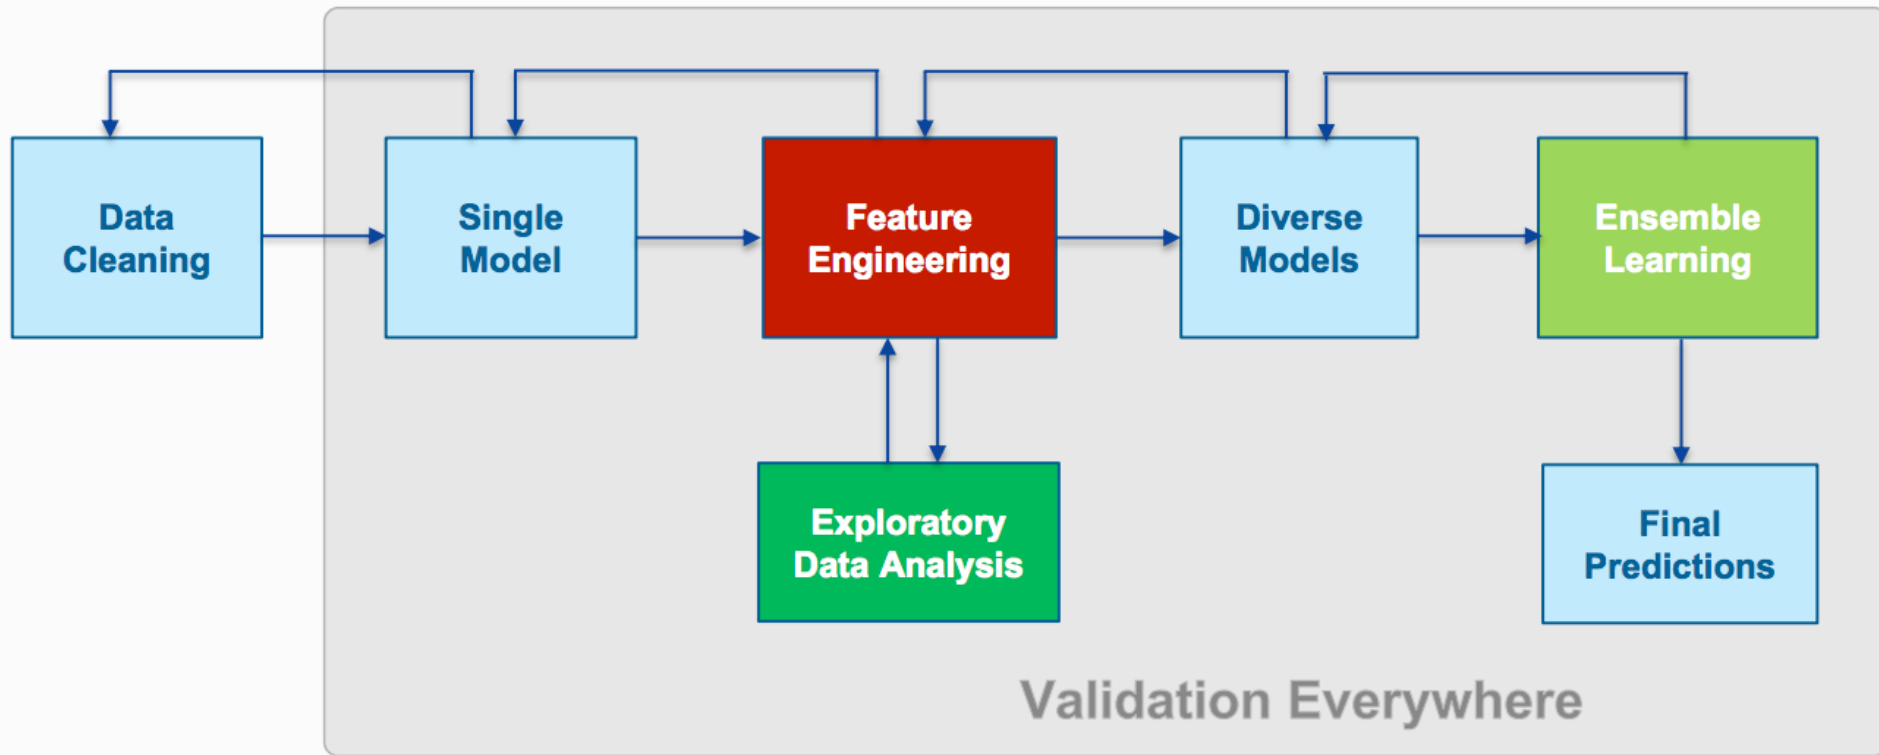
# Which past competition to check?

- Well, all of them of course. But if I would need to choose I'd selected these 4:
  - Greek Media Monitoring Multilabel Classification (WISE 2014). Why?
    - Interesting problem, a lot of classes.
  - Allstate Claims Severity. Why?
    - A great dataset to learn and polish you ensembling techniques
  - Caterpillar Tube Pricing. Why?
    - A good example of business problem to solve
    - Need some work with data before you can build good models
    - Has a “data leakage” you can find and exploit
  - National Data Science Bowl. Why?
    - Good Deep Learning competition to start

# Kaggle Competition Dataset and Rules



# Pipeline



# Validation

- Main idea you need to model competition train/test split, so you can test all your ideas locally:
  - Train/test is random:
    - Dataset is big or no access to good hardware:
      - Train/test split
    - Dataset is small and/or we have enough computational power:
      - N-fold cross validation
  - Train/test split is not random (good example is time-series):
    - Train/test split

# Train/Test split

Original Data




Training Set

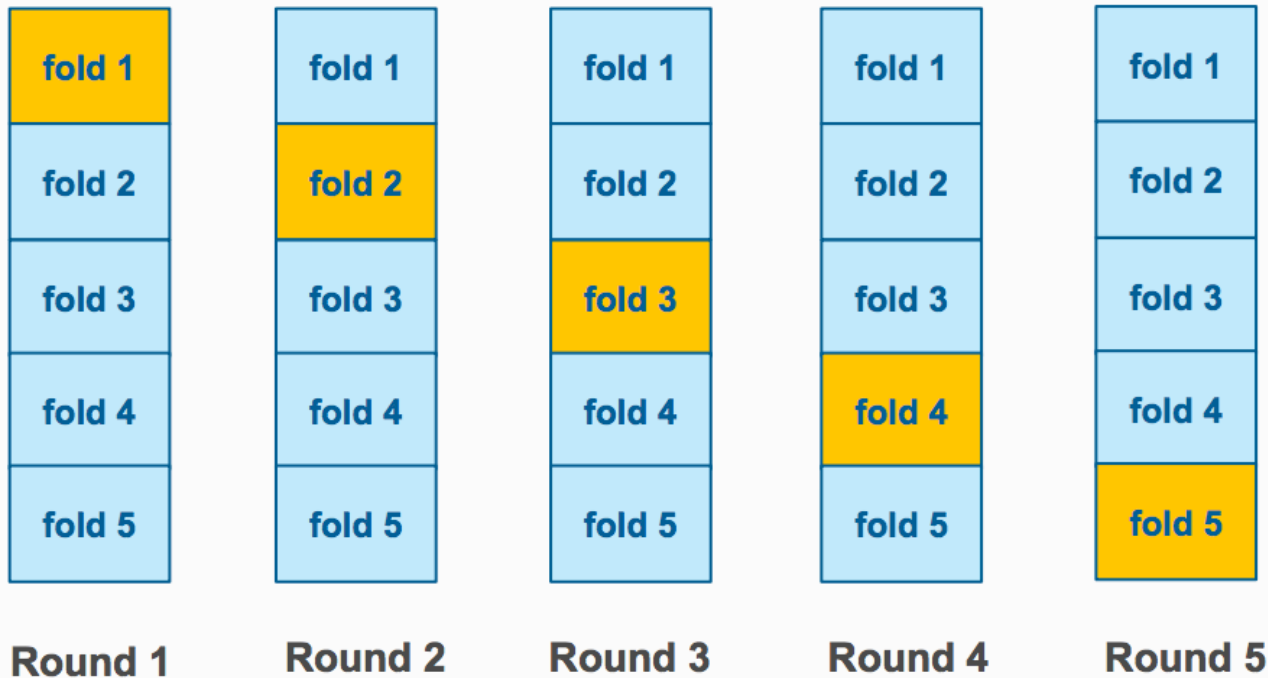

Test Set

--	--	--	--	--



# Train/validation split tips

- Rule-of-thumb: 80/20 or 70/30
- How to identify best validation set?
  - LB feedback
  - Adversarial validation. Combine train and test sets into one and train classifier to classify train and test examples correctly. When choose a number of misclassified examples that the model was most certain about. It means that they look like test examples but in reality are training examples
    - [fastml.com/adversarial-validation-part-one](https://fastml.com/adversarial-validation-part-one)
    - [fastml.com/adversarial-validation-part-two](https://fastml.com/adversarial-validation-part-two)

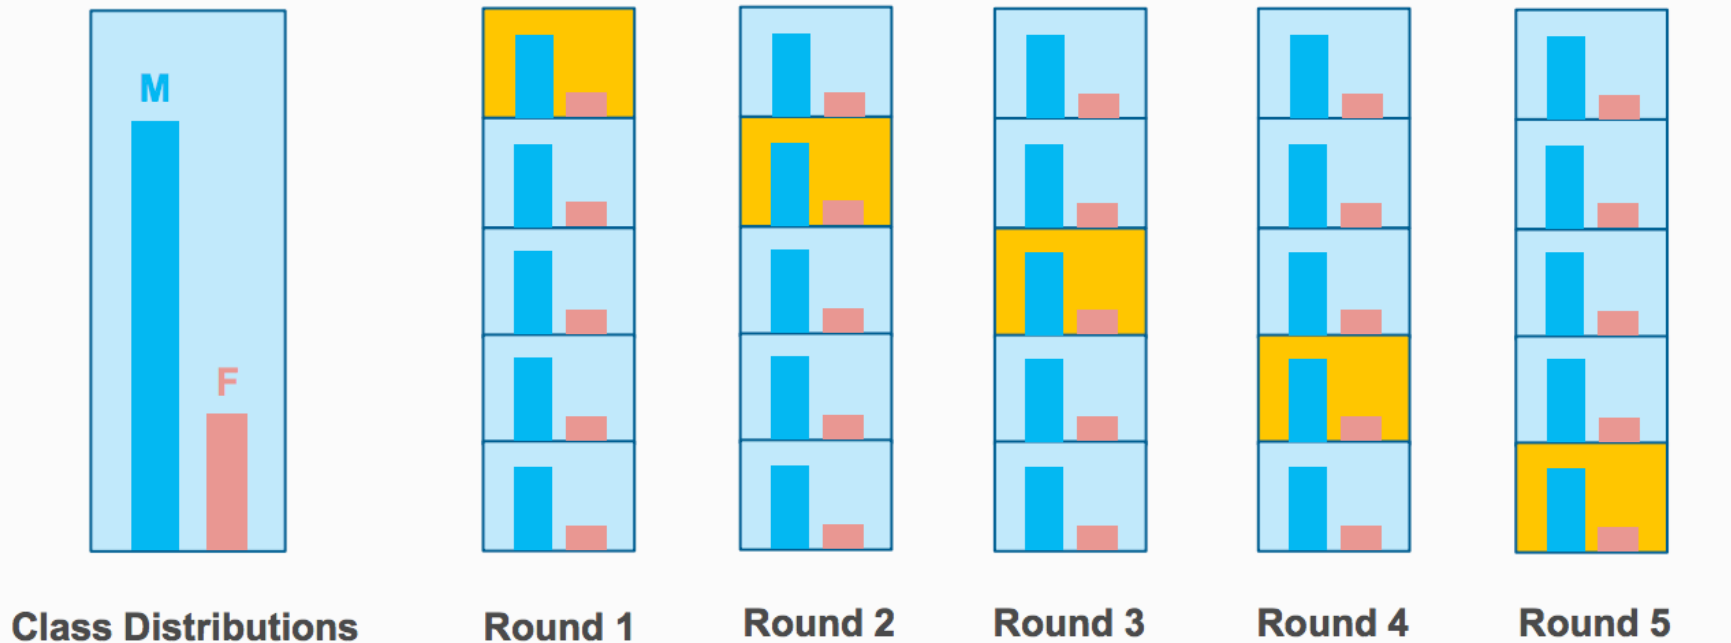
# 5-fold Cross-Validation





*score(CV) = the average of evaluation scores from each fold*  
*You can also repeat the process many times!*

 Training Data  
 Validation Data

# Stratified 5-foldCross-Validation



*Keep the distribution of classes in each fold*

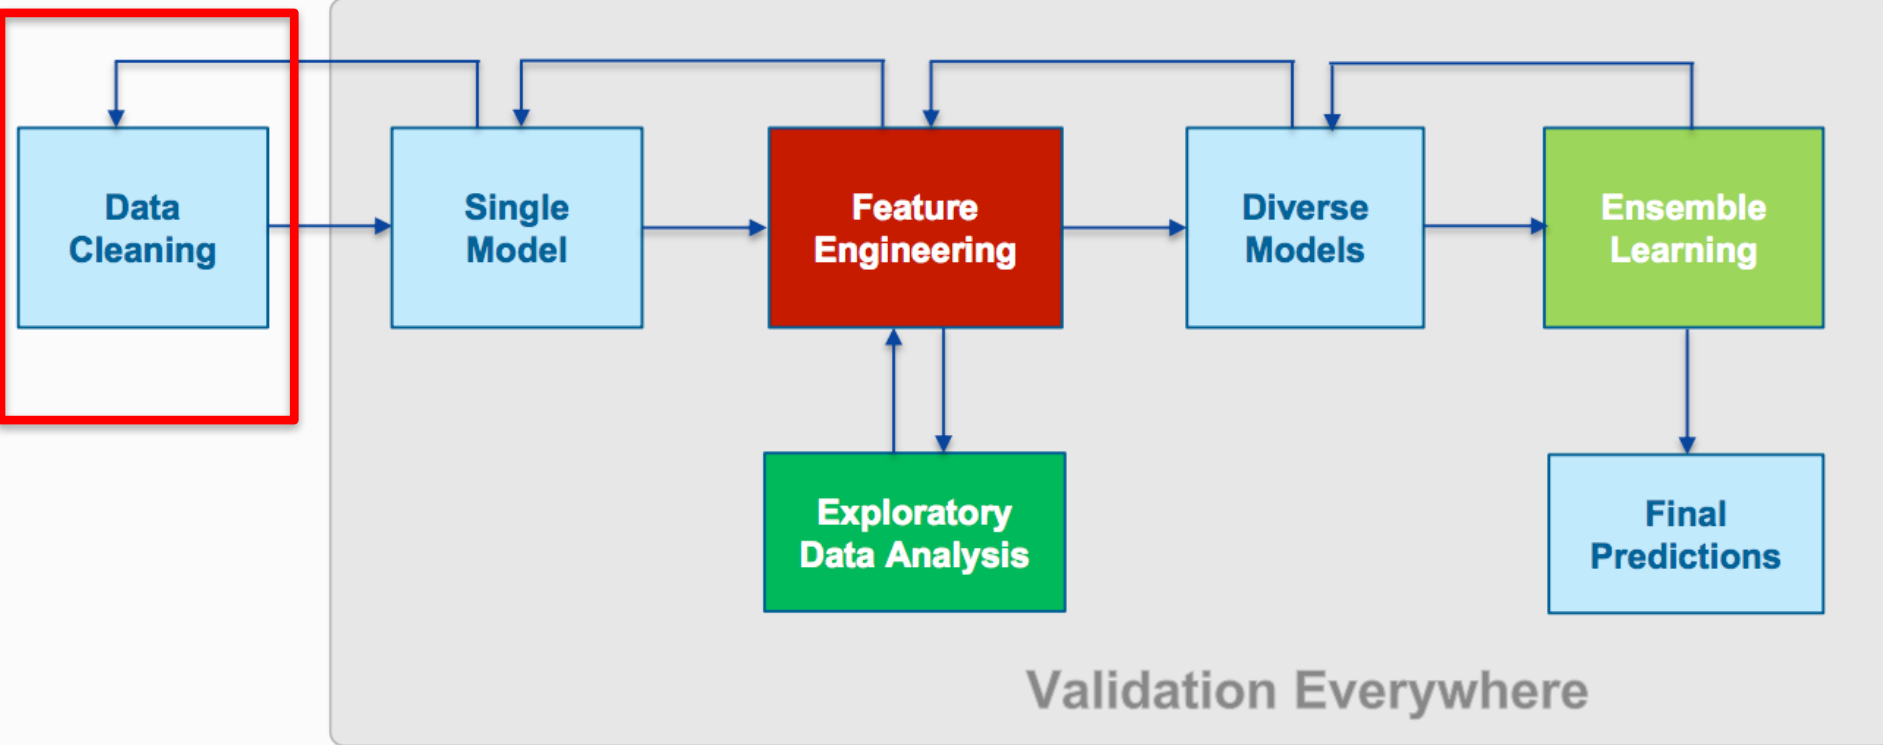
 **Training Data**  
 **Validation Data**



# Cross Validation tips

- It is normal to experience big shake up on private LB if not using local CV correctly
- 5 or 10 folds are not always the best choices (you need to consider the cost of computation time for training models)
  - Which K to choose?
    - Depends on your data
    - Mimic the ratio of training and testing in validation process
    - Find a K with lowest gap between local CV and public LB scores
    - Standard deviation of K-fold CV score matters more than mean:
      - To find balanced cross-validation: train the model, measure scores and select CV with smaller standard deviation
  - Stratified K-fold CV is important for imbalanced dataset, especially for classification problems.

# Data cleaning



# Data cleaning

- Data cleaning is the removal of duplicates, useless data, or fixing of missing data
- Reduce dimensional complexity of dataset
  - Make training faster without (significant) hurting the performance
- Apply imputation methods to help (hopefully) utilize incomplete rows
  - Incomplete rows may contain relevant features (don't just drop them!)
  - In the risk of distorting original data, so be cautious!

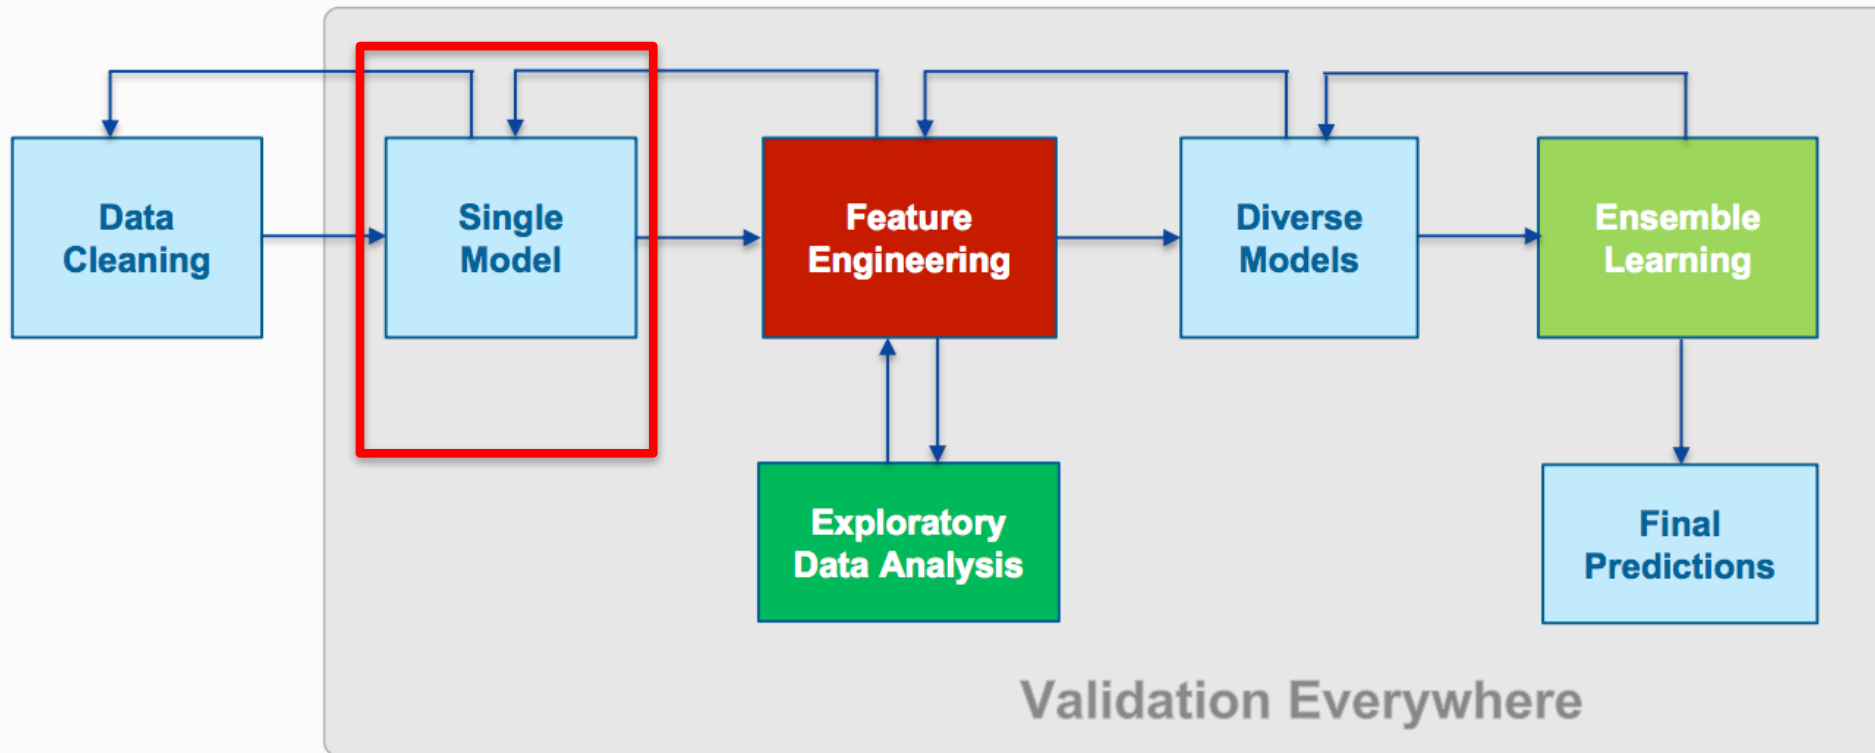
# Data cleaning

- Remove duplicate features
- Columns having the same value distribution or variance
  - Only need to keep one of them
- Remove constant features
  - Columns with only one unique value
  - Remove features with near-zero variance
- Remove features with significant amount of missing values
  - Try to use it first in the models, of course

# Data cleaning

- Some machine learning tools cannot accept NAs in the input
- Encode missing values to avoid NAs
- Binary features
  - -1 for negatives, 0 for missing values and 1 for positives
- Categorical features
  - Encode as an unique category
    - “Unknown”, “Missing”, ....
- Numeric features
  - Tree-based methods
    - Encode as a big positive or negative number
      - 999, -99999, ....
      - $\max(x) + 1$ ,  $\min(x) - 1$
  - Linear, neural nets, etc. methods
    - Encode by splitting into 2 columns:
      - Binary column isNA (0 if not and 1 if yes)
      - In original column replace NAs by mean or median

## Building first model



# Mostly Used ML models

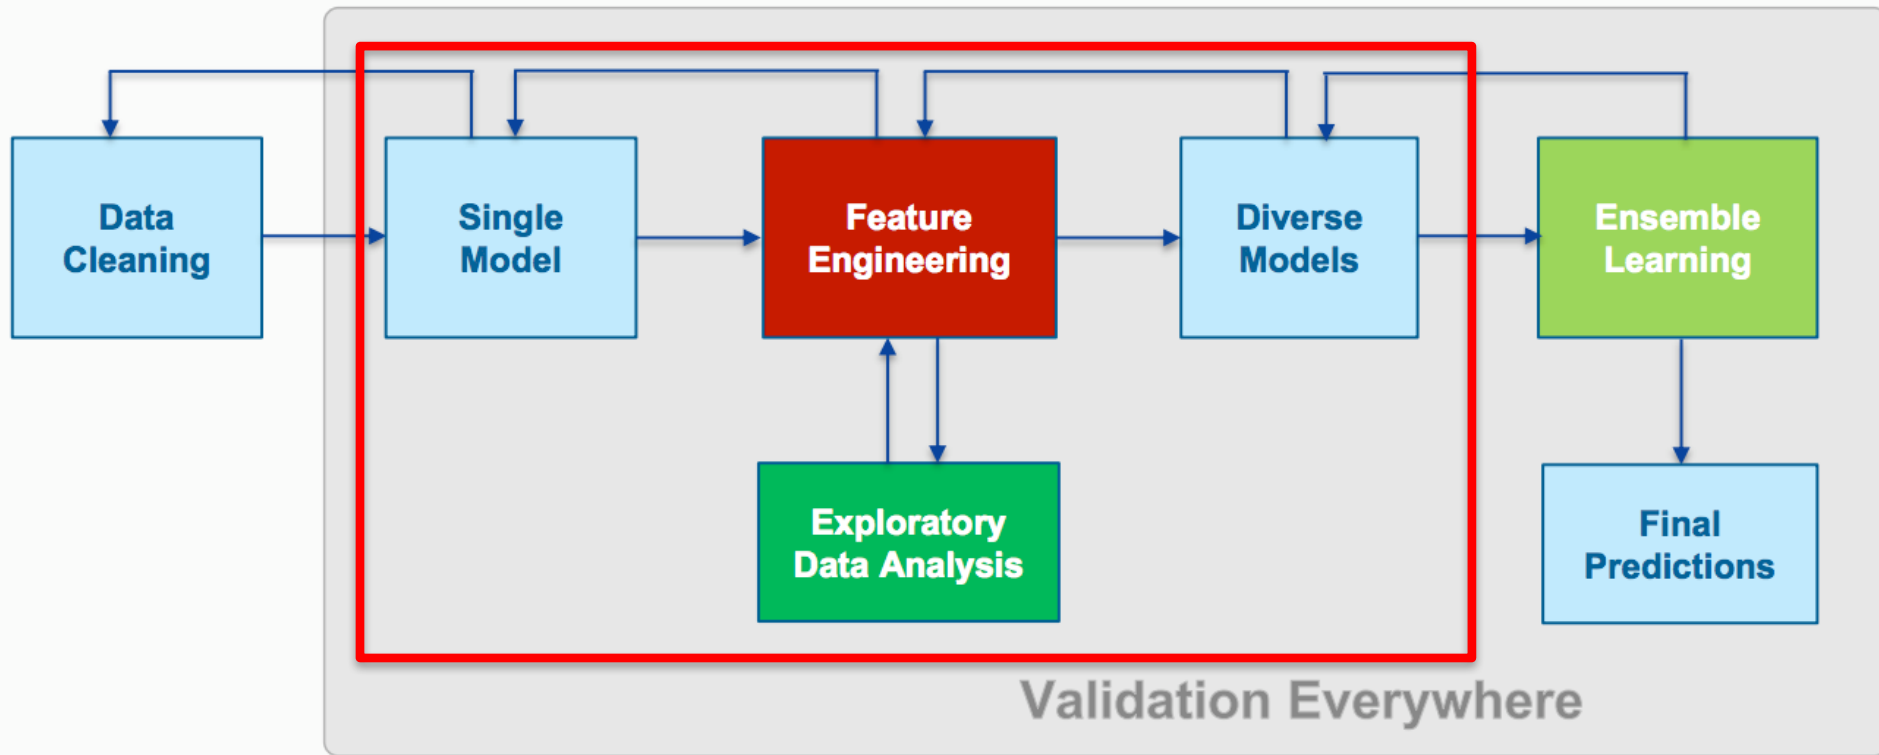
Model type	Name	Package
Regression	Linear Regression	sklearn.linear_model.LinearRegression
	Ridge Regression	sklearn.linear_model.Ridge
	Lasso Regression	sklearn.linear_model.Lasso
Instance-based	K Nearest Neighbors	sklearn.neighbors.KNeighborsClassifier sklearn.neighbors.KNeighborsRegressor
	Support Vector Machines (SVM)	sklearn.svm.SVC, sklearn.svm.SVR sklearn.svm.LinearSVC, sklearn.svm.LinearSVR
Hyperplane-based	Naive Bayes	sklearn.naive_bayes.GaussianNB sklearn.naive_bayes.MultinomialNB sklearn.naive_bayes.BernoulliNB
	Logistic Regression	sklearn.linear_model.LogisticRegression
Ensemble trees	Random Forests	H2O Distributed Random Forest
	Extremely Randomized Trees	sklearn.ensemble.ExtraTreesClassifier sklearn.ensemble.ExtraTreesRegressor
	Gradient Boosting Machines (GBM)	H2O GBM; Xgboost; LightGBM RGF

# Mostly Used ML models cont'd

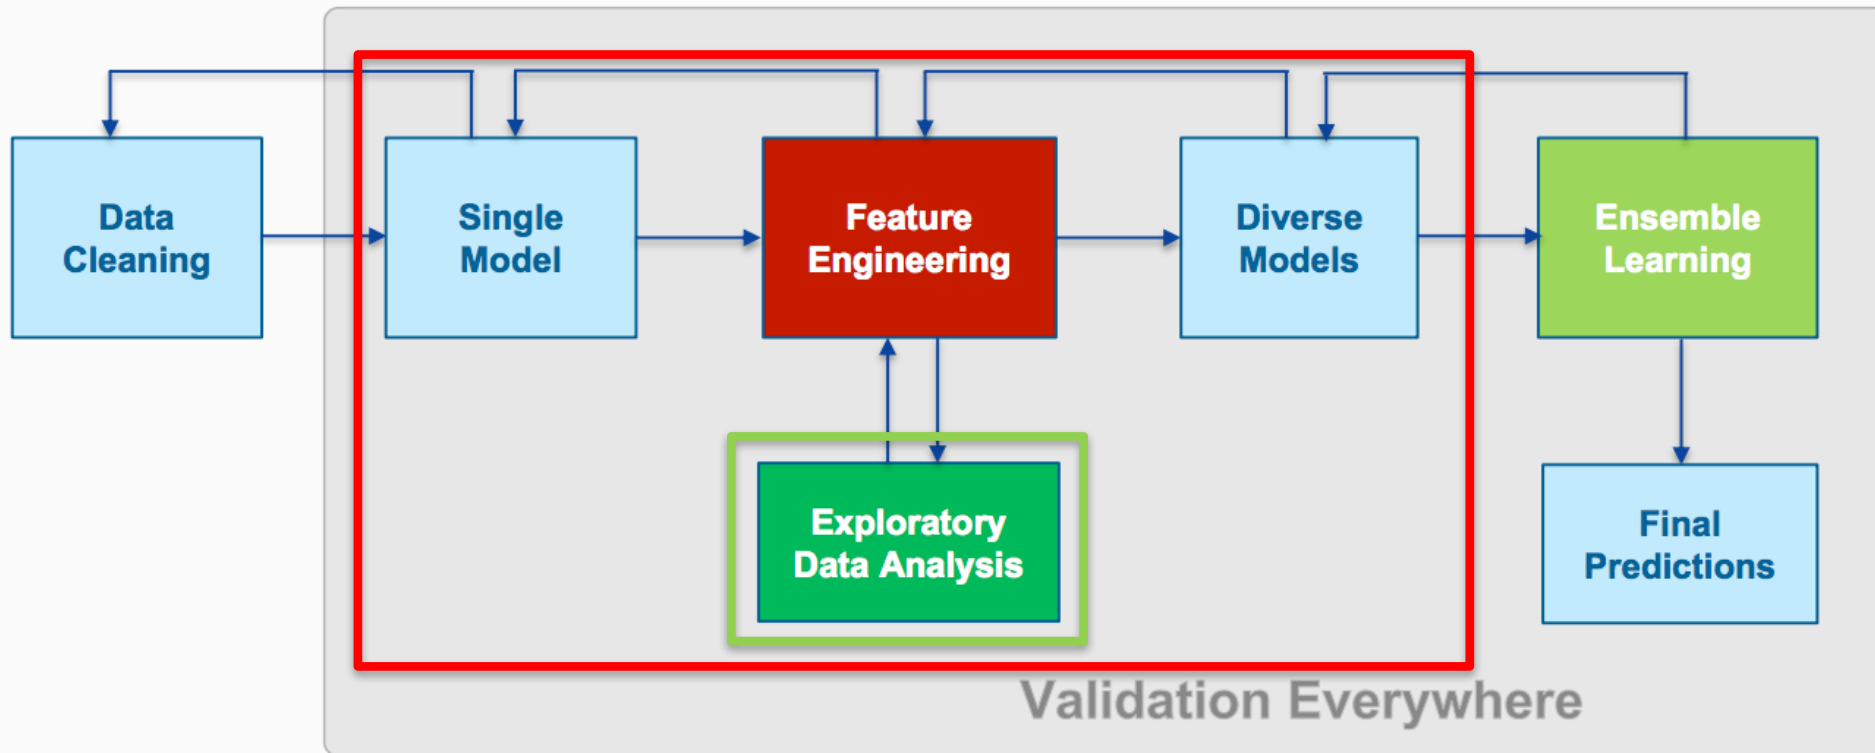
Model type	Name	Package
Neural Network	MLP CNN RNN	Keras MXNet
Recommendation	Matrix Factorization	sklearn.decomposition.NMF
	Factorization machines	pyFM
Clustering	K-Means	sklearn.cluster.MiniBatchKMeans H2O k-means (auto k!!!)
	(H)DBSCAN	sklearn.cluster.DBSCAN hdbscan
Dimensionality reduction	tSNE	tsne
	PCA	sklearn.decomposition.PCA
	Autoencoder	H2O autoencoder



# Main cycle



# Exploratory data analysis



# Correlation Graph

The nodes of this graph are the variables in a data set. The weights between the nodes are defined by the absolute value of their pairwise Pearson correlation.

To create:

- calculate Pearson correlation between columns/variables
- build undirected graph where each node is a column/variable
- connection weights between nodes are defined by Pearson correlation absolute values; weights below a certain threshold are not displayed
- node size is determined by number of connections (node degree)
- node color is determined by a graph communities calculation
- node position is defined by a graph force field algorithm

Free graph software: <https://gephi.org/>



# 2-D projections



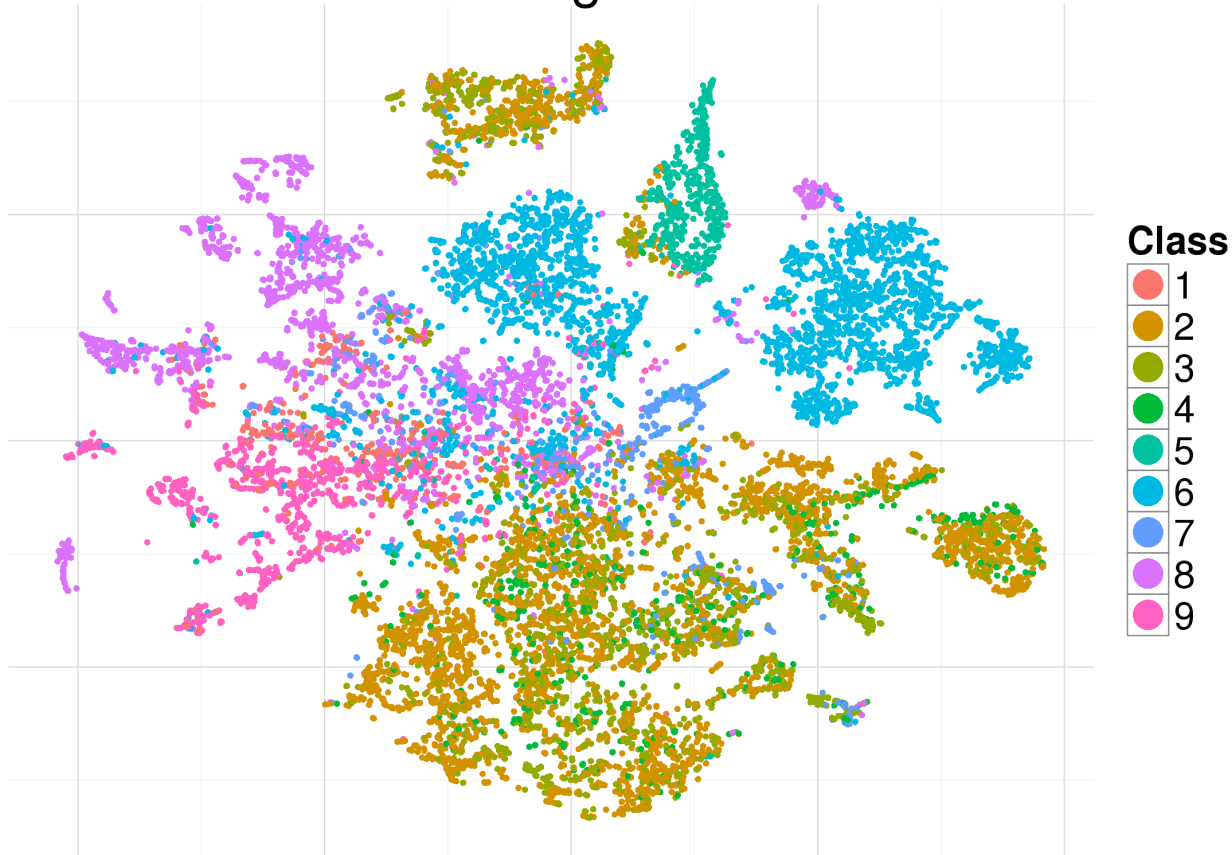
786 dimensions to 2 dimensions  
with PCA



786 dimensions to 2 dimensions  
with autoencoder

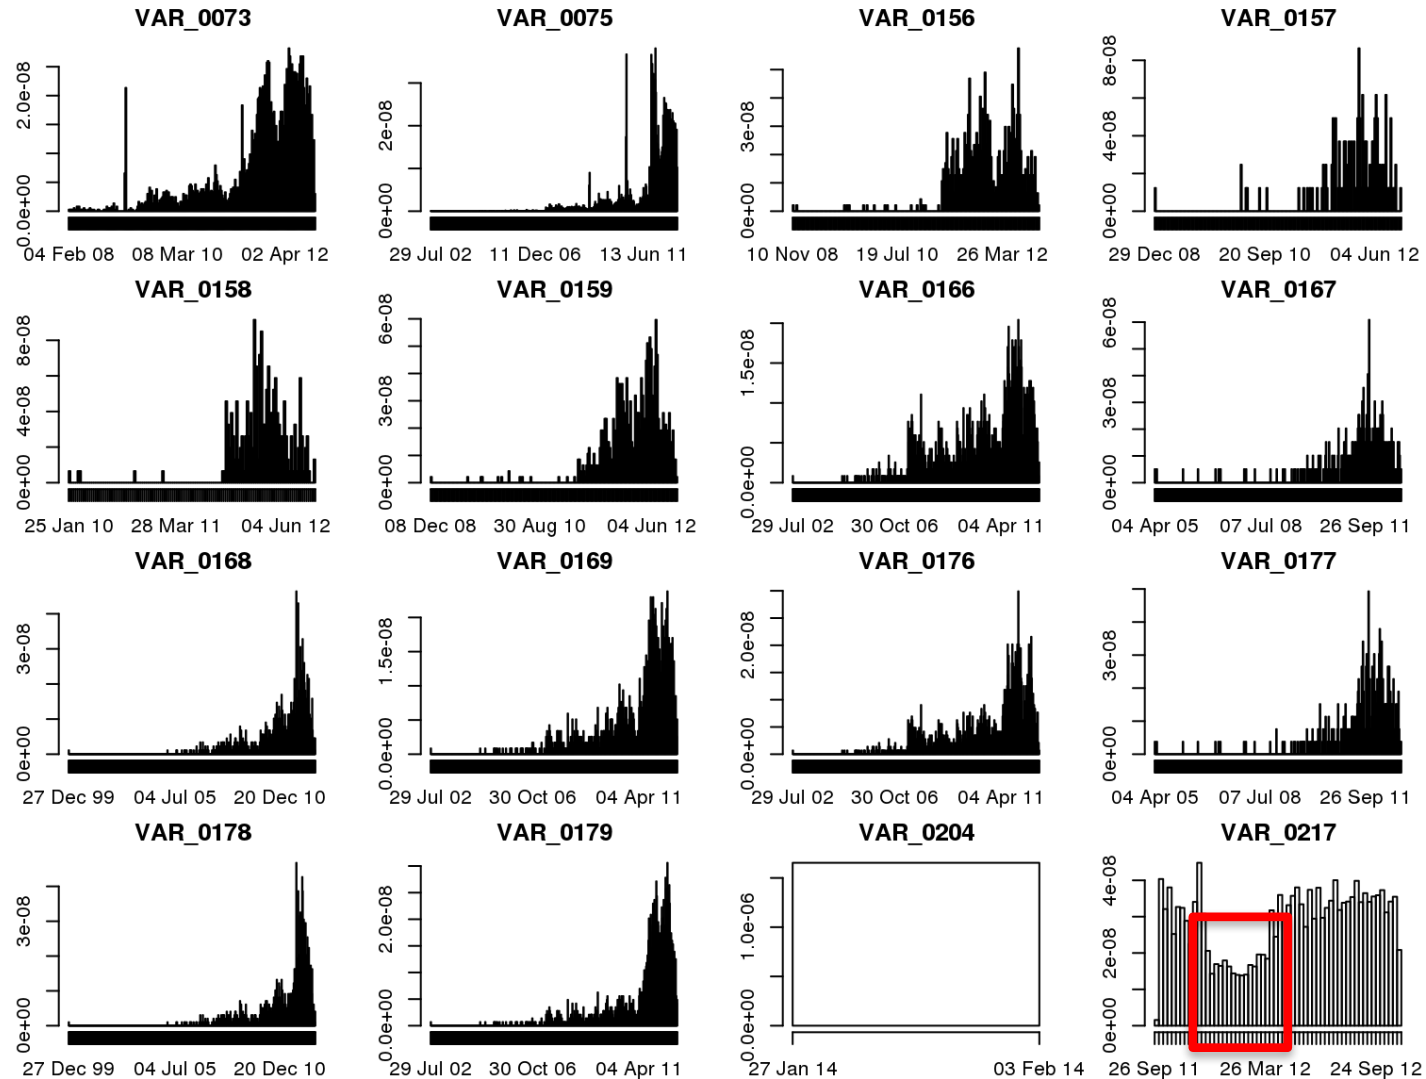
# 2-D projections

t-SNE 2D Embedding of Products Data



<https://www.kaggle.com/benhamner/otto-group-product-classification-challenge/t-sne-visualization>

# Distributions

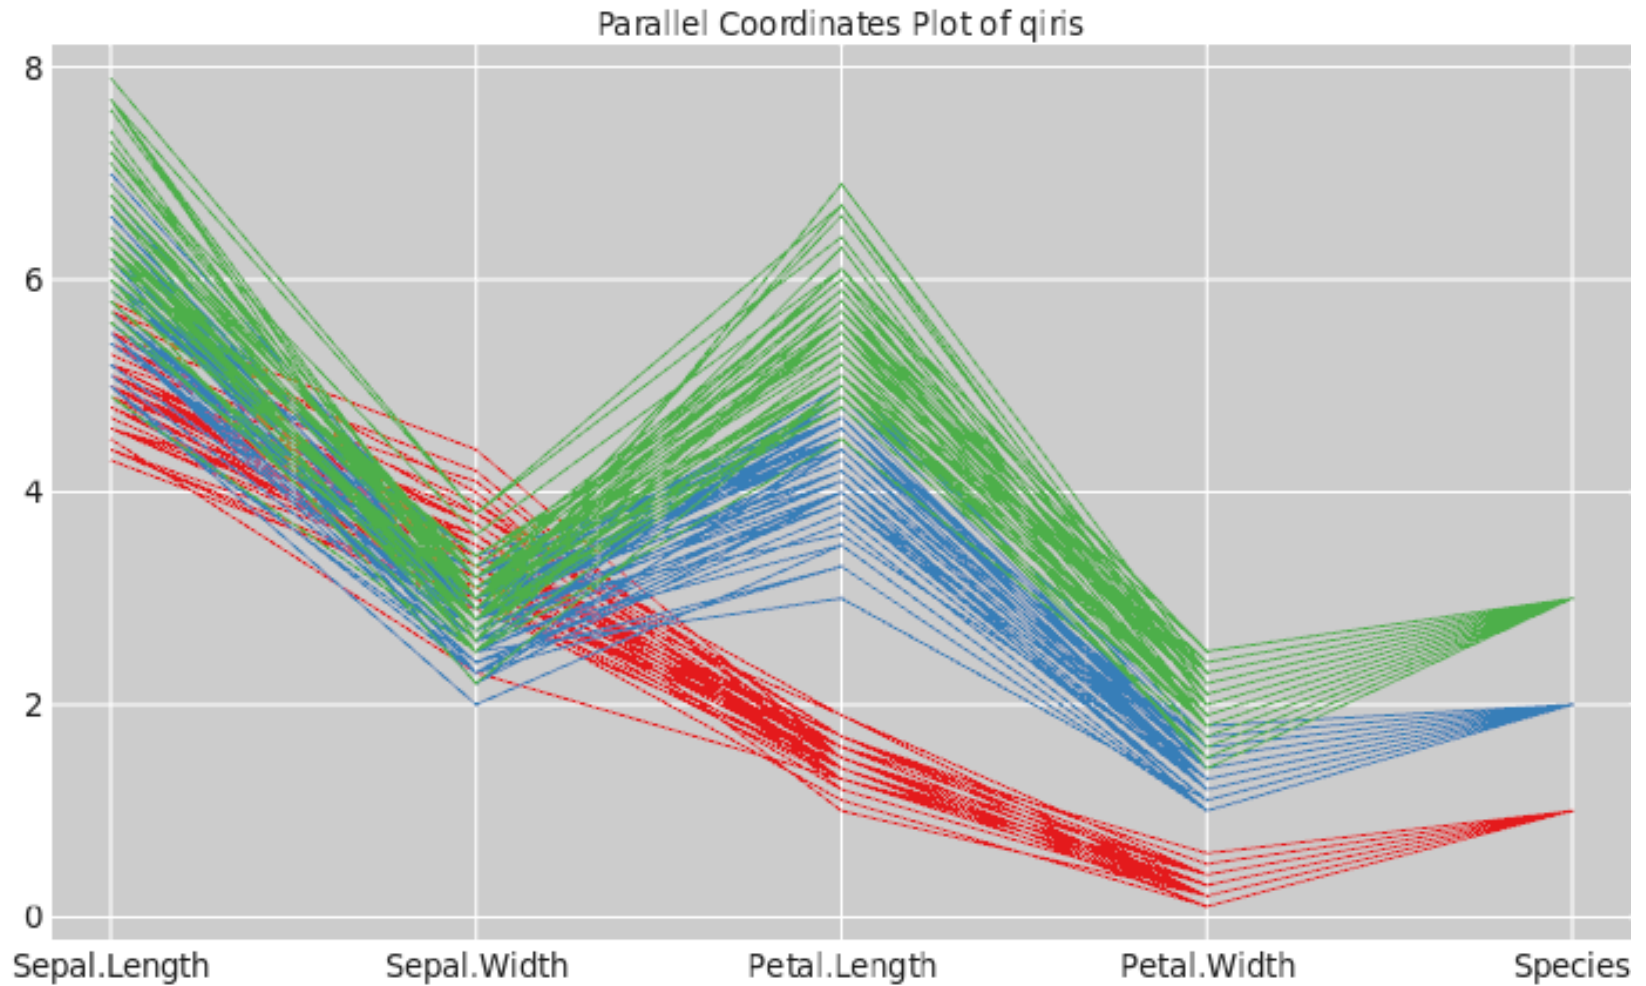


<https://www.kaggle.com/darraghdog/springleaf-marketing-response/explore-springleaf/notebook>

# Pair plots

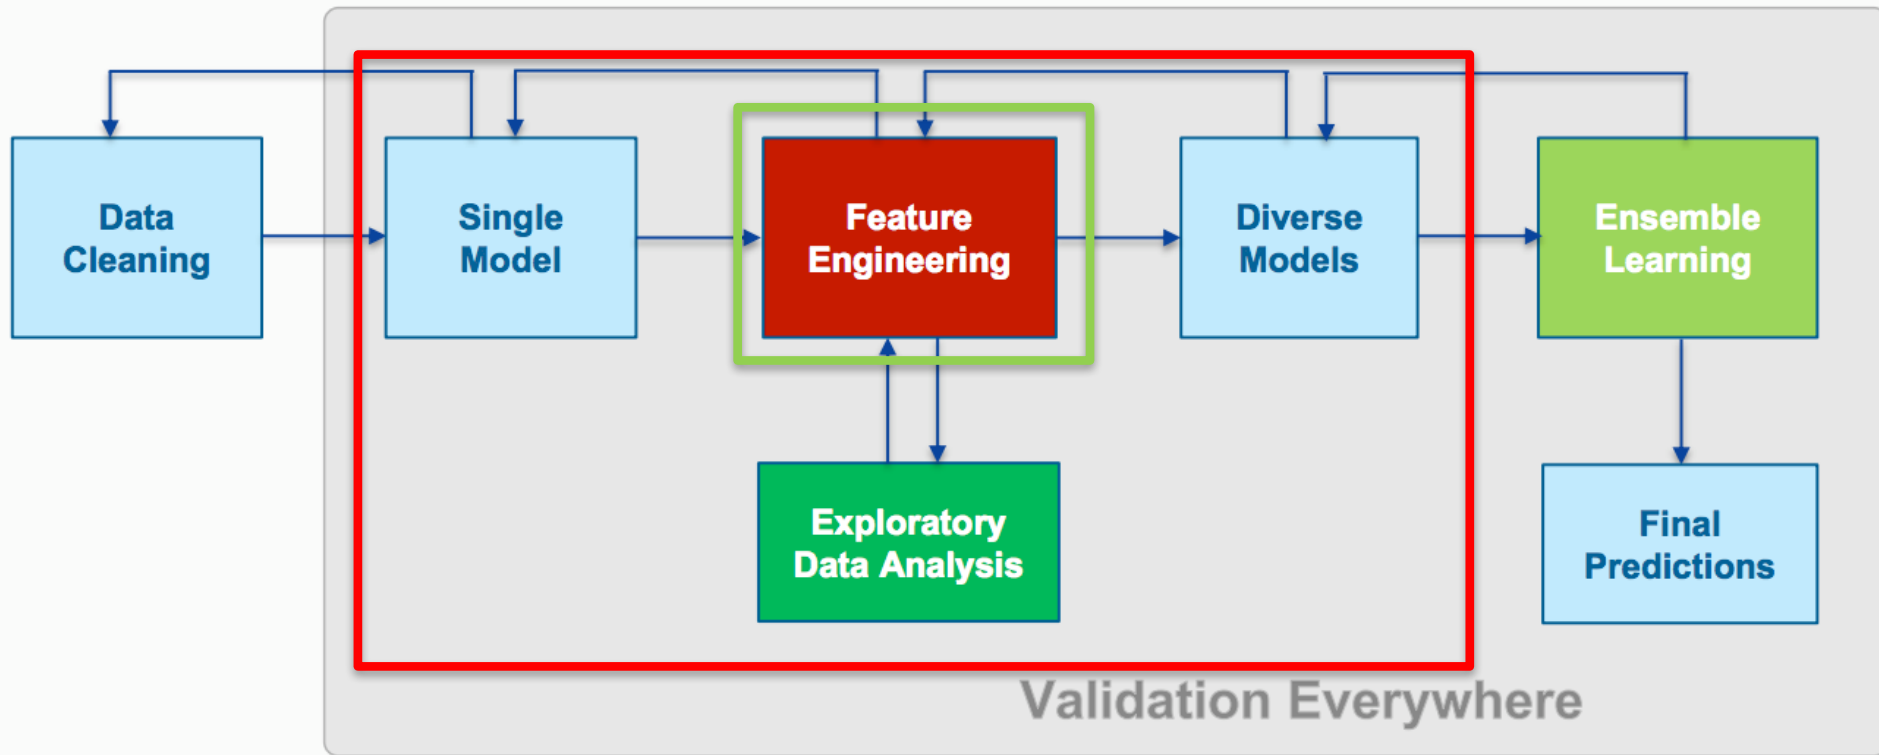


# Parallel coordinates plot





# Feature Engineering



# Feature Engineering

- Extract more new gold features, remove irrelevant or noisy features
  - Simpler models with better results
- The most important factor for the success of machine learning
- Key Elements
  - Data Transformation
  - Feature Encoding
  - Feature Extraction
  - Feature Selection

# Data Transformation

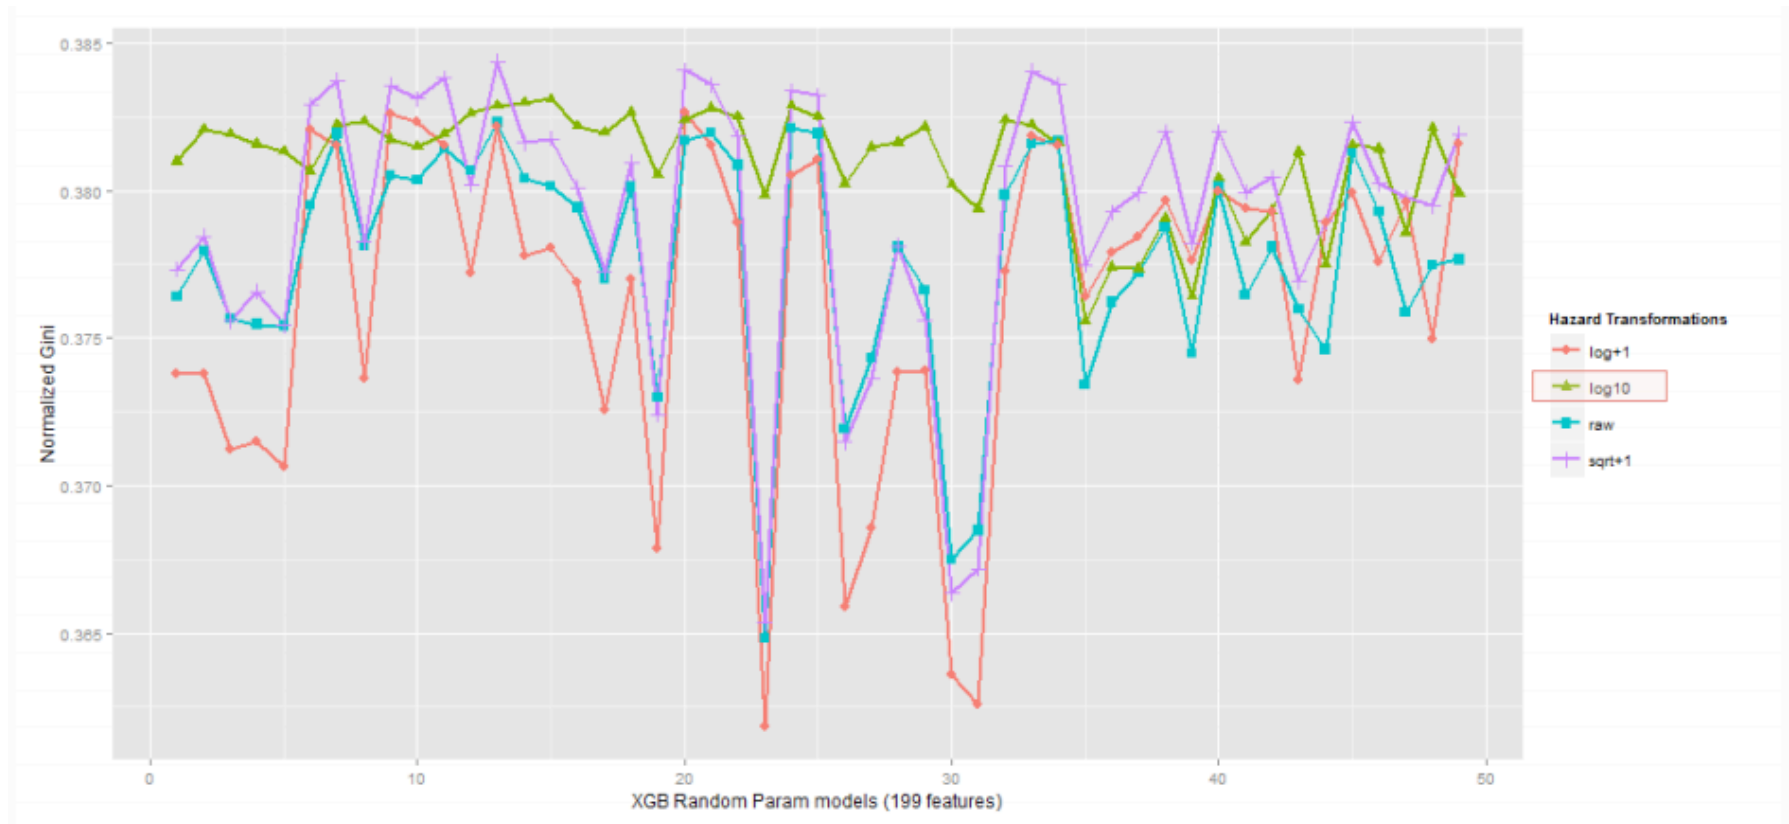
- Feature Scaling
  - Rescaling
    - Turn numeric features into the same scale (e.g., [-1,+1], [0,1], ...)
    - Python scikit-learn: MinMaxScaler
  - Standardization
    - Removing the mean ( $\mu = 0$ ) and scaling to unit variance ( $\sigma = 1$ )
    - Python scikit-learn: StandardScaler
  - To avoid features in greater numeric ranges dominating those in smaller numeric ranges
  - Critical for regularized linear models, KNN, SVM, K-Means, etc.
  - Can make a big difference to the performance of some models
  - Also speed up the training of gradient decent
  - But not necessary to do it for tree-based models

# Data Transformation

- Predictor/Response Variable Transformation
  - Use it when variable shows a skewed distribution make the residuals more close to “normal distribution” (bell curve)
  - Can improve model fit
    - $\log(x)$ ,  $\log(x+1)$ ,  $\sqrt{x}$ ,  $\sqrt{x+1}$ , etc.

# Variable Transformation

In Liberty Mutual Group: Property Inspection Prediction



Different transformations might lead to different results

# Feature Encoding

- Turn categorical features into numeric features to provide more fine-grained information
  - Help explicitly capture non-linear relationships and interactions between the values of features
  - Some machine learning tools only accept numbers as their input
    - xgboost, gbm, glmnet, libsvm, liblinear, etc.

# Feature Encoding

- Labeled Encoding
  - Interpret the categories as ordered integers (mostly wrong)
  - Python scikit-learn: LabelEncoder
  - Ok for tree-based methods
- One Hot Encoding
  - Transform categories into individual binary (0 or 1) features
  - Python scikit-learn: DictVectorizer, OneHotEncoder
  - Ok for K-means, Linear, NNs, etc.

# Feature Encoding

- Target mean encoding:
  - Instead of dummy encoding categorical variables and increasing the number of features we can encode each level as the mean of the response.
  - To avoid overfitting it is better to calculate weighted average of the overall mean of the training set and the mean of the level:

$$\lambda(n) * \text{mean}(\text{level}) + (1 - \lambda(n)) * \text{mean}(\text{dataset})$$

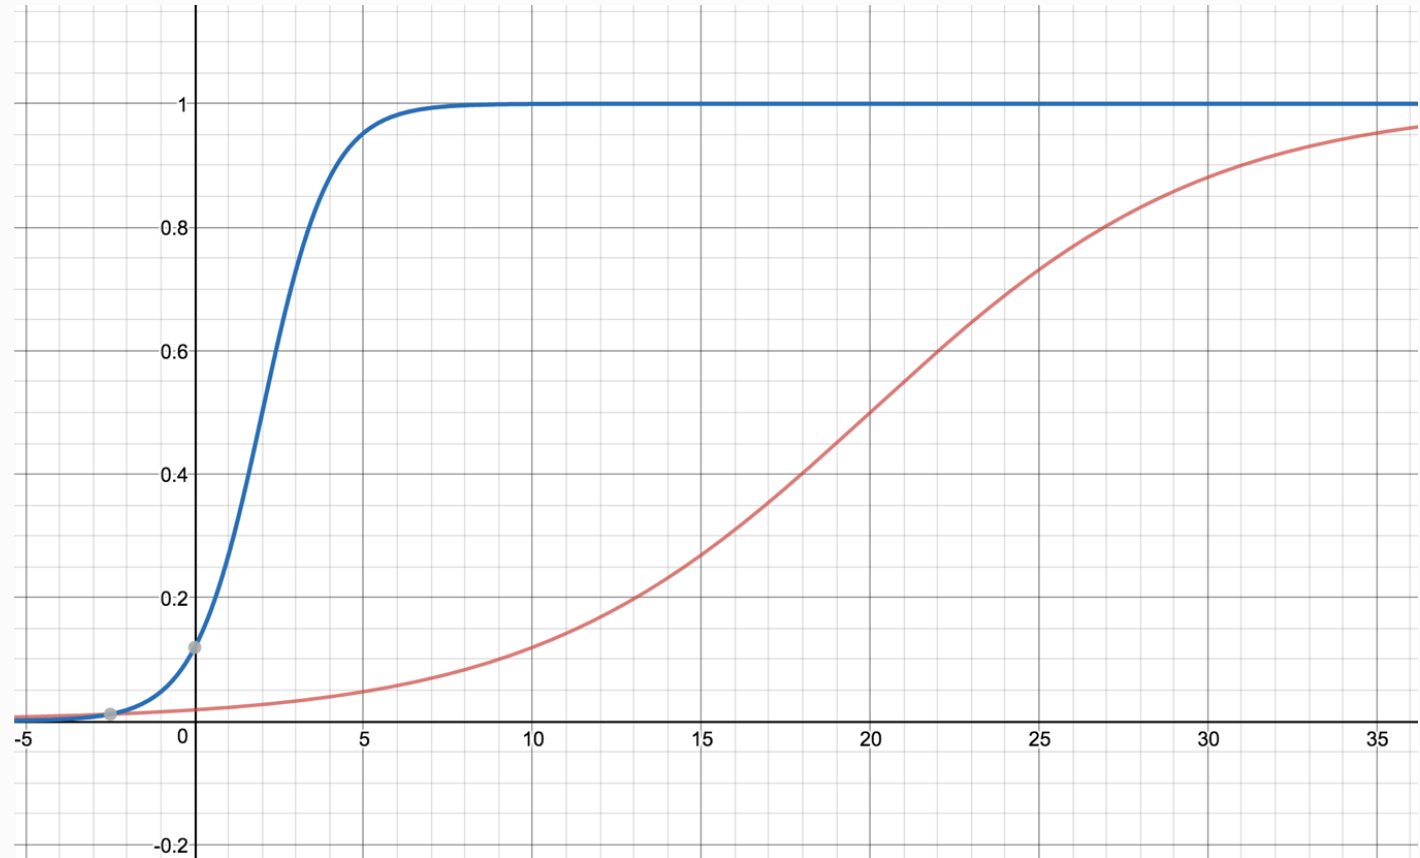
- The weights are based on the frequency of the levels i.e. if a category only appears a few times in the dataset then its encoded value will be close to the overall mean instead of the mean of that level.



# Feature Encoding – target mean encoding cont'd

$$\frac{1}{1 + \exp\left(\frac{-(x-k)}{f}\right)}$$

x = frequency  
k = inflection point  
f = steepness



# Feature Extraction: HTML Data

- HTML files are often used as the data source for classification problems
  - For instance, to identify whether a given html is an AD or not
- Possible features inside
  - html attributes (id, class, href, src, ....)
  - tag names
  - inner content
  - javascript function calls, variables, ....
  - script comment text
- Parsing Tools
  - BeautifulSoup (Python), etc

# Feature Extraction: Textual Data

- Bag-of-Words: extract tokens from text and use their occurrences (or TF/IDF weights) as features
- Require some NLP techniques to aggregate token counts more precisely
  - Split token into sub-tokens by delimiters or case changes
  - N-grams at word (often 2-5 grams) or character level
  - Stemming for English words
  - Remove stop words (not always necessary)
  - Convert all words to lower case
- Bag-of-Words Tools
  - Python: CountVectorizer, TfidfTransformer in scikit-learn package

# Feature Extraction: Textual Data

- Deep Learning for textual data
  - Turn each token into a vector of predefined size
  - Help compute “semantic distance” between tokens/words
    - For example, the semantic distance between user query and product titles in search results (how relevant?)
  - Greatly reduce the number of text features used for training
    - Use average vector of all words in given text
    - Vector size: 100~300 is often enough
  - Tools
    - Word2vec, Doc2vec, GloVe

# Feature Extraction

- There usually have some meaningful features inside existing features, you need to extract them manually
- Again you can use counts as features
- Some examples
  - Location
    - Address, city, state and zip code .... (categorical or numeric)
  - Time
    - Year, month, day, hour, minute, time ranges, .... (numeric)
    - Weekdays or weekend (binary)
    - Morning, noon, afternoon, evening, ... (categorical)
  - Numbers
    - Turn age numbers into ranges (ordinal or categorical)

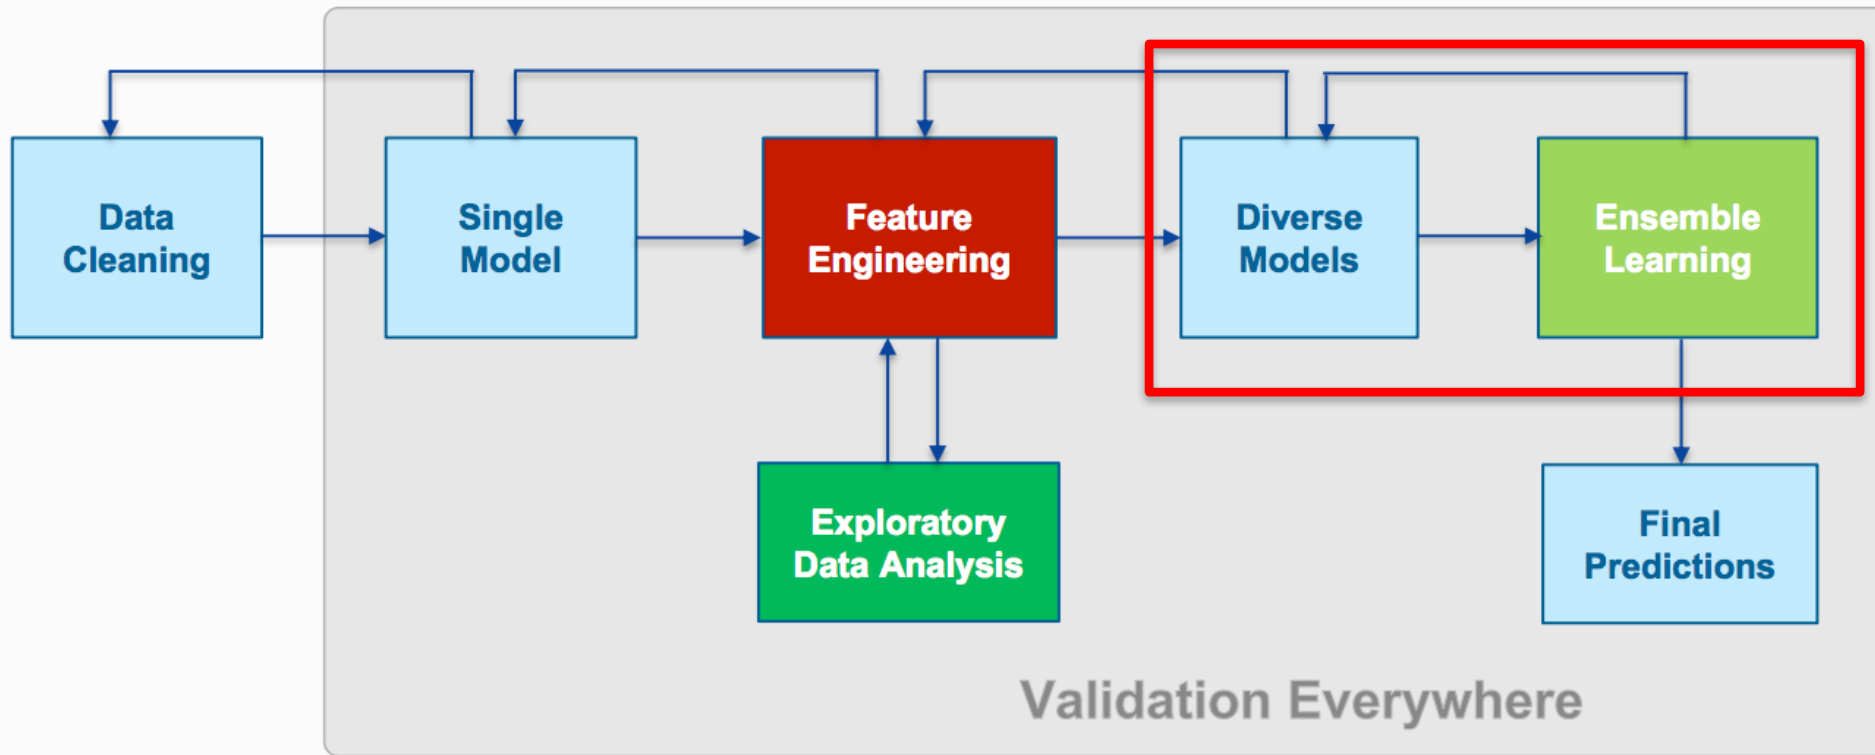
# Feature Selection

- Reduce the number of features by removing redundant, irrelevant or noisy features
- Feature Explosion after Feature Extraction!
  - More than 100,000 unique token features from textual data is common
  - Hard to put all of them in memory!
  - PCA or truncated SVD can help to select top-N informative features
    - With the risk of ignoring non-linear relationships between features and dropping important features
    - Use them only if there is no other choice

# Feature Selection

Type	Name	Python
Feature Importance Ranking	Gini Impurity	sklearn.ensemble.RandomForestClassifier sklearn.ensemble.RandomForestRegressor sklearn.ensemble.GradientBoostingClassifier sklearn.ensemble.GradientBoostingRegressor
	Chi-square	sklearn.feature_selection.chi2
	Correlation	scipy.stats.pearsonr scipy.stats.spearmanr
	Information Gain	sklearn.ensemble.RandomForestClassifier sklearn.ensemble.RandomForestRegressor sklearn.ensemble.GradientBoostingClassifier sklearn.ensemble.GradientBoostingRegressor xgboost
	L1-based Non-zero Coefficients	sklearn.linear_model.Lasso sklearn.linear_model.LogisticRegression
Feature Subset Selection	Recursive Feature Elimination (RFE)	sklearn.feature_selection.RFE
	Boruta Feature Selection	Boruta

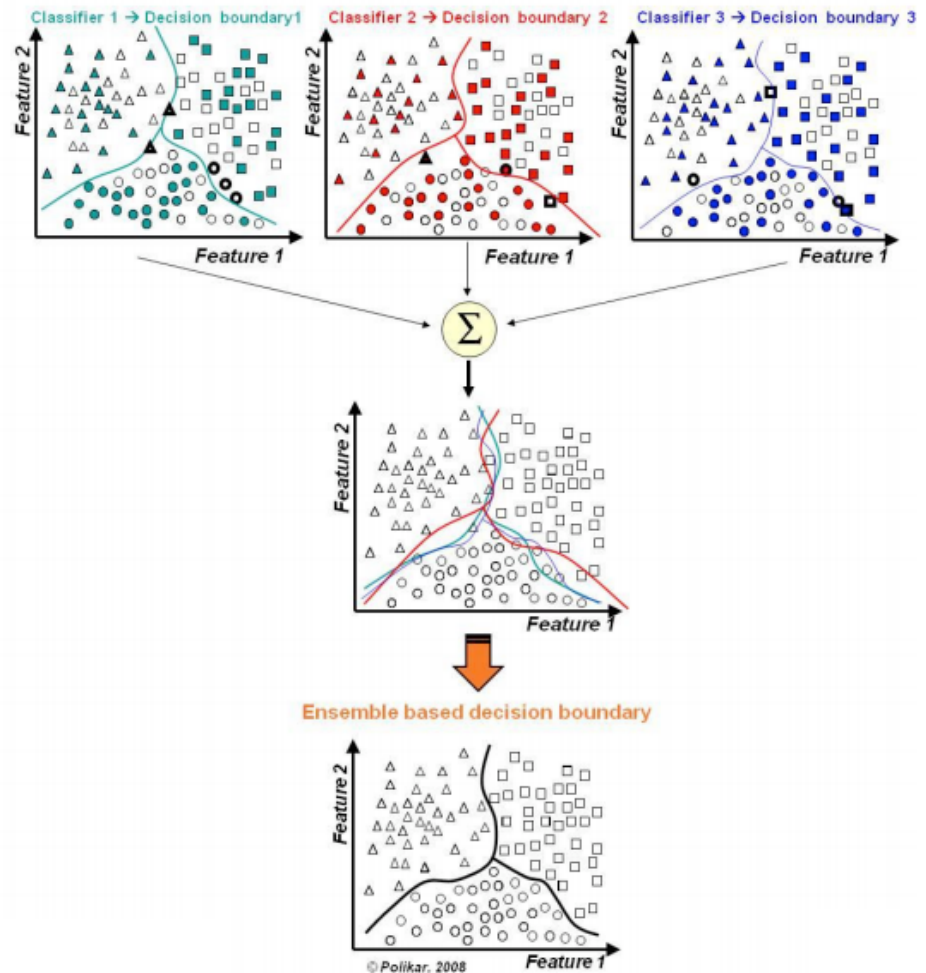
# Feature Engineering





# Ensemble Learning - Illustration

- Assume that diverse models see different aspects of the problem space and are wrong in different ways
- Simplest way: the average (possibly weighted) of model predictions
  - Less variance
  - Less generalization error
  - Less chance of overfitting
  - Better chance to win!



# Ensemble Learning - Blending

- Simplest way: Weighted average:

$$f(x) = \alpha * f_1(x) + \beta * f_2(x) + \gamma * f_3(x)$$

- Not so simple way:

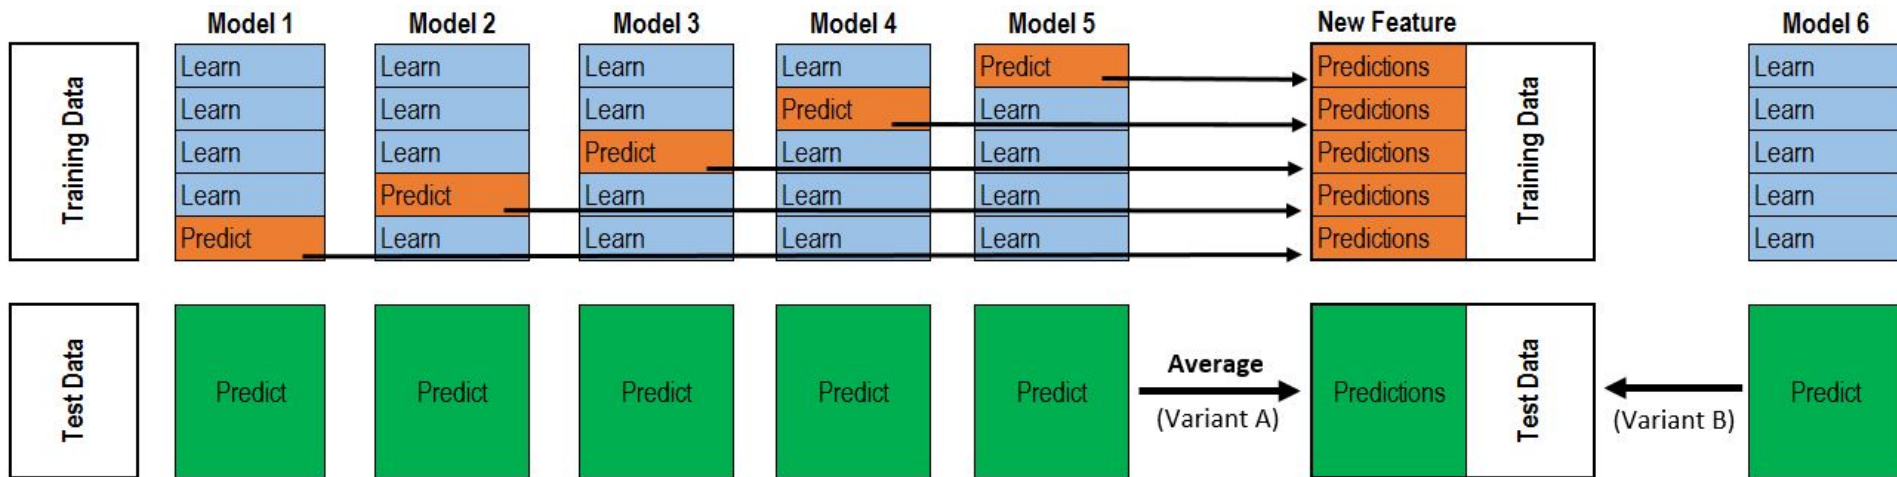
$$f(x) = \alpha_1 * f_1(x)^{\alpha_2} + \beta_1 * f_2(x)^{\beta_2} + \gamma_1 * f_3(x)^{\gamma_2}$$

- And everything in between
- How to find coefficients?
  - By using any optimization package:
    - hyperopt
    - `scipy.optimize.fmin`

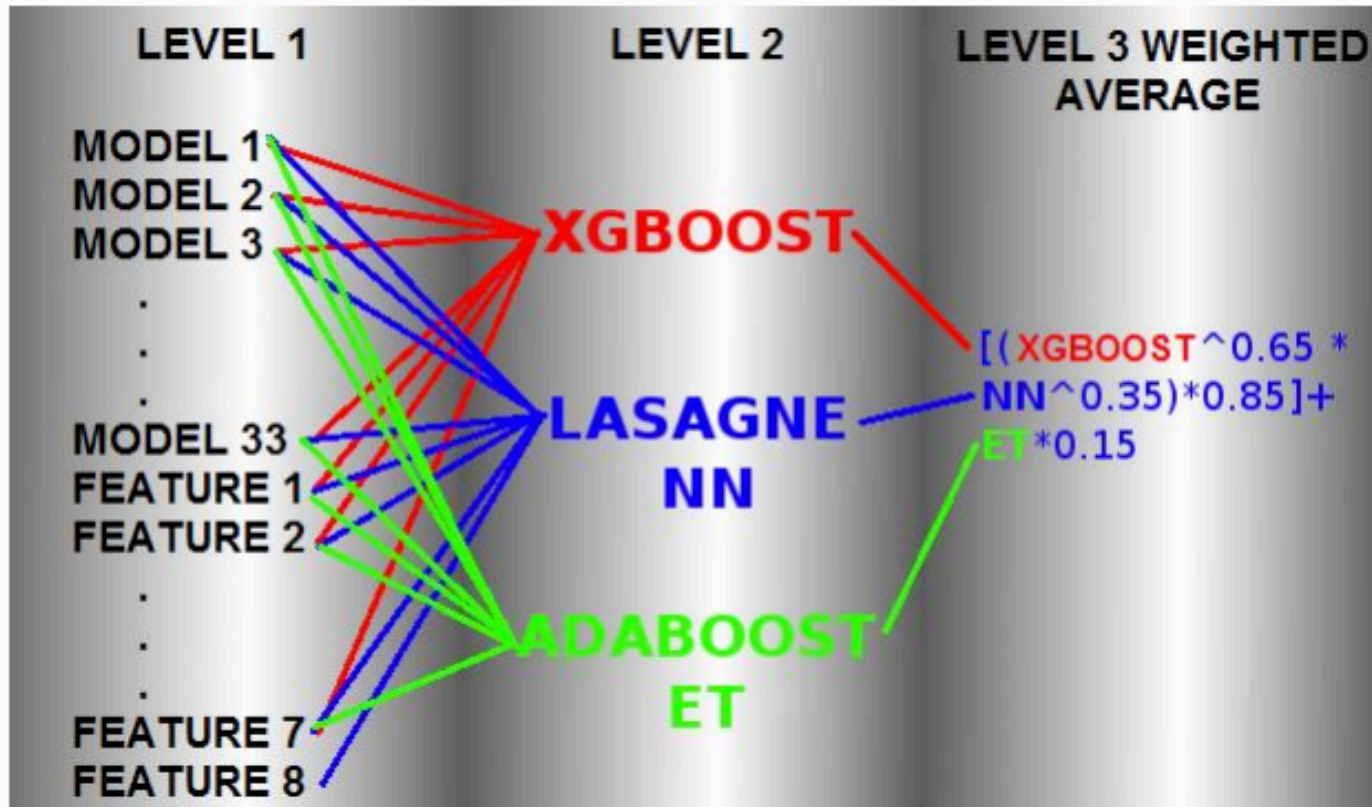
# Stacking

- Using the predictions of current level models as features for training next level models
- What we called “out-of-fold predictions” or “meta-features”
  - Usually we only build 2 levels of models, 4 levels at most
  - Can combine meta-features with original feature sets for training
  - Potential risk of overfitting if cross validation process is not correct
- Also called “Stacking”

# Stacking

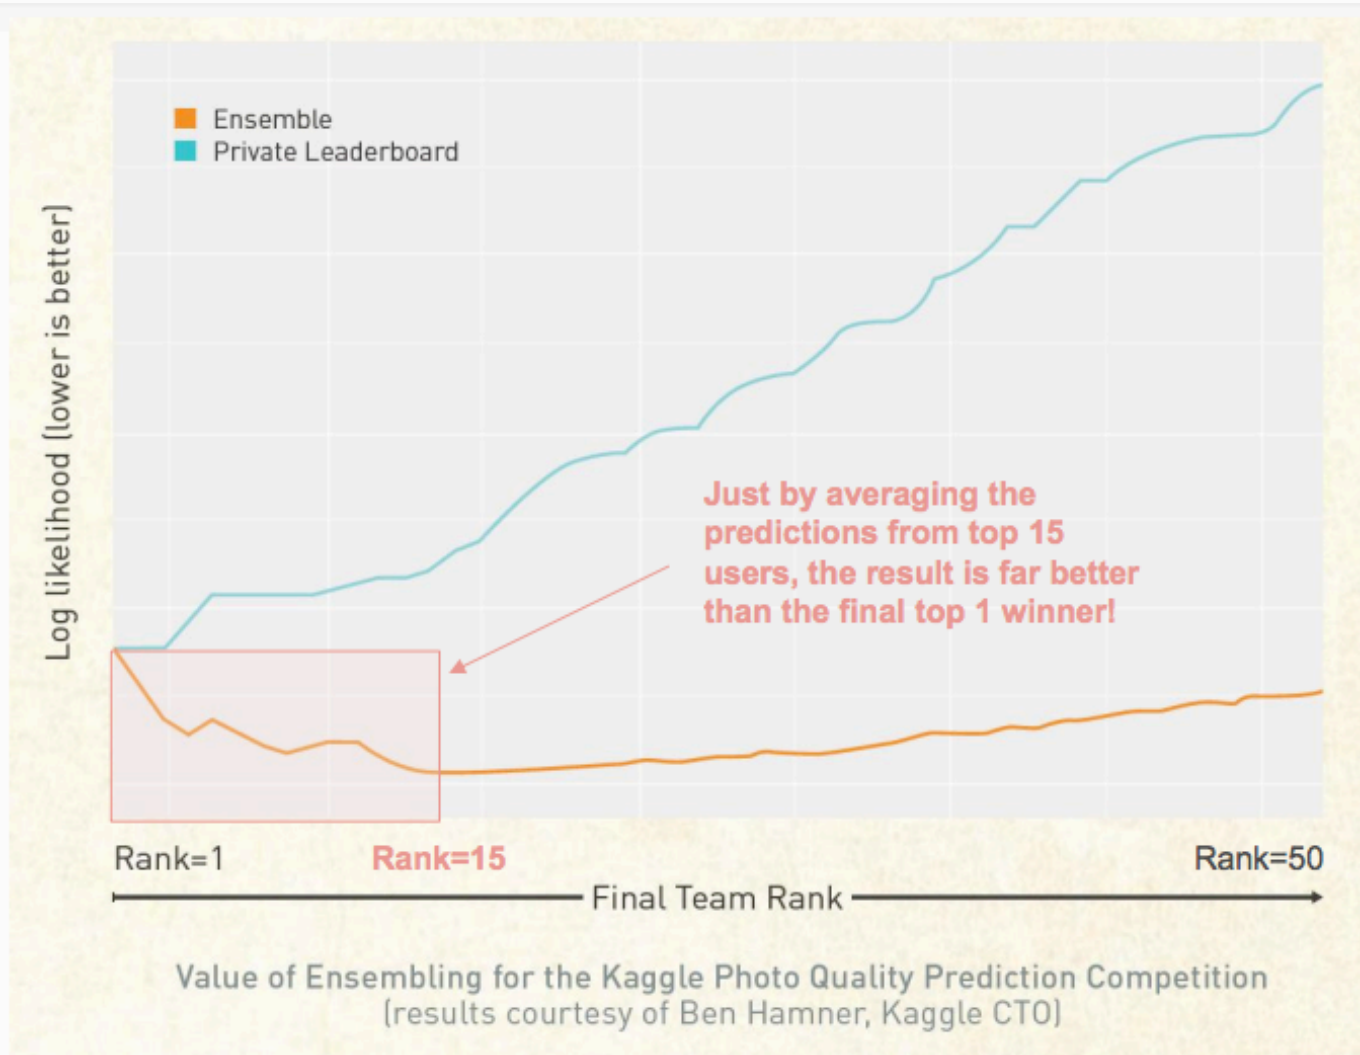


**Stacked Generalization**  
**Used by Gilberto Titericz Junior (Kaggle #1)**  
**in Otto Product Classification Challenge**



Reference: <https://www.kaggle.com/c/otto-group-product-classification-challenge/forums/t/14335/1st-place-winner-solution-gilberto-titericz-stanislav-semenov>

# The Power of Ensemble Learning



Reference: Booz Allen Hamilton. (2015). "Field Guide to Data Science," page 93.

# Ensemble Learning - Tips

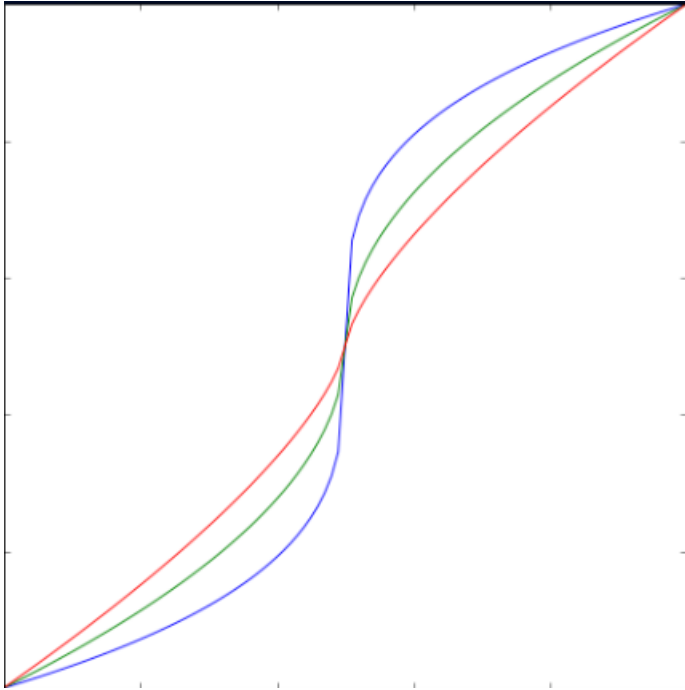
- Always start from simple blending first
- Even simple averaging of multiple models of the same type can help reduce variance
- Generally work well with less-correlated good models
  - Models with similar performance but their correlation lies in between 0.85~0.95 (that's ideal)
  - Gradient Boosting Machines usually blend well with Neural Nets and Linear models
- Must Read: [Kaggle Ensembling Guide](#) by MLWave

# Teaming strategy

- Usually we agree to team up at pretty early phase but keep working independently to avoid any bias in our solutions
- After teaming up:
  - Code sharing can help you to learn new things but a waste of time from competition perspective;
  - Share data – especially some engineered features;
  - Combine your models' outcomes using k-fold stacking and build a second level meta-learner (stacking)
  - Continue iteratively add new features and building new models
- Teaming up right before competition end:
  - “Black-box” ensembling – linear and not so linear blending using LB as validation set.
    - Linear:  $f(x) = \alpha * f_1(x) + \beta * f_2(x) + \gamma * f_3(x)$
    - Non linear:  $f(x) = \alpha_1 * f_1(x)^{\alpha_2} + \beta_1 * f_2(x)^{\beta_2} + \gamma_1 * f_3(x)^{\gamma_2}$



# Metrics tips



- Depending on the competition metric you can do some tricks as well.

- Logloss:

- $-\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \log(p_{ij})$ 
  - where N is the number of examples, M is the number of classes, and  $y_{ij}$  is a binary variable indicating whether class j was correct for example i.

- Calibrate result by using formula:

- $y_{new} = 0.5 * \left( (2 * \text{abs}(y - 0.5))^{beta} \right) * \text{sign}(y - 0.5) + 0.5$
- *beta* coefficient can be found by using LB feedback or CV

- Using geometric mean instead average

# Metrics tips

- MAE:

- $\frac{1}{N} \sum_{i=1}^N |y_i - y'_i|$

- where N is the number of examples

- Most of the ML algorithms optimize RMSE instead of MAE, so small adjustment can be useful:

- $y_{new} = \alpha * y + \beta$

- Also, using median instead of averaging (mean) might be beneficial

# Metrics tips

- RMSLE:

$$\epsilon = \sqrt{\frac{1}{n} \sum_{i=1}^n (\log(p_i + 1) - \log(a_i + 1))^2}$$

- where  $n$  is the number of examples

- Most of the ML algorithms optimize RMSE, so log-transforming of target variable will make it to optimize RMSLE instead:

- $y_{new} = \log(y + 1)$

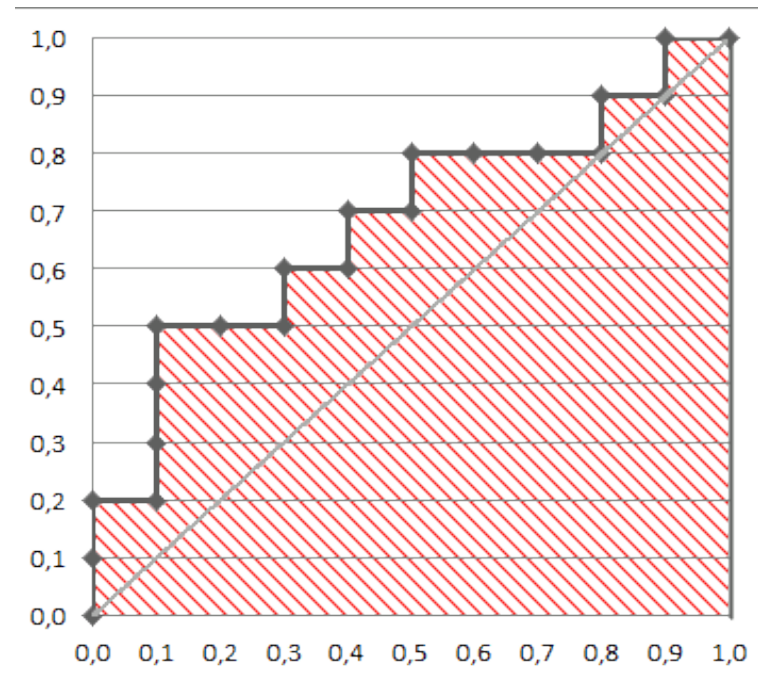
- Don't forget to transform it back:

- $y_{old} = \exp(y_{new}) - 1$

# Metrics tips

- AUC: Area Under ROC Curve, total area is 100%:

- so AUC = 1 is for a perfect classifier for which all positive come after all negatives
- AUC = 0.5 - randomly ordered
- AUC = 0 - all negative come before all positive
- so  $AUC \in [0,1] \in [0,1]$
- typically we don't have classifiers with  $AUC < 0.5$



- AUC is sensitive to order (ranking), not to specific values:

- $y_{new} = \frac{rank(y)}{N}$ ; N – number of rows
- by this transformation we can blend in ensemble any classifiers, no matter how different their distributions were

# Summary

- Know your tools
- Know your data
- Know your metric
- Right validation schema is a key
- Never stick to favorite models only, try all of them
  - Always include linear models into ensemble
  - ... and Neural Nets
  - ... and kNN
- Good features beat ensemble
  - My team got 10<sup>th</sup> place on BNP Paribas Competition with ensemble of 120 models
  - Branden Murray's single model would be on 8<sup>th</sup> (his team placed 4<sup>th</sup>)

**Thank you!**

Q & A