NLP with H_2O

Agenda

- Our Use Case
- H2O Overview
- Natural Language Processing
- Demo

Our Use Case

The Data

The Amazon Fine Food Reviews dataset consists of 568,454 food reviews Amazon users left up to October 2012

- J. McAuley and J. Leskovec. From amateurs to connoisseurs: modeling the evolution of user expertise through online reviews

Column	Example
Product ID	B006K2ZZ7K
User ID	A1UQRSCLF8GW1T
Helpfulness Numerator	1
Helpfulness Denominator	1
Score	5
Time	1350777600
Summary	Great taffy
Text	*"Great taffy at a great price. There was a wide assortment of yummy taffy. Delivery was very quick. If your a taffy lover, this is a deal."*

Goal

- Predict whether a food product has a good rating based on the reviews

“Great taffy at a great price. There was a wide assortment of yummy taffy. Delivery was very quick. If your a taffy lover, this is a deal.”

H2O Overview

What is H2O?

Math Platform

Open source in-memory AI engine

- Parallelized and distributed algorithms
- GLM, Random Forest, GBM, Deep Learning, etc.

Tech and API

Easy to use and adopt

- Written in Java – perfect for Java Programmers
- Install is lightweight
- REST API (Java) – run H2O from R, Python, WebUI, Excel, Tableau, Tibco

Big Data

More data? Or better models? BOTH

- Use all of your data – model without sampling
- More Data + Better Models = Better Predictions

H2O Cluster

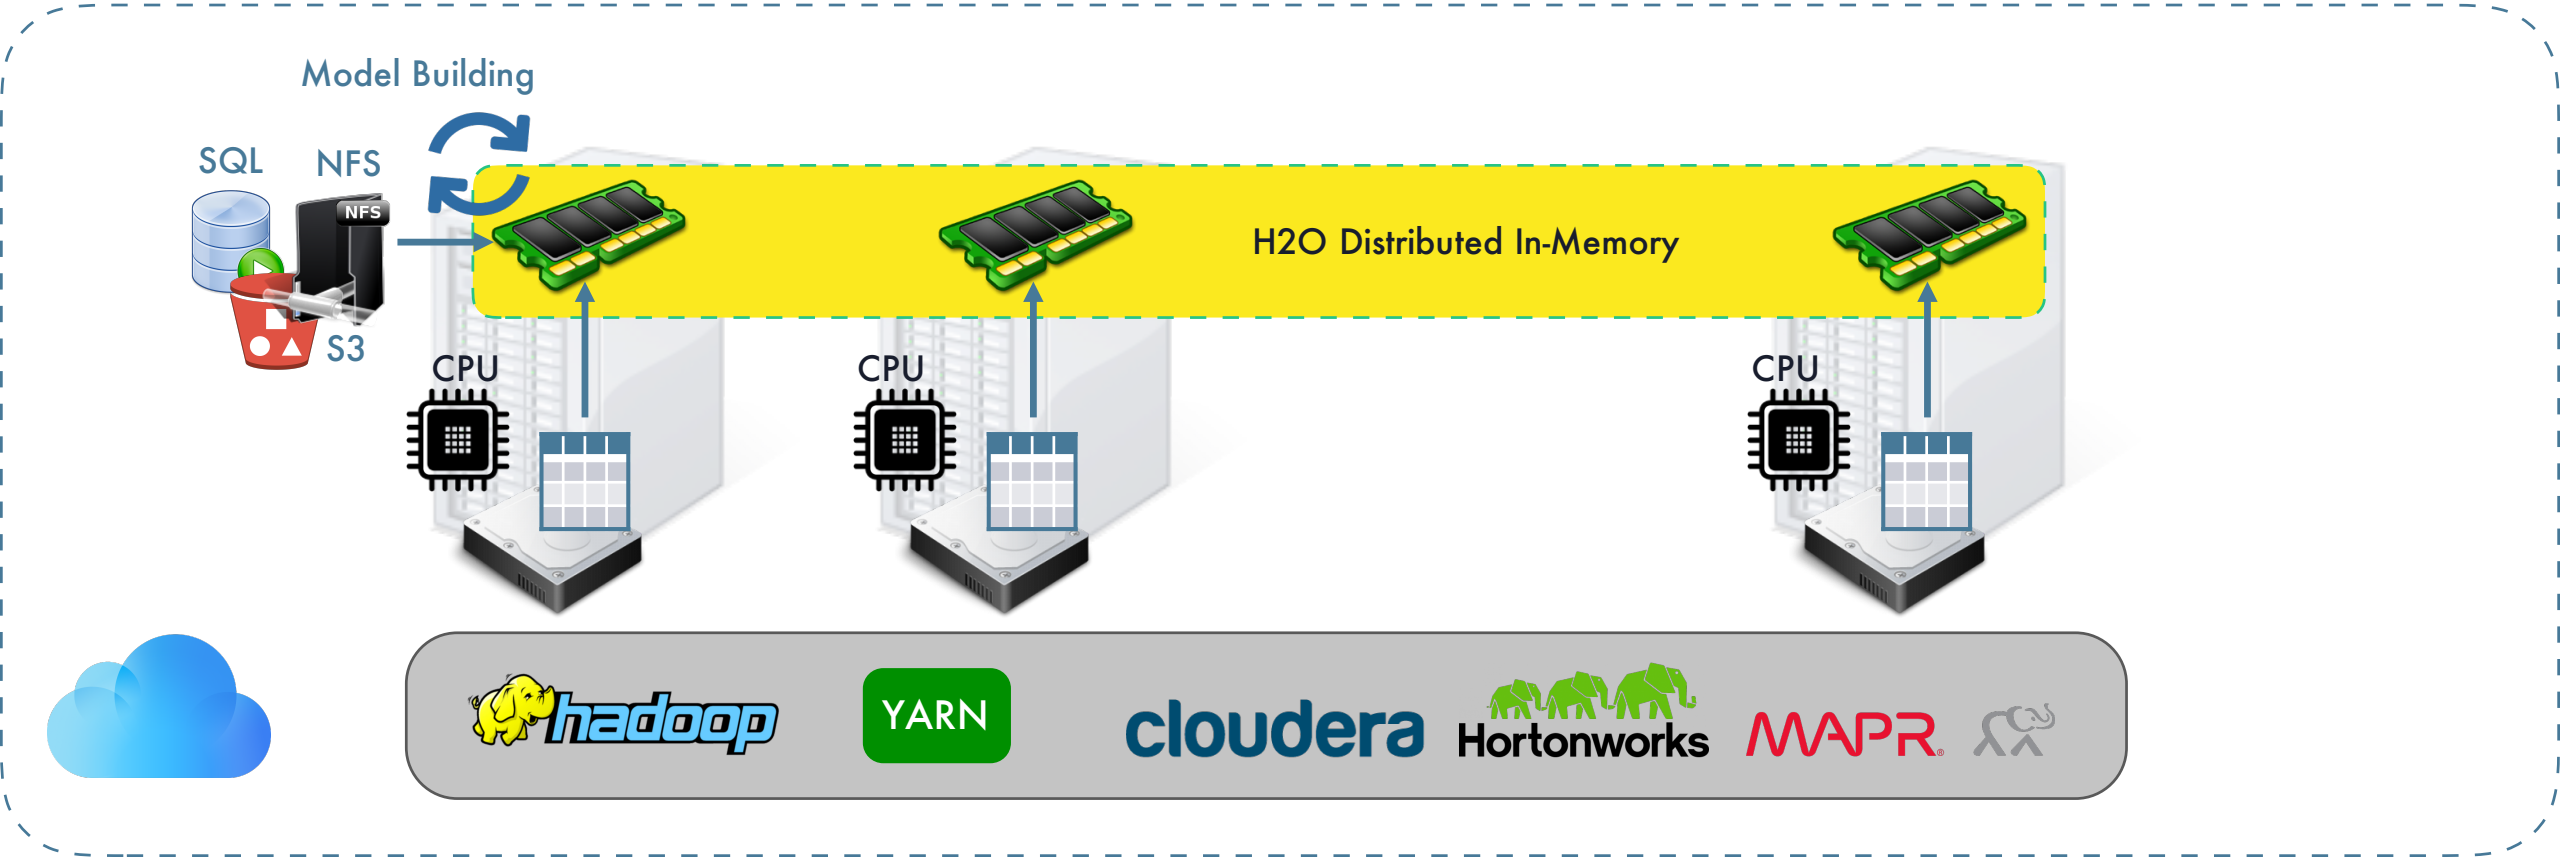
Model Building
SQL
NFS
NFS
S3
H2O Distributed In-Memory
CPU
CPU
CPU
hadoop
YARN
cloudera
Hortonworks
MAPR

H2O Cluster

H2O Clients

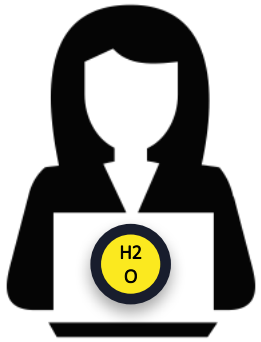

H2O Cluster

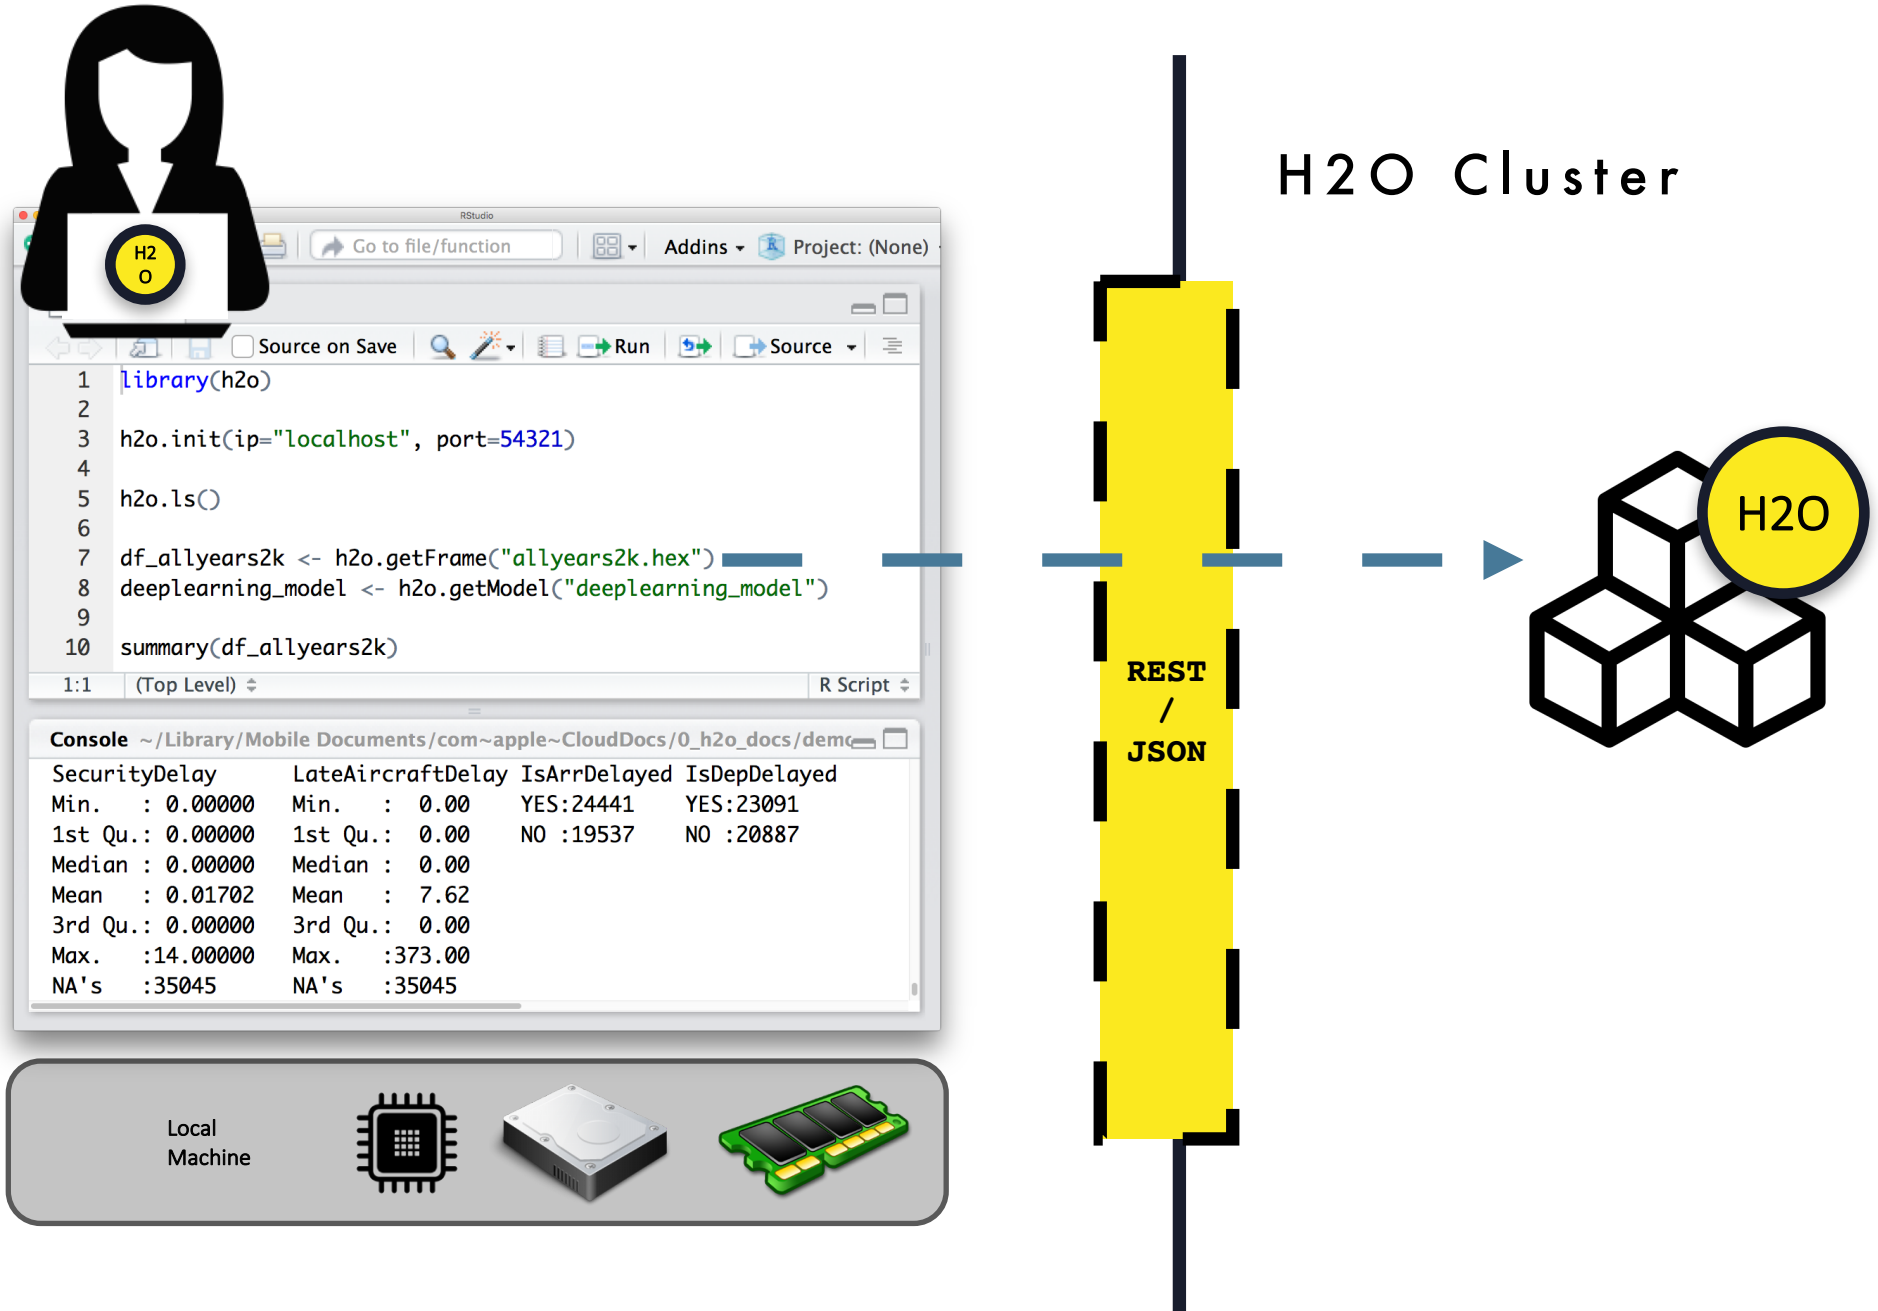
H2O Cluster
H2O
RStudio
Go to file/function
Addins
Project: (None)
Source on Save
Run
Source
1 library(h2o)
2
3 h2o.init(ip="localhost", port=54321)
4
5 h2o.ls()
6
7 df_allyears2k <- h2o.getFrame("allyears2k.hex")
8 deeplearning_model <- h2o.getModel("deeplearning_model")
9
10 summary(df_allyears2k)
1:1 (Top Level)
R Script
Console ~/Library/Mobile Documents/com~apple~CloudDocs/0_h2o_docs/dem
SecurityDelay LateAircraftDelay IsArrDelayed IsDepDelayed
Min. : 0.00000 Min. : 0.00 YES:24441 YES:23091
1st Qu.: 0.00000 1st Qu.: 0.00 NO :19537 NO :20887
Median : 0.00000 Median : 0.00
Mean : 0.01702 Mean : 7.62
3rd Qu.: 0.00000 3rd Qu.: 0.00
Max. :14.00000 Max. :373.00
NA's :35045 NA's :35045
Local
Machine
REST
/
JSON
H2O

Natural Language Processing

What is NLP?

- A way for computers to understand text and language

- Used For:
  - Sentiment Analysis
  - Topic Identification
  - **Improving Supervised Learning Models**

Methods for Natural Language Processing

Bag of Words

- does the word exist in the document?

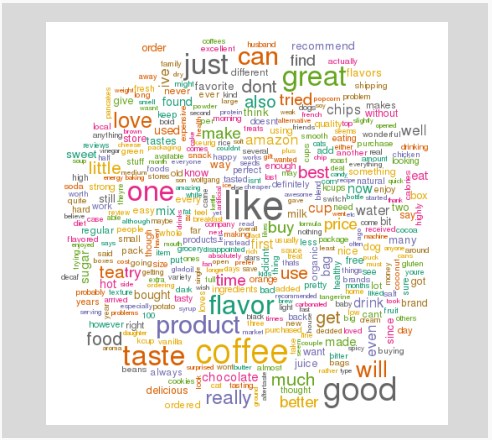

Count Based

- how often do words occur in each document?
- TF-IDF

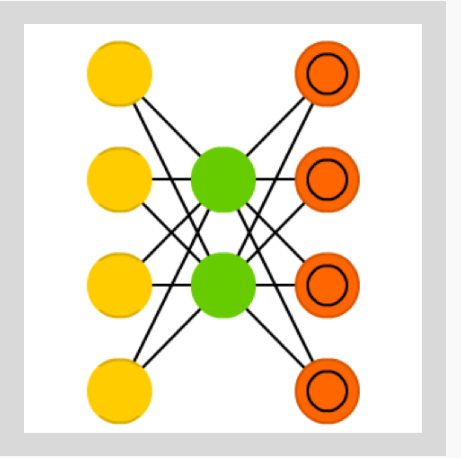

Predictive

- train models predicting a word or sentence from its neighbors
- Word2Vec, RNN

H2O.ai

Word Embeddings

- What?
  - Mapping of words to vectors from a high dimensional space (100 – 1000)

- Why?
  - Embeddings capture the meaning of the word
  - Semantically similar words are close to each other

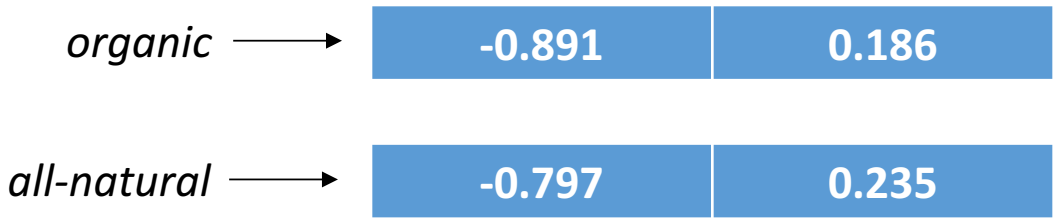

Word Embeddings

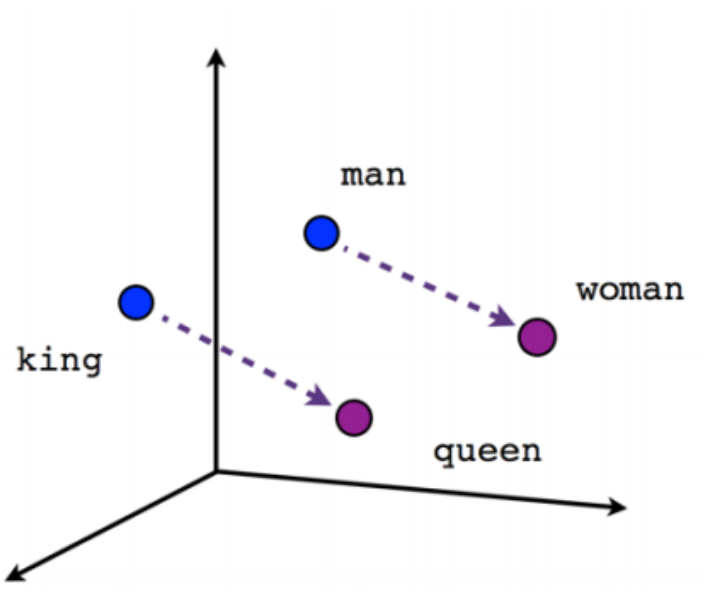

Male-Female

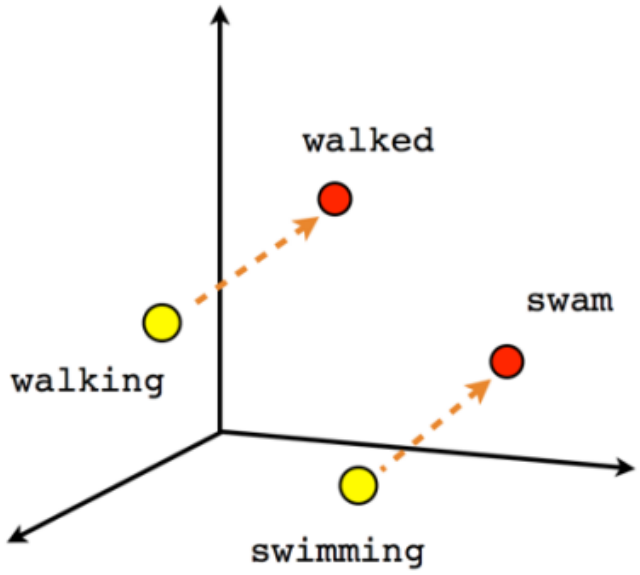

Verb tense

https://www.tensorflow.org/tutorials/word2vec

Word2Vec Algorithm

How do we use a neural network to capture the semantic meaning of words?

- Frame the problem as a supervised learning problem
  - Given an input word predict the neighboring words

*“It's even better than the **organic, all-natural brands** I have tried.”*

Given: *all-natural*

Predict: *organic, brands*

Word2Vec Algorithm

*"It's even better than the **organic, all-natural brands** I have tried."*

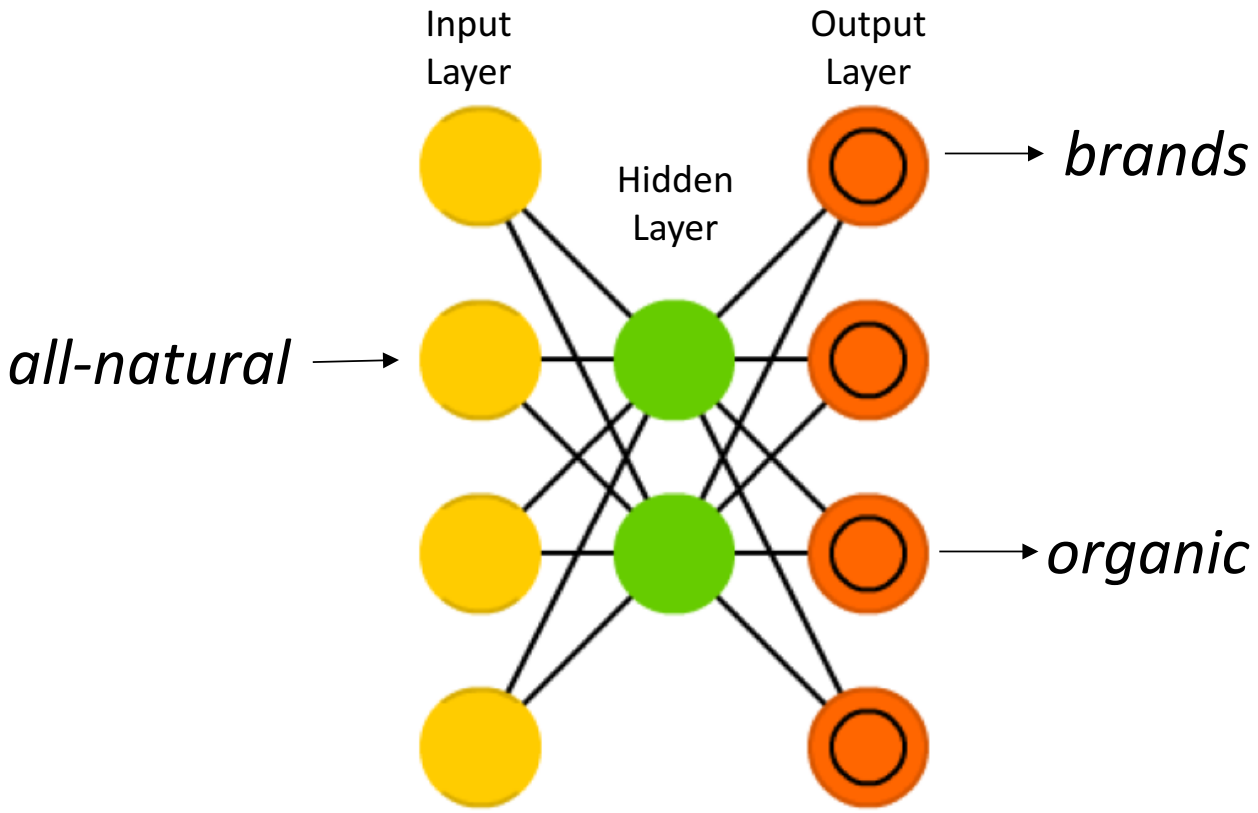

Word2Vec Algorithm

Matrix from Hidden Layer = Word Embeddings

Word	C1	C2
brands	0.647	0.235
all-natural	-0.797	0.235
organic	-0.891	0.186
tried	-0.751	0.409

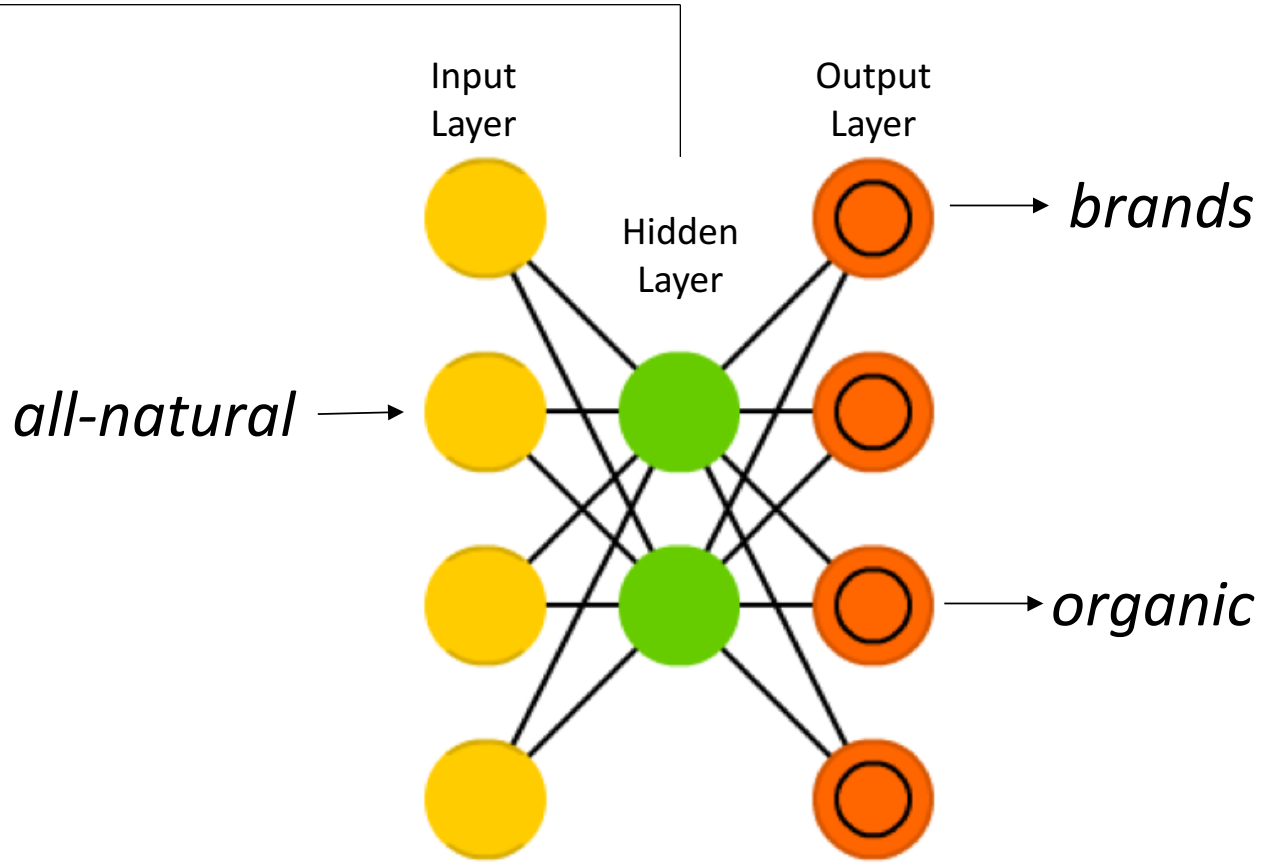

Word Embeddings

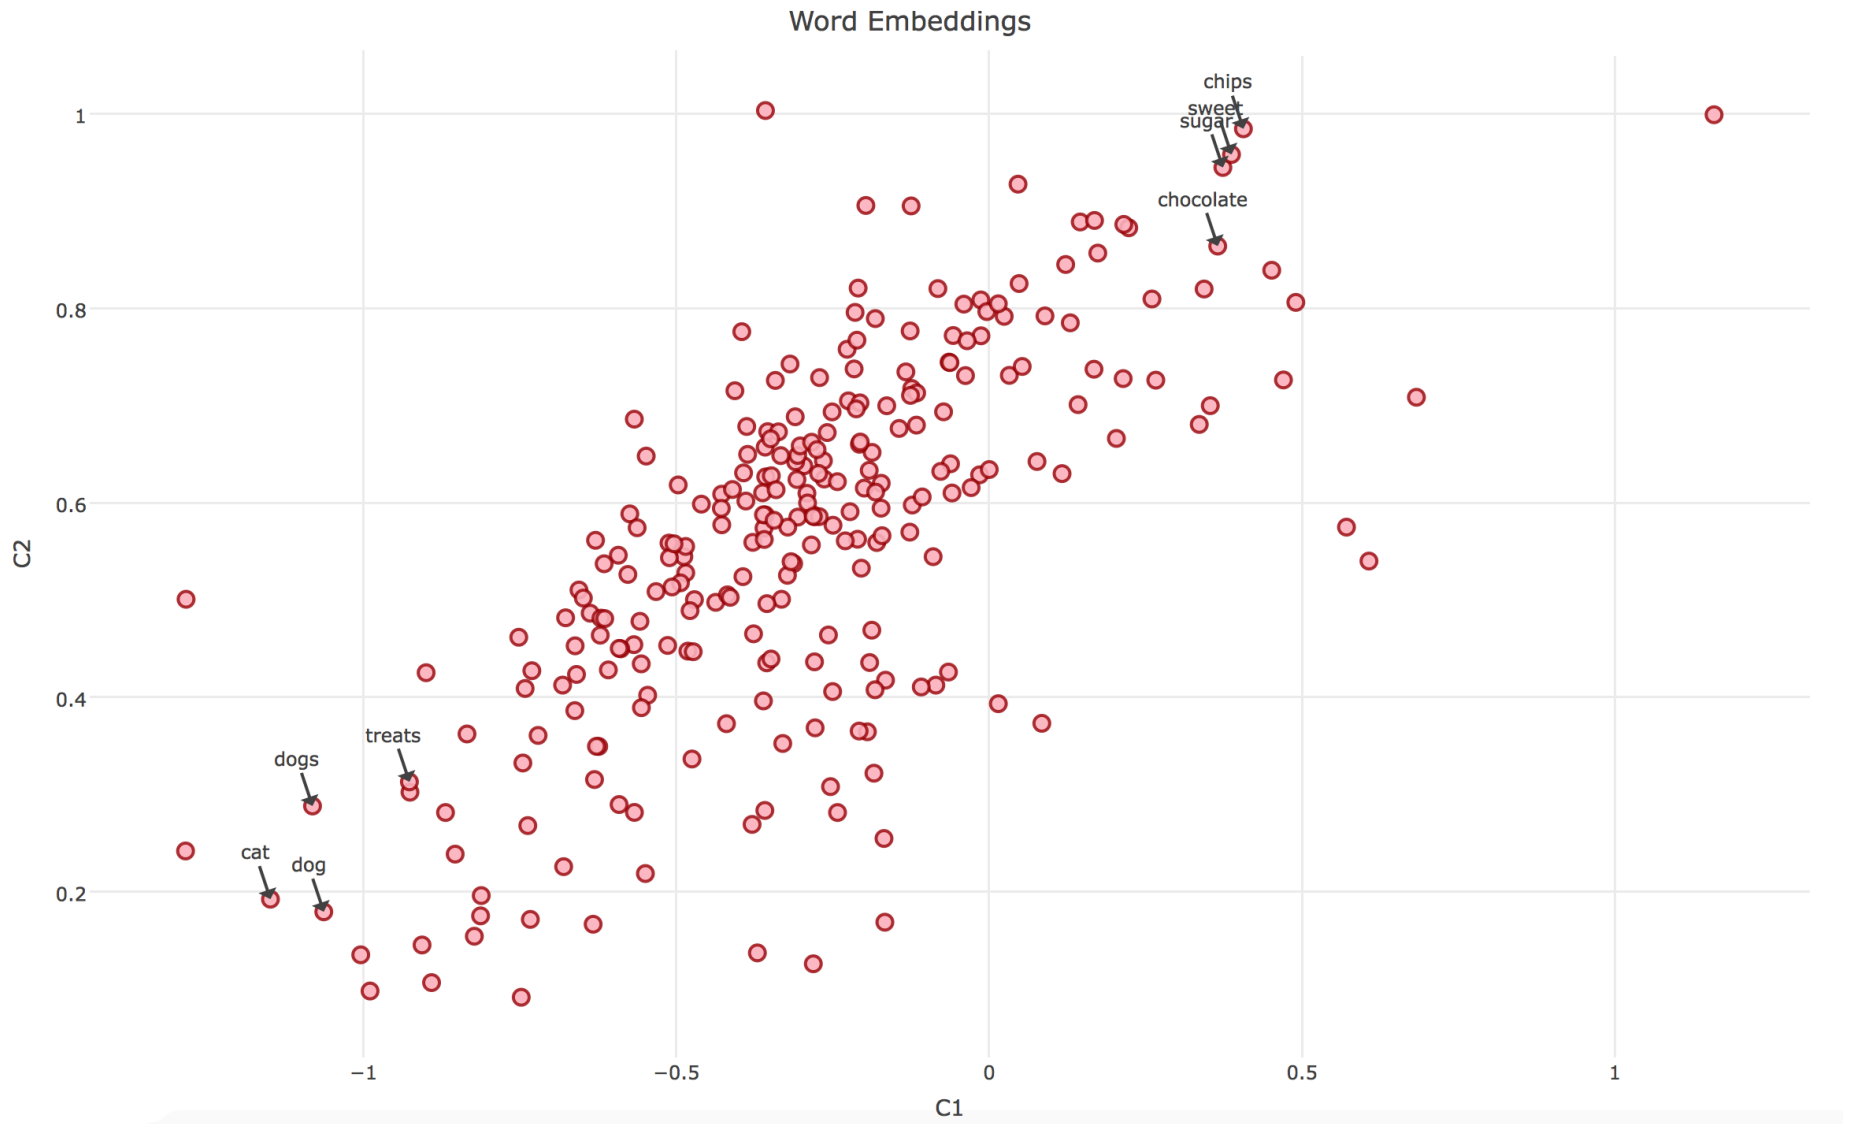

Word2Vec Usage

- Word2Vec Embeddings are typically used as a pre-processing step to a supervised learning algorithm

*"It's even better than the **organic, all-natural brands** I have tried."*

Word	C1	C2
even	0.602	-0.499
better	0.765	-0.437
brands	0.647	0.235
all-natural	-0.797	0.235
organic	-0.891	0.186
tried	-0.751	0.409

Aggregate

C1	C2
-0.071	0.0215

Rule Based Model

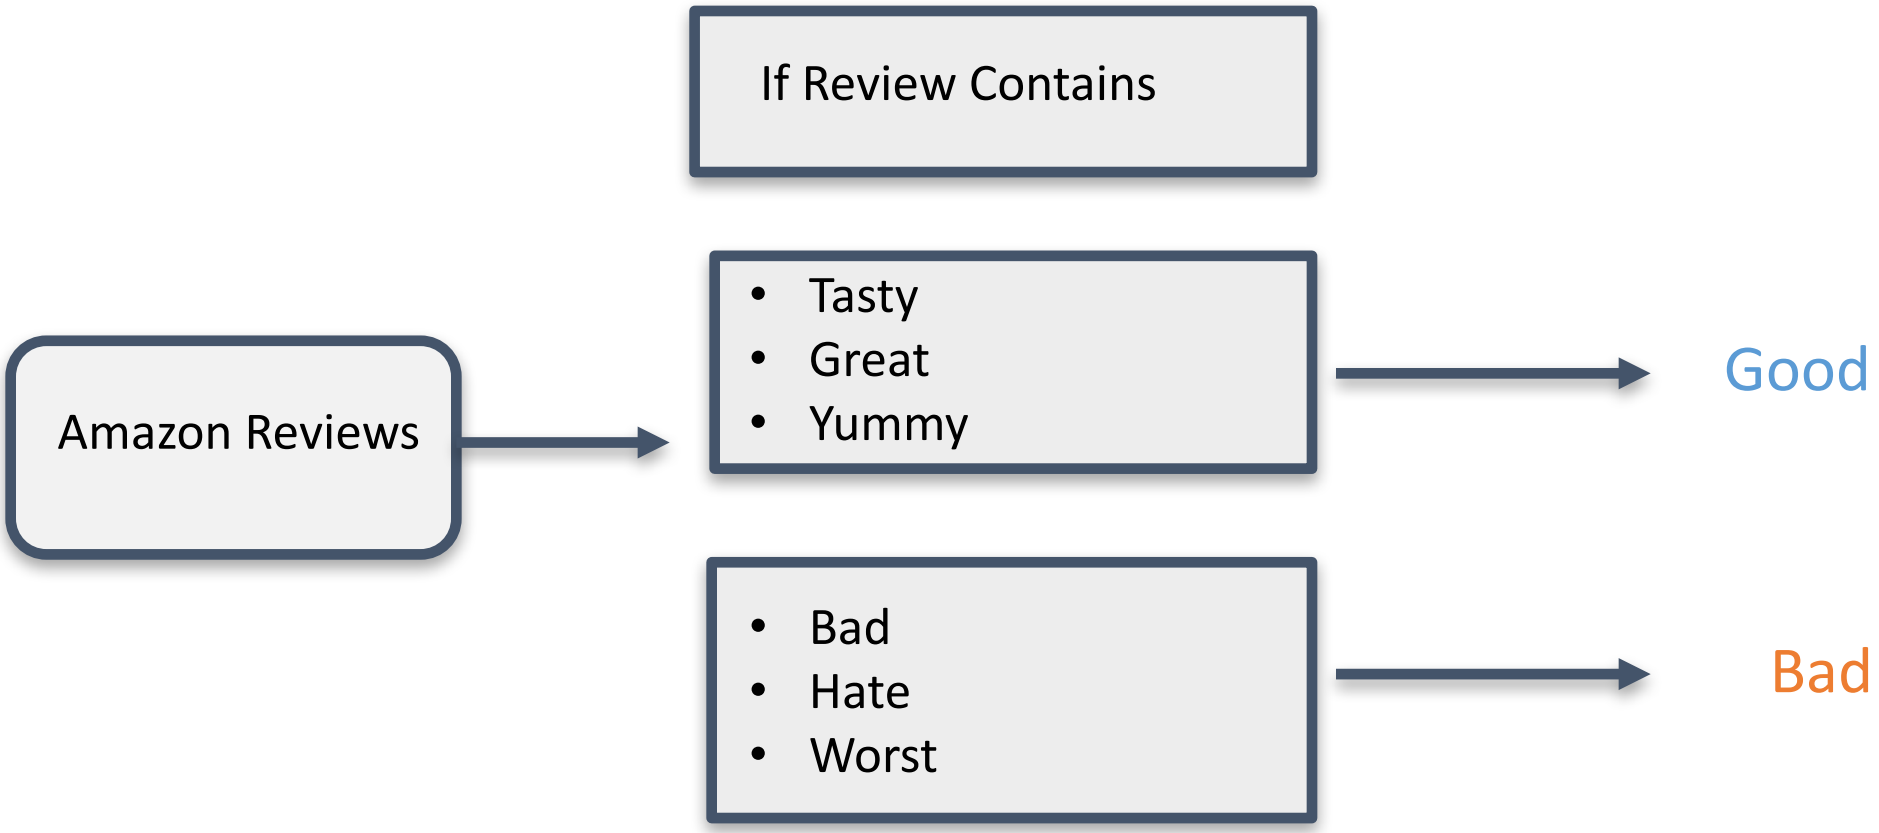

Machine Learning Model

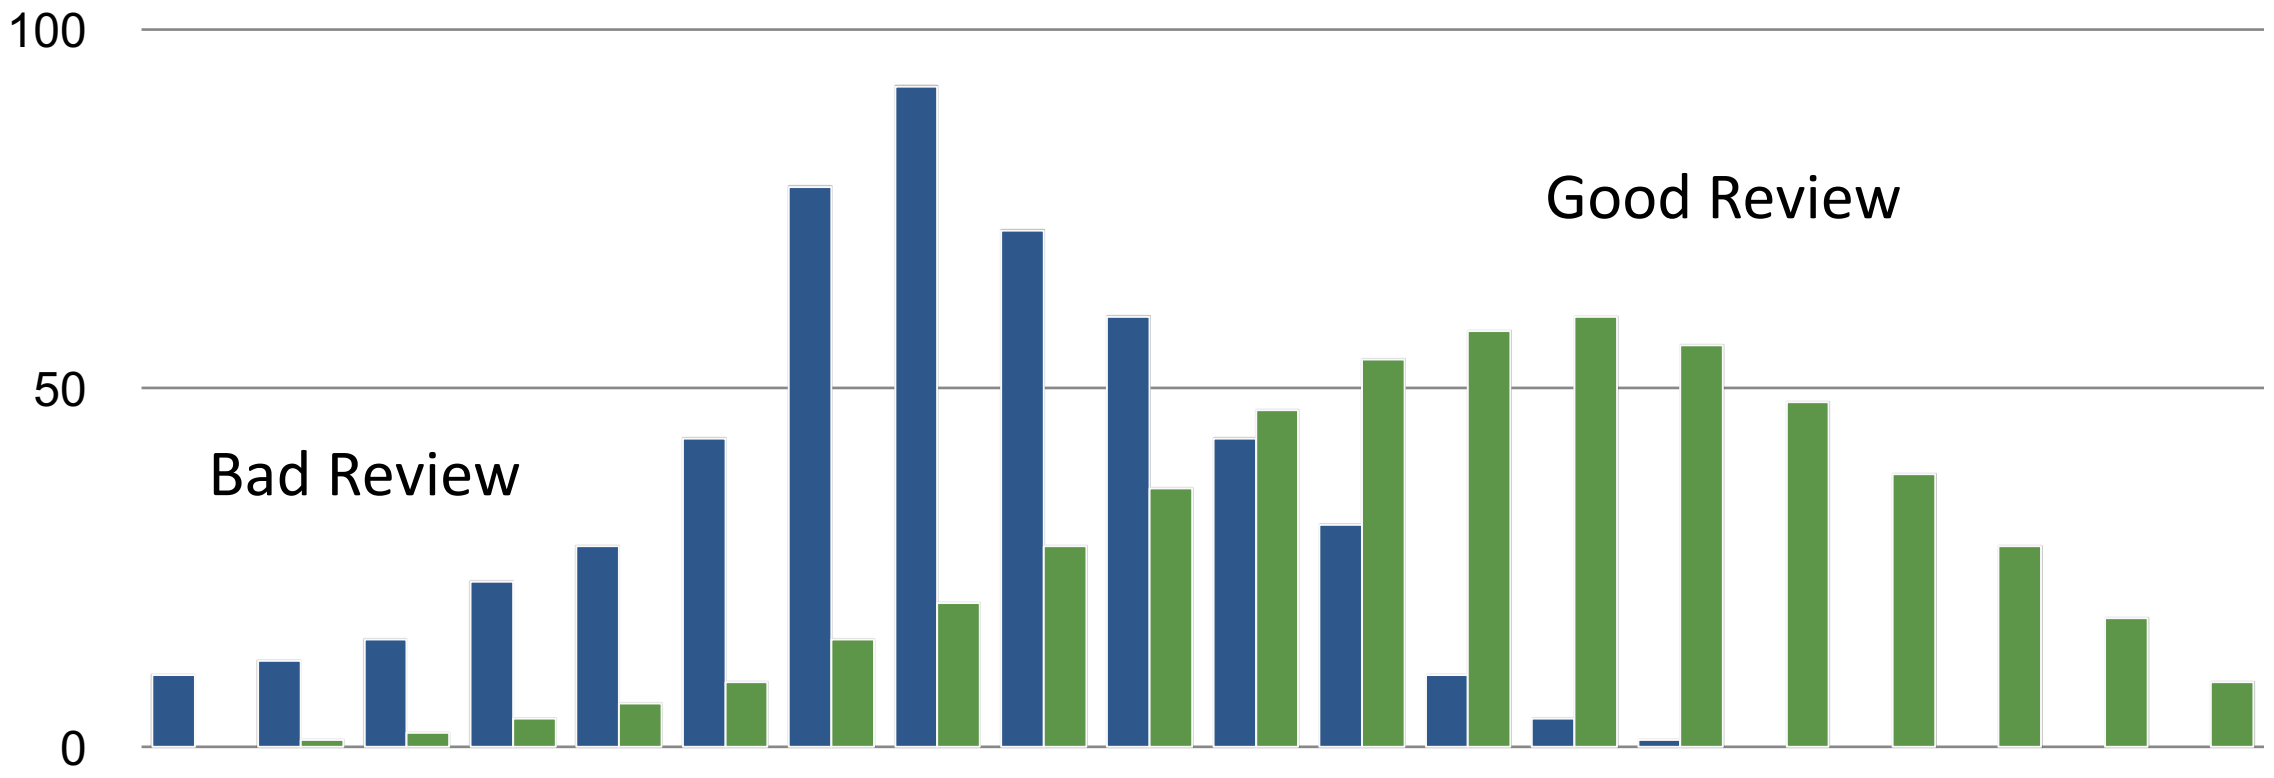

Machine Learning Model

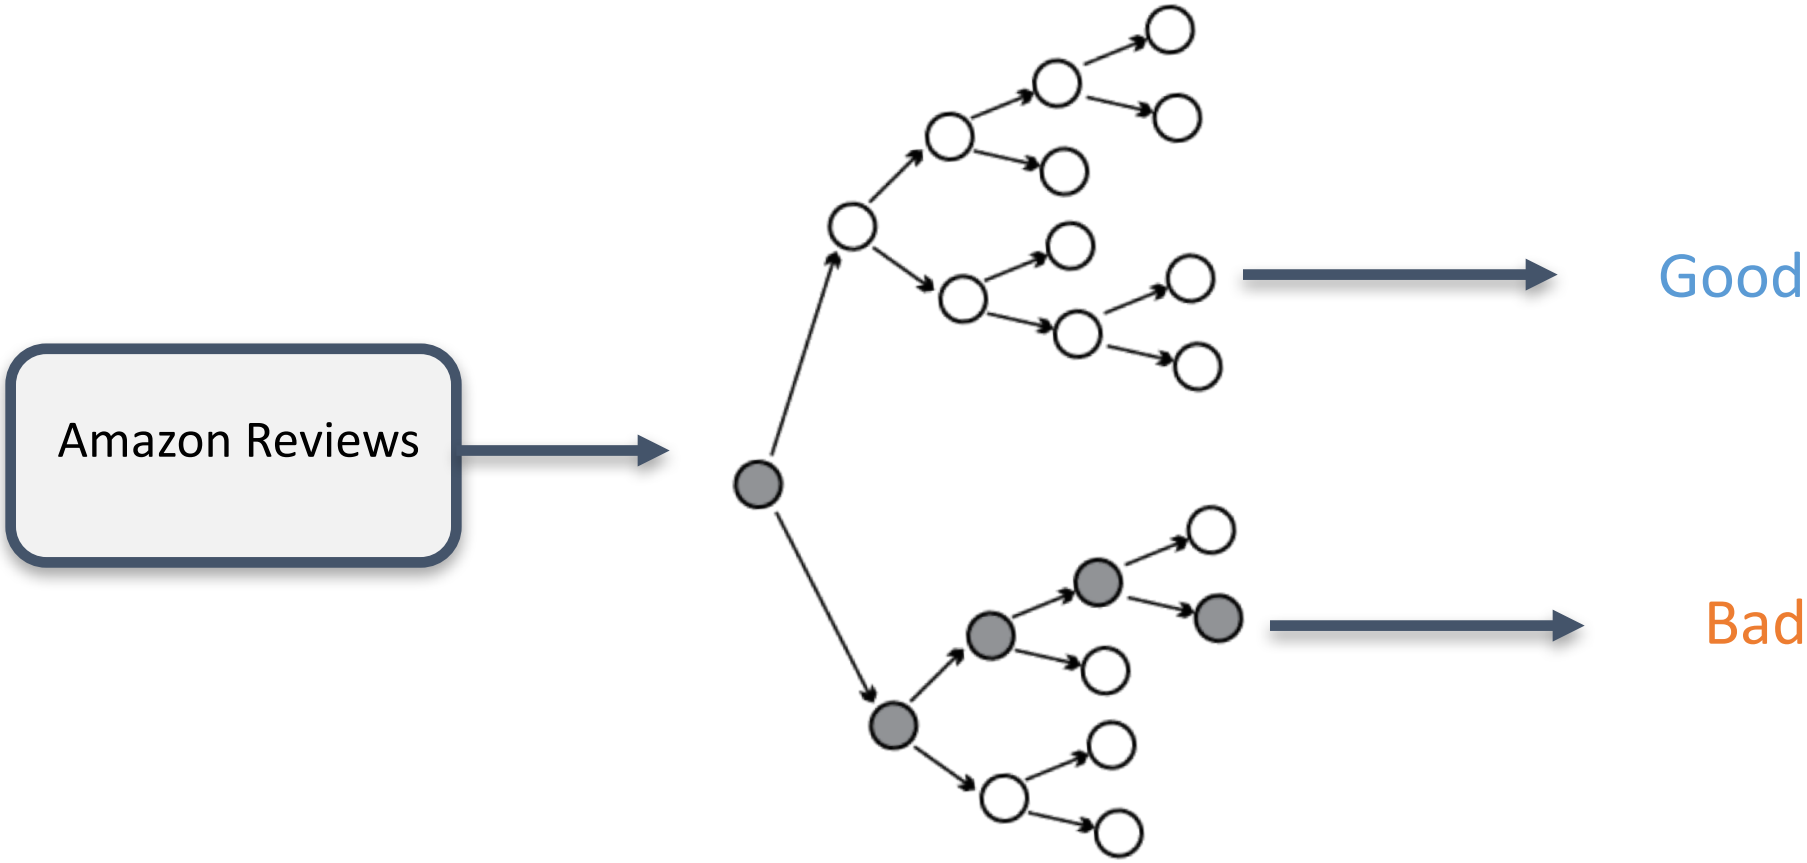

Demo

Workflow

Use Case: Predict whether a food product has a good rating based on the review.

1. Tokenize Reviews
   - Break up reviews into separate words
   - Filter words: remove stop words like “the” and “if”
2. Train a Word2Vec Model
3. Use model to transform reviews to vectors
4. Train a supervised learning model to predict good rating

End