

# Convex Optimization

**Stephen Boyd**

Electrical Engineering  
Computer Science  
Management Science and Engineering  
Institute for Computational Mathematics & Engineering  
Stanford University

H2O World, 12/4/17

# Outline

Mathematical Optimization

Convex Optimization

Examples

Large-Scale Distributed Optimization

Summary

# Outline

Mathematical Optimization

Convex Optimization

Examples

Large-Scale Distributed Optimization

Summary

## Optimization problem

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & g_i(x) = 0, \quad i = 1, \dots, p\end{array}$$

- ▶  $x \in \mathbf{R}^n$  is (vector) variable to be chosen
- ▶  $f_0$  is the *objective function*, to be minimized
- ▶  $f_1, \dots, f_m$  are the *inequality constraint functions*
- ▶  $g_1, \dots, g_p$  are the *equality constraint functions*
- ▶ variations: maximize objective, multiple objectives, ...

## Finding good (or best) actions

- ▶  $x$  represents some *action*, e.g.,
  - ▶ trades in a portfolio
  - ▶ airplane control surface deflections
  - ▶ schedule or assignment
  - ▶ resource allocation
  - ▶ transmitted signal
- ▶ constraints limit actions or impose conditions on outcome
- ▶ the smaller the objective  $f_0(x)$ , the better
  - ▶ total cost (or negative profit)
  - ▶ deviation from desired or target outcome
  - ▶ risk
  - ▶ fuel use

# Engineering design

- ▶  $x$  represents a design (of a circuit, device, structure, ...)
- ▶ constraints come from
  - ▶ manufacturing process
  - ▶ performance requirements
- ▶ objective  $f_0(x)$  is combination of cost, weight, power, ...

## Finding good models

- ▶  $x$  represents the *parameters* in a model
- ▶ constraints impose requirements on model parameters (e.g., nonnegativity)
- ▶ objective  $f_0(x)$  is the prediction error on some observed data (and possibly a term that penalizes model complexity)

# Inversion

- ▶  $x$  is something we want to estimate/reconstruct, given some measurement  $y$
- ▶ constraints come from prior knowledge about  $x$
- ▶ objective  $f_0(x)$  measures deviation between predicted and actual measurements



## Worst-case analysis (pessimization)

- ▶ variables are actions or parameters out of our control (and possibly under the control of an adversary)
- ▶ constraints limit the possible values of the parameters
- ▶ minimizing  $-f_0(x)$  finds *worst possible parameter values*
- ▶ if the worst possible value of  $f_0(x)$  is tolerable, you're OK
- ▶ it's good to know what the worst possible scenario can be

## Optimization-based models

- ▶ model an entity as taking actions that solve an optimization problem
  - ▶ an individual makes choices that maximize expected utility
  - ▶ an organism acts to maximize its reproductive success
  - ▶ reaction rates in a cell maximize growth
  - ▶ currents in a circuit minimize total power

## Optimization-based models

- ▶ model an entity as taking actions that solve an optimization problem
  - ▶ an individual makes choices that maximize expected utility
  - ▶ an organism acts to maximize its reproductive success
  - ▶ reaction rates in a cell maximize growth
  - ▶ currents in a circuit minimize total power
- ▶ (except the last) these are *very crude* models
- ▶ and yet, they often work very well

## Summary

- ▶ **summary:** optimization arises *everywhere*

## Summary

- ▶ **summary:** optimization arises *everywhere*
- ▶ **the bad news:** most optimization problems are *intractable*  
*i.e., we cannot solve them*

## Summary

- ▶ **summary:** optimization arises *everywhere*
- ▶ **the bad news:** most optimization problems are *intractable*  
*i.e., we cannot solve them*
- ▶ **an exception:** *convex optimization problems are tractable*  
*i.e., we (generally) can solve them*

# Outline

Mathematical Optimization

**Convex Optimization**

Examples

Large-Scale Distributed Optimization

Summary

## Convex optimization

convex optimization problem:

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & Ax = b\end{array}$$

- ▶ variable  $x \in \mathbf{R}^n$
- ▶ equality constraints are linear
- ▶  $f_0, \dots, f_m$  are **convex**: for  $\theta \in [0, 1]$ ,

$$f_i(\theta x + (1 - \theta)y) \leq \theta f_i(x) + (1 - \theta)f_i(y)$$

*i.e.*,  $f_i$  have nonnegative (upward) curvature



## Why

- ▶ beautiful, nearly complete theory
  - ▶ duality, optimality conditions, ...

## Why

- ▶ beautiful, nearly complete theory
  - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
  - ▶ get **global solution** (and optimality certificate)
  - ▶ polynomial complexity

## Why

- ▶ beautiful, nearly complete theory
  - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
  - ▶ get **global solution** (and optimality certificate)
  - ▶ polynomial complexity
- ▶ conceptual unification of many methods

## Why

- ▶ beautiful, nearly complete theory
  - ▶ duality, optimality conditions, ...
- ▶ effective algorithms, methods (in theory and practice)
  - ▶ get **global solution** (and optimality certificate)
  - ▶ polynomial complexity
- ▶ conceptual unification of many methods
- ▶ **lots of applications** (many more than previously thought)

## Application areas

- ▶ machine learning, statistics
- ▶ finance
- ▶ supply chain, revenue management, advertising
- ▶ control
- ▶ signal and image processing, vision
- ▶ networking
- ▶ circuit design
- ▶ combinatorial optimization
- ▶ quantum mechanics
- ▶ flux-based analysis

## The approach

- ▶ try to formulate your optimization problem as convex
- ▶ if you succeed, you can (usually) solve it (numerically)

## The approach

- ▶ try to formulate your optimization problem as convex
- ▶ if you succeed, you can (usually) solve it (numerically)
- ▶ some tricks:
  - ▶ change of variables
  - ▶ approximation of true objective, constraints
  - ▶ *relaxation*: ignore terms or constraints you can't handle

# Outline

Mathematical Optimization

Convex Optimization

Examples

Large-Scale Distributed Optimization

Summary



## Radiation treatment planning

- ▶ radiation beams with intensities  $x_j$  are directed at patient
- ▶ radiation dose  $y_i$  received in voxel  $i$
- ▶  $y = Ax$
- ▶  $A \in \mathbf{R}^{m \times n}$  comes from beam geometry, physics
- ▶ goal is to choose  $x$  to deliver prescribed radiation dose  $d_i$ 
  - ▶  $d_i = 0$  for non-tumor voxels
  - ▶  $d_i > 0$  for tumor voxels
- ▶  $y = d$  not possible, so we'll need to compromise
- ▶ typical problem has  $n = 10^3$  beams,  $m = 10^6$  voxels

## Radiation treatment planning via convex optimization


$$\begin{array}{ll}\text{minimize} & \sum_i f_i(y_i) \\ \text{subject to} & x \geq 0, \quad y = Ax\end{array}$$

- ▶ variables  $x \in \mathbf{R}^n$ ,  $y \in \mathbf{R}^m$
- ▶ objective terms are

$$f_i(y_i) = w_i^{\text{over}}(y_i - d_i)_+ + w_i^{\text{under}}(d_i - y_i)_+$$


- ▶  $w_i^{\text{over}}$  and  $w_i^{\text{under}}$  are positive weights
- ▶ i.e., we charge linearly for over- and under-dosing
- ▶ a convex optimization problem

## Example



- ▶ real patient case with  $n = 360$  beams,  $m = 360000$  voxels
- ▶ optimization-based plan essentially the same as plan used

## Example



- ▶ real patient case with  $n = 360$  beams,  $m = 360000$  voxels
- ▶ optimization-based plan essentially the same as plan used
- ▶ (but we computed the plan in a few seconds, not many hours)

## Image in-painting

- ▶ guess pixel values in obscured/corrupted parts of image
- ▶ *total variation in-painting*: choose pixel values  $x_{ij} \in \mathbf{R}^3$  to minimize *total variation*


$$\text{TV}(x) = \sum_{ij} \left\| \begin{bmatrix} x_{i+1,j} - x_{ij} \\ x_{i,j+1} - x_{ij} \end{bmatrix} \right\|_2$$

- ▶ a convex problem

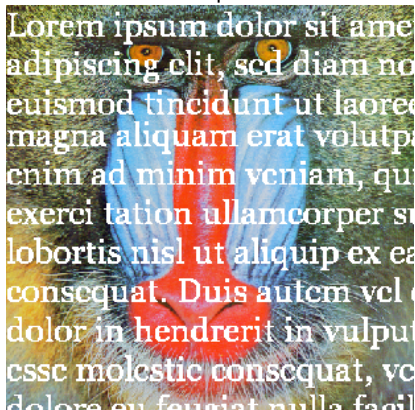
## Example

512 × 512 color image ( $n \approx 800000$  variables)

Original



Corrupted



## Example

Original



Recovered



## Support vector machine

- ▶ goal: predict a Boolean outcome from a set of  $n$  features
  - ▶ e.g., spam filter, fraud detection, customer purchase



## Support vector machine

- ▶ goal: predict a Boolean outcome from a set of  $n$  features
  - ▶ e.g., spam filter, fraud detection, customer purchase
- ▶ data  $(a_i, b_i)$ ,  $i = 1, \dots, m$ 
  - ▶  $a_i \in \mathbf{R}^n$  feature vectors;  $b_i \in \{-1, 1\}$  Boolean outcomes
- ▶ linear predictor:  $\hat{b} = \text{sign}(w^T a - v)$ 
  - ▶  $w \in \mathbf{R}^n$  is weight vector;  $v \in \mathbf{R}$  is threshold

## Support vector machine


- ▶ goal: predict a Boolean outcome from a set of  $n$  features
  - ▶ e.g., spam filter, fraud detection, customer purchase
- ▶ data  $(a_i, b_i)$ ,  $i = 1, \dots, m$ 
  - ▶  $a_i \in \mathbf{R}^n$  feature vectors;  $b_i \in \{-1, 1\}$  Boolean outcomes
- ▶ linear predictor:  $\hat{b} = \text{sign}(w^T a - v)$ 
  - ▶  $w \in \mathbf{R}^n$  is weight vector;  $v \in \mathbf{R}$  is threshold
- ▶ SVM: choose  $w, v$  to minimize (convex) objective

$$(1/m) \sum_{i=1}^m \left(1 - b_i(w^T a_i - v)\right)_+ + (\lambda/2) \|w\|_2^2$$

where  $\lambda > 0$  is parameter

# SVM

$$w^T z - v = 0 \text{ (solid); } |w^T z - v| = 1 \text{ (dashed)}$$



## Sparsity via $\ell_1$ regularization

- ▶ adding  $\ell_1$ -norm regularization

$$\lambda \|x\|_1 = \lambda(|x_1| + |x_2| + \cdots + |x_n|)$$

to objective results in **sparse**  $x$

- ▶  $\lambda > 0$  controls trade-off of sparsity versus main objective
- ▶ **preserves convexity, hence tractability**
- ▶ used for many years, in many fields
  - ▶ sparse design
  - ▶ feature selection in machine learning (lasso, SVM, ...)
  - ▶ total variation reconstruction in signal processing
  - ▶ compressed sensing

## Lasso

- ▶ regression problem with  $\ell_1$  regularization:

$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with  $A \in \mathbf{R}^{m \times n}$

- ▶ useful even when  $n \gg m$  (!!); does **feature selection**

## Lasso

- ▶ regression problem with  $\ell_1$  regularization:

$$\text{minimize} \quad (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with  $A \in \mathbf{R}^{m \times n}$

- ▶ useful even when  $n \gg m$  (!!); does **feature selection**
- ▶ *cf.*  $\ell_2$  regularization ('ridge regression'):

$$\text{minimize} \quad (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_2^2$$

## Lasso

- ▶ regression problem with  $\ell_1$  regularization:

$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_1$$

with  $A \in \mathbf{R}^{m \times n}$


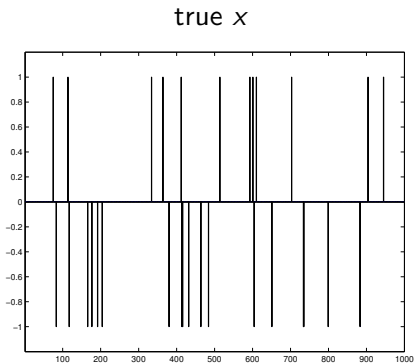
- ▶ useful even when  $n \gg m$  (!!); does **feature selection**
- ▶ cf.  $\ell_2$  regularization ('ridge regression'):

$$\text{minimize } (1/2)\|Ax - b\|_2^2 + \lambda\|x\|_2^2$$

- ▶ lasso, ridge regression have **same computational cost**

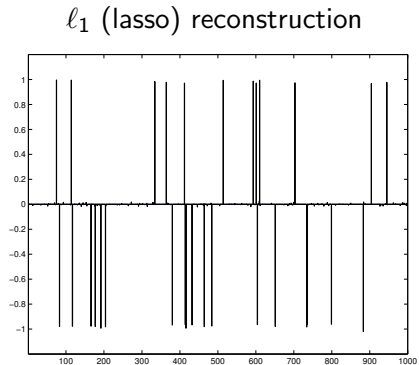
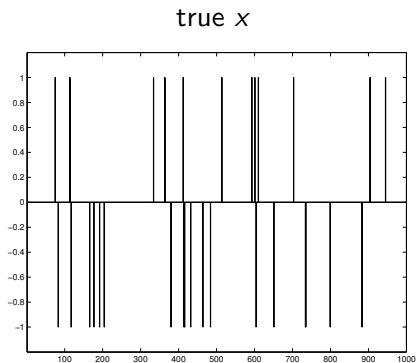
## Example

- ▶  $m = 200$  examples,  $n = 1000$  features
- ▶ examples are noisy linear measurements of true  $x$
- ▶ true  $x$  is sparse (30 nonzeros)





## Example



## State of the art — Medium scale solvers

- ▶ 1000s–10000s variables, constraints
- ▶ reliably solved by interior-point methods on single machine
- ▶ exploit problem sparsity
- ▶ not quite a technology, but getting there

## State of the art — Modeling languages

- ▶ (new) high level language support for convex optimization
  - ▶ describe problem in high level language
  - ▶ description is automatically transformed to cone problem
  - ▶ solved by standard solver, transformed back to original form

## State of the art — Modeling languages

- ▶ (new) high level language support for convex optimization
  - ▶ describe problem in high level language
  - ▶ description is automatically transformed to cone problem
  - ▶ solved by standard solver, transformed back to original form
- ▶ enables rapid prototyping (for small and medium problems)
- ▶ ideal for teaching (can do a lot with short scripts)

## CVXPY

- ▶ parser/solver written in Python (S. Diamond, 2013)
- ▶ SVM: minimize

$$(1/m) \sum_{i=1}^m \left(1 - b_i(w^T a_i - v)\right)_+ + (\lambda/2) \|w\|_2^2$$

- ▶ CVXPY specification:

```
w = Variable(n); v = Variable()    # weight, offset
losses = pos(1-mul_elemwise(b, A*w-v))
L = (1/m)*sum_entries(losses)      # avg. loss
obj = Minimize(L+(lambda/2)*sum_squares(w))
Problem(obj).solve()
```

# Outline

Mathematical Optimization

Convex Optimization

Examples

Large-Scale Distributed Optimization

Summary

## Large-scale distributed optimization

- ▶ *large-scale* optimization problems arise in many applications
  - ▶ machine learning/statistics with huge datasets
  - ▶ dynamic optimization on large-scale networks
  - ▶ image, video processing

# Large-scale distributed optimization

- ▶ *large-scale* optimization problems arise in many applications
  - ▶ machine learning/statistics with huge datasets
  - ▶ dynamic optimization on large-scale networks
  - ▶ image, video processing
- ▶ we'll use *distributed optimization*
  - ▶ split variables/constraints/objective terms among a set of *agents/processors/devices*
  - ▶ agents coordinate to solve large problem, by passing relatively small messages
  - ▶ can target modern large-scale computing platforms
  - ▶ long history, going back to 1950s



## Consensus optimization

- ▶ want to solve problem with  $N$  objective terms

$$\text{minimize } \sum_{i=1}^N f_i(x)$$

e.g.,  $f_i$  is the loss function for  $i$ th block of training data

- ▶ consensus form:

$$\begin{array}{ll} \text{minimize} & \sum_{i=1}^N f_i(x_i) \\ \text{subject to} & x_i - z = 0 \end{array}$$

- ▶  $x_i$  are **local variables**
- ▶  $z$  is the **global variable**
- ▶  $x_i - z = 0$  are **consistency** or **consensus** constraints

## Consensus optimization via ADMM

with  $\bar{x}^k = (1/N) \sum_{i=1}^N x_i^k$  (average over local variables)


$$\begin{aligned}x_i^{k+1} &:= \operatorname{argmin}_{x_i} \left( f_i(x_i) + (\rho/2) \|x_i - \bar{x}^k + u_i^k\|_2^2 \right) \\ u_i^{k+1} &:= u_i^k + (x_i^{k+1} - \bar{x}^{k+1})\end{aligned}$$

- ▶ get **global** minimum, under very general conditions
- ▶  $u^k$  is running sum of inconsistencies (PI control)
- ▶ minimizations carried out independently and in parallel
- ▶ coordination is via averaging of local variables  $x_i$


## Example — Consensus SVM

- ▶ baby problem with  $n = 2$ ,  $m = 400$  to illustrate
- ▶ examples split into  $N = 20$  groups, in worst possible way:  
each group contains only positive or negative examples


## Iteration 1



## Iteration 5



## Iteration 40



# Outline

Mathematical Optimization

Convex Optimization

Examples

Large-Scale Distributed Optimization

Summary

## Summary

- ▶ convex optimization problems **arise in many applications**
- ▶ convex optimization problems **can be solved effectively**
  - ▶ small problems at microsecond/millisecond time scales
  - ▶ medium-scale problems using general purpose methods
  - ▶ arbitrary-scale problems using distributed optimization
- ▶ high level language support makes prototyping easy



## References

*many* researchers have worked on the topics covered

- ▶ *Convex Optimization* (Boyd & Vandenberghe)
- ▶ *CVXPY: A Python-embedded modeling language for convex optimization* (Diamond & Boyd)
- ▶ *Distributed optimization and statistical learning via the alternating direction method of multipliers* (Boyd, Parikh, Chu, Peleato, & Eckstein)

all available (with code) on-line